

# 11. Basic processor design – introduction to pipelining

---

EECS 370 – Introduction to Computer Organization  
Fall 2013

**Profs. Valeria Bertacco, Robert Dick, and Satish Narayanasamy**

EECS Department  
University of Michigan in Ann Arbor, USA

# Announcements

---

- ❑ Project 2
  - Due 25 Oct.
- ❑ Competition
  - Due 28 Oct.
  - Win a Raspberry Pi.
- ❑ Homework 4
  - Due 24 Oct.



# The problem with single-cycle datapaths

---

1 ns – Register read/write time

2 ns – ALU/adder

2 ns – memory access

0 ns – MUX, PC access, sign extend, ROM

Instr.	Get instr.	Read reg.	ALU	Mem.	Write reg.	Total
add	2 ns	1 ns	2 ns	0 ns	1 ns	6 ns
beq	2 ns	1 ns	2 ns	0 ns	0 ns	5 ns
sw	2 ns	1 ns	2 ns	2 ns	0 ns	7 ns
lw	2 ns	1 ns	2 ns	2 ns	1 ns	8 ns

## Review: execution time

---

Assume 100 instructions executed

25% of instructions are loads,  
10% of instructions are stores,  
45% of instructions are adds, and  
20% of instructions are branches.

Single-cycle execution:

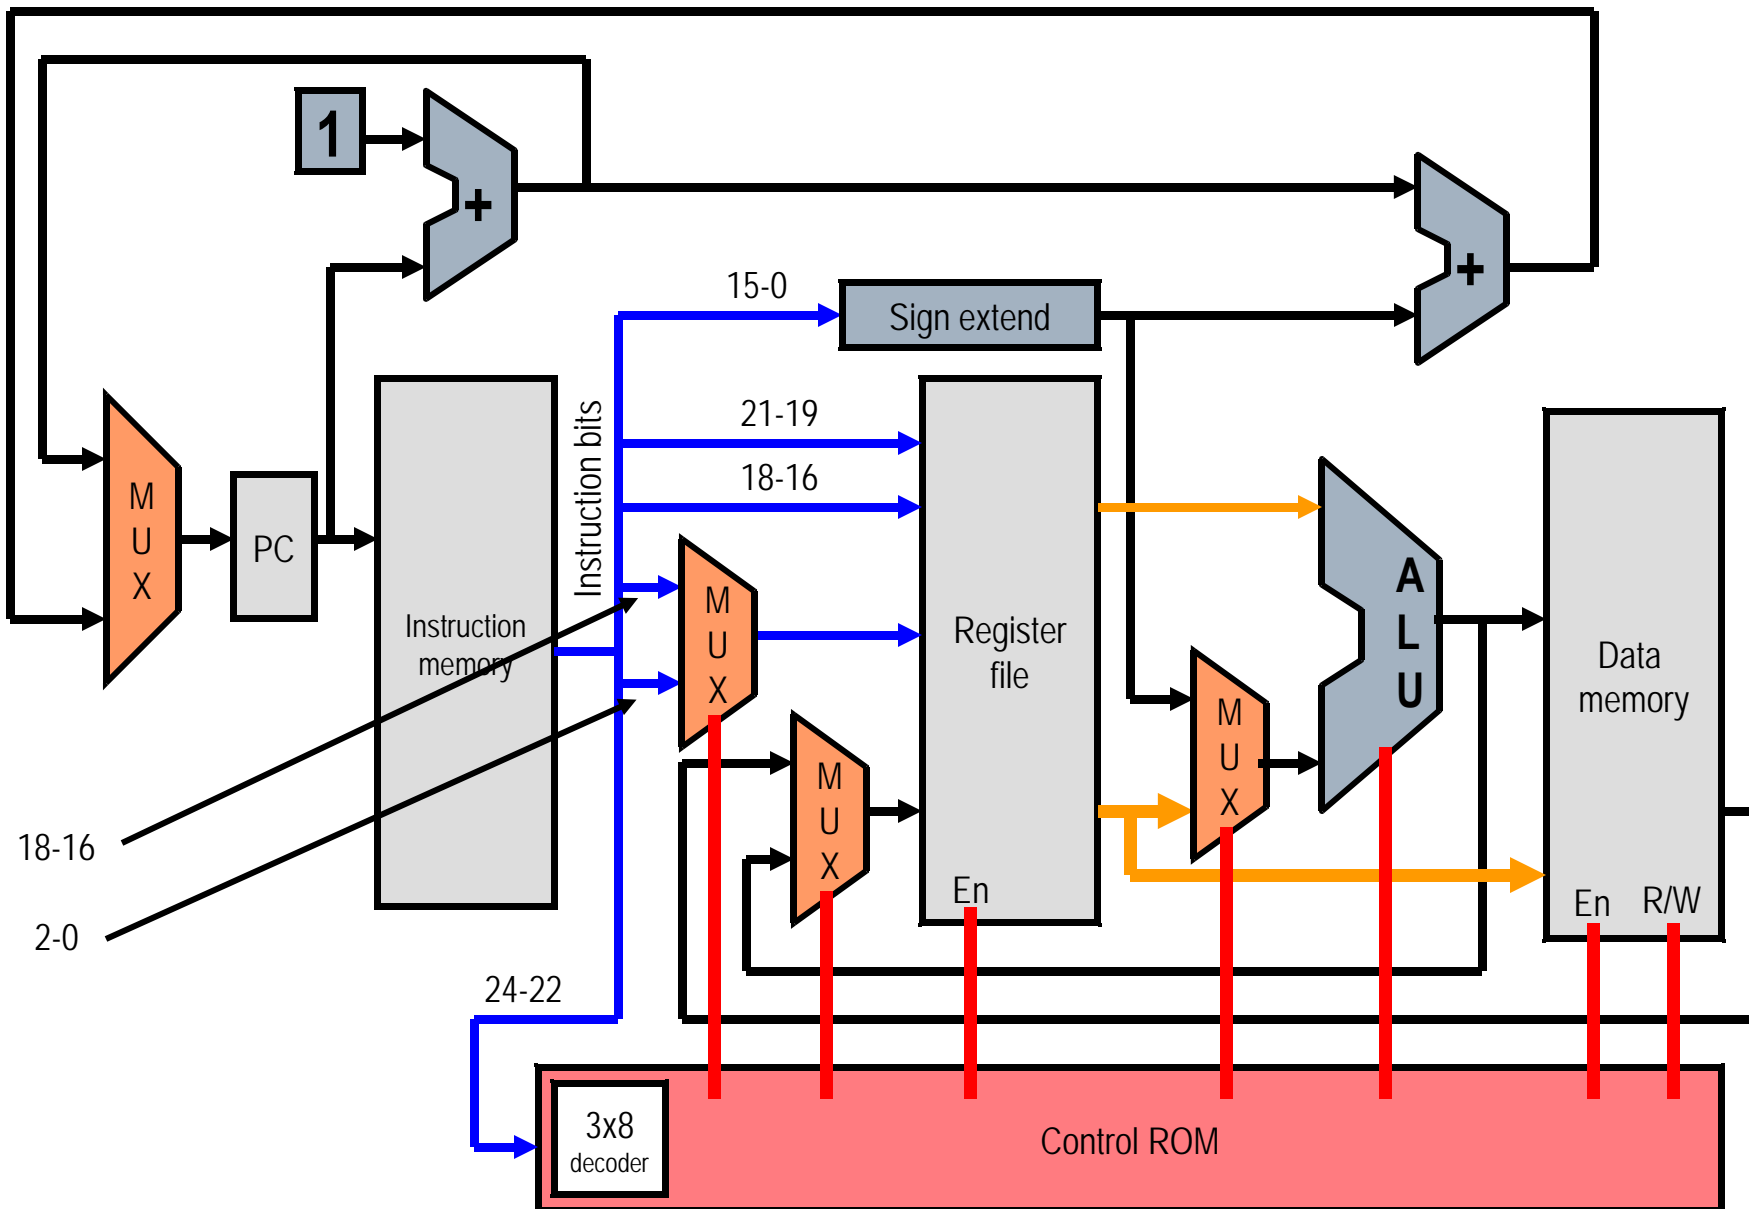
$$100 \cdot 8\text{ns} = \underline{\mathbf{800}} \text{ ns}$$

Optimal execution:

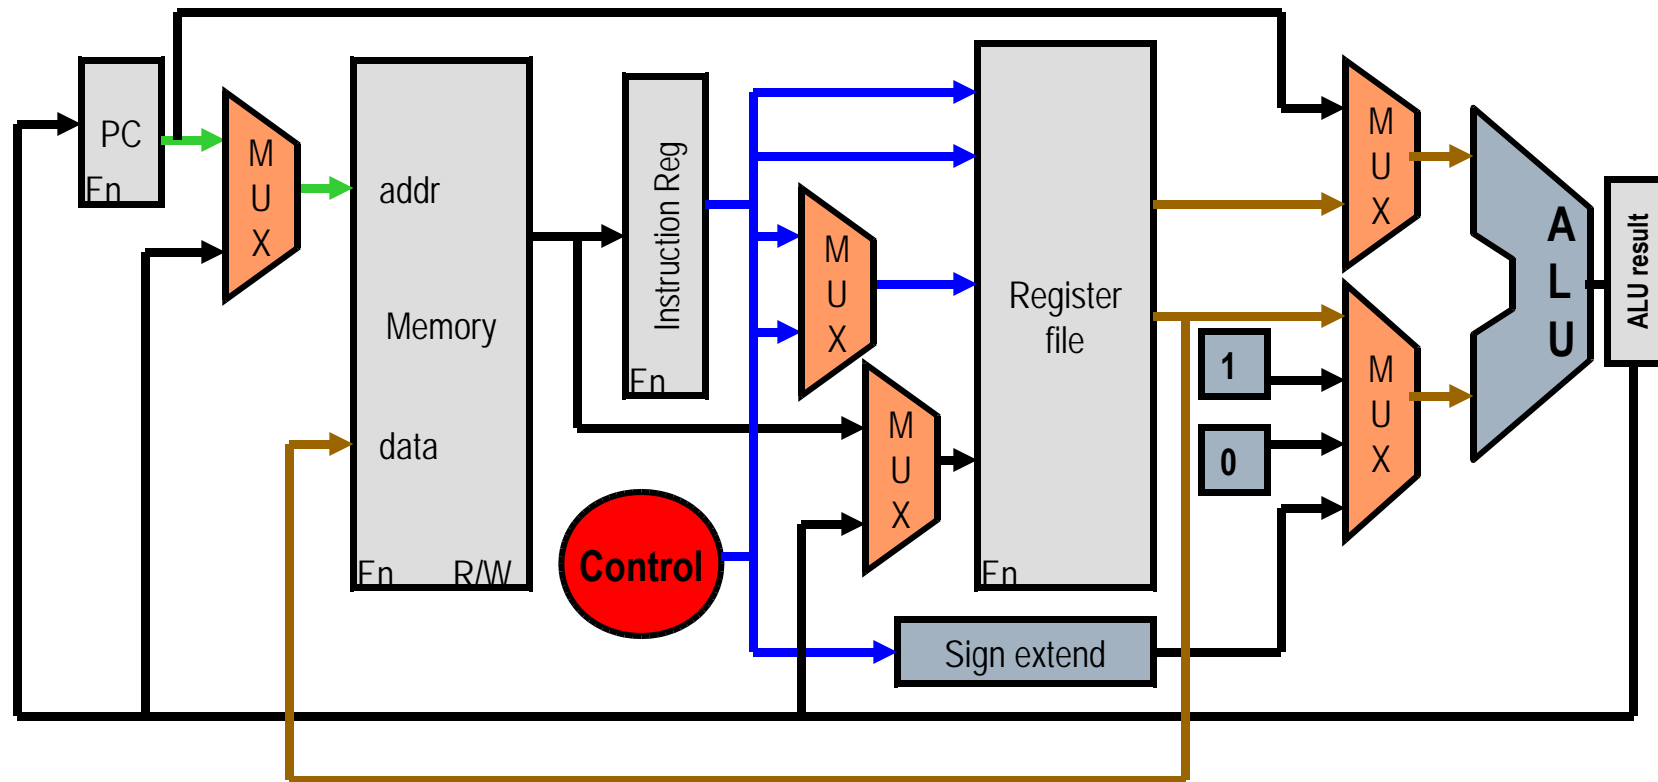
$$25 \cdot 8\text{ns} + 10 \cdot 7\text{ns} + 45 \cdot 6\text{ns} + 20 \cdot 5\text{ns} = \underline{\mathbf{640}} \text{ ns}$$

In reality, overhead to support multicycle.

# Review: single-cycle LC2Kx datapath



# Review: multi-cycle LC2Kx datapath



# Review: single-cycle and multi-cycle performance

---

Given:

2 ns – Register read/write time

1 ns – ALU/adder

2 ns – memory access

0 ns – MUX, PC access, sign extend, ROM

For program with 100 instructions:

25 lw, 10 sw, 45 add, 20 beq

For SC and MC datapaths:

What is the cycle time

How many cycles to execute?

# Performance metrics

---


1. **Response time:** From job submission to job completion, for individual job.
  - When is my job done (time)?
  - How long will this program/instruction take?
  
1. **Throughput:** In the steady state, what rate are jobs completed at?
  - How many jobs can be finished in two hours (when individual jobs are short relative to two hours)?
  - How many programs/instructions complete per hour?
  - Improved relatively easily by using multiprocessors.



# Relating response time to throughput

---

More throughput  $\leftrightarrow$  higher response time (*Little's law*)  
Works when individual jobs are short relative to response time.



# Performance metrics – execution time

---

- ❑ Response time for a program is its **execution time**.

## Execution time (for an application):

**= total instructions executed x CPI x clock period**

- Called the “Iron Law” of performance

- ❑ CPI = **avg** number of clock cycles per instruction *for an application*.

- ❑ For multi-cycle processor implementations we need

- Cycles necessary for each type of instruction.
- Mix of instructions executed in the application

(dynamic instruction execution profile)

- Cache state can also influence this.

# What are we building to?

---

- ❑ Single-cycle processor implementations
  - $\text{CPI} = ?$
  - clock period = ?
  
- ❑ Multi-cycle processor implementations
  - $\text{CPI} = ?$
  - clock period = ?
  
- ❑ Next step: improve CPI without much hurting clock period
  - Work on different parts of multiple instructions at the same time.

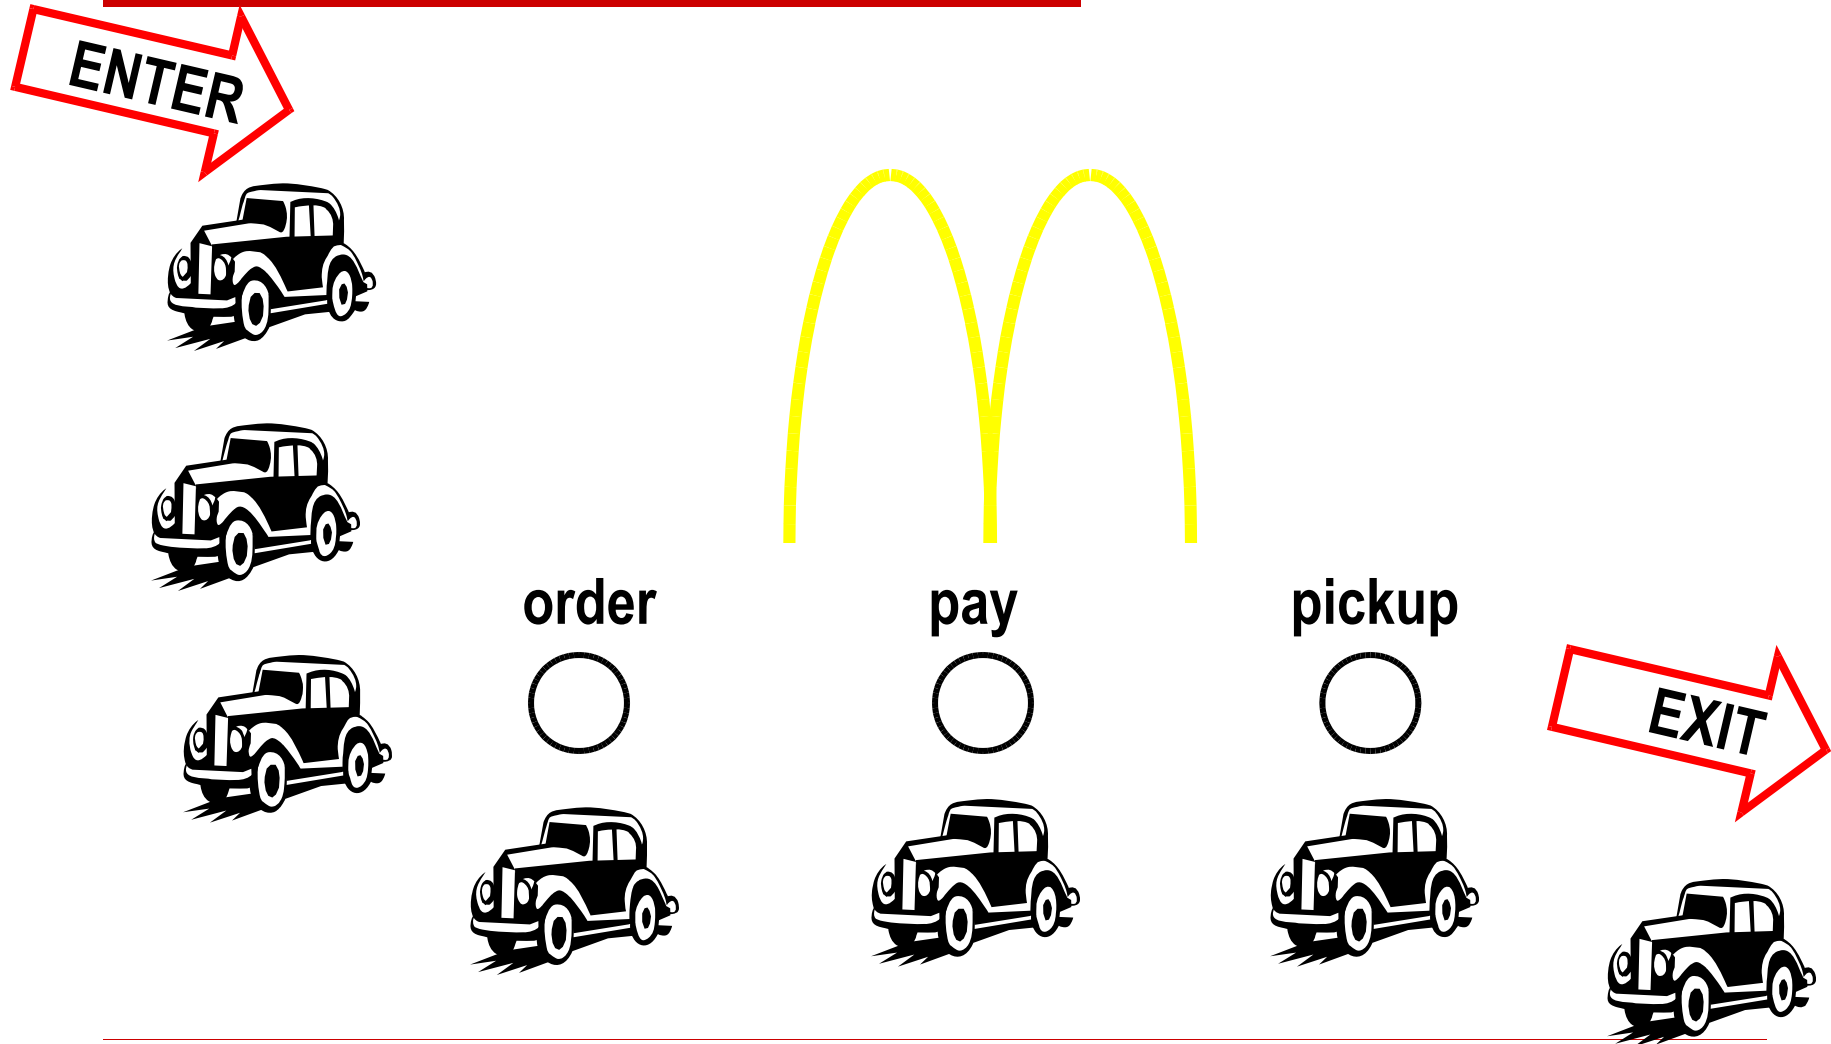
# Pipelining

---


- ❑ Want to execute an instruction?
  - Build a processor (multi-cycle)
  - Find instructions
  - Line up instructions (1, 2, 3, ...)
  - Overlap execution
    - Cycle #1: Fetch 1
    - Cycle #2: Decode 1    Fetch 2
    - Cycle #3: ALU 1       Decode 2       Fetch 3
    - . . . . .
  - This is called pipelining instruction execution.
  - Used extensively for the first time on IBM 360 (1960s).
  - CPI approaches 1.

# Pipelining

---



# Pipelining



# Pipelining trends

---

- ❑ Execute numerous instructions at the same time.
  - Pipelining: 12-20+ cycles.
  - Multiple pipelines.
- ❑ Pentium
  - 2 pipelines, 5 cycles each (10 instructions “in flight”).
- ❑ Pentium Pro/II/III
  - Roughly 3 pipelines, 12 cycles deep.
  - Instructions can execute out of their original program order.
- ❑ Pentium IV
  - 4 pipelines, 20 cycles deep.

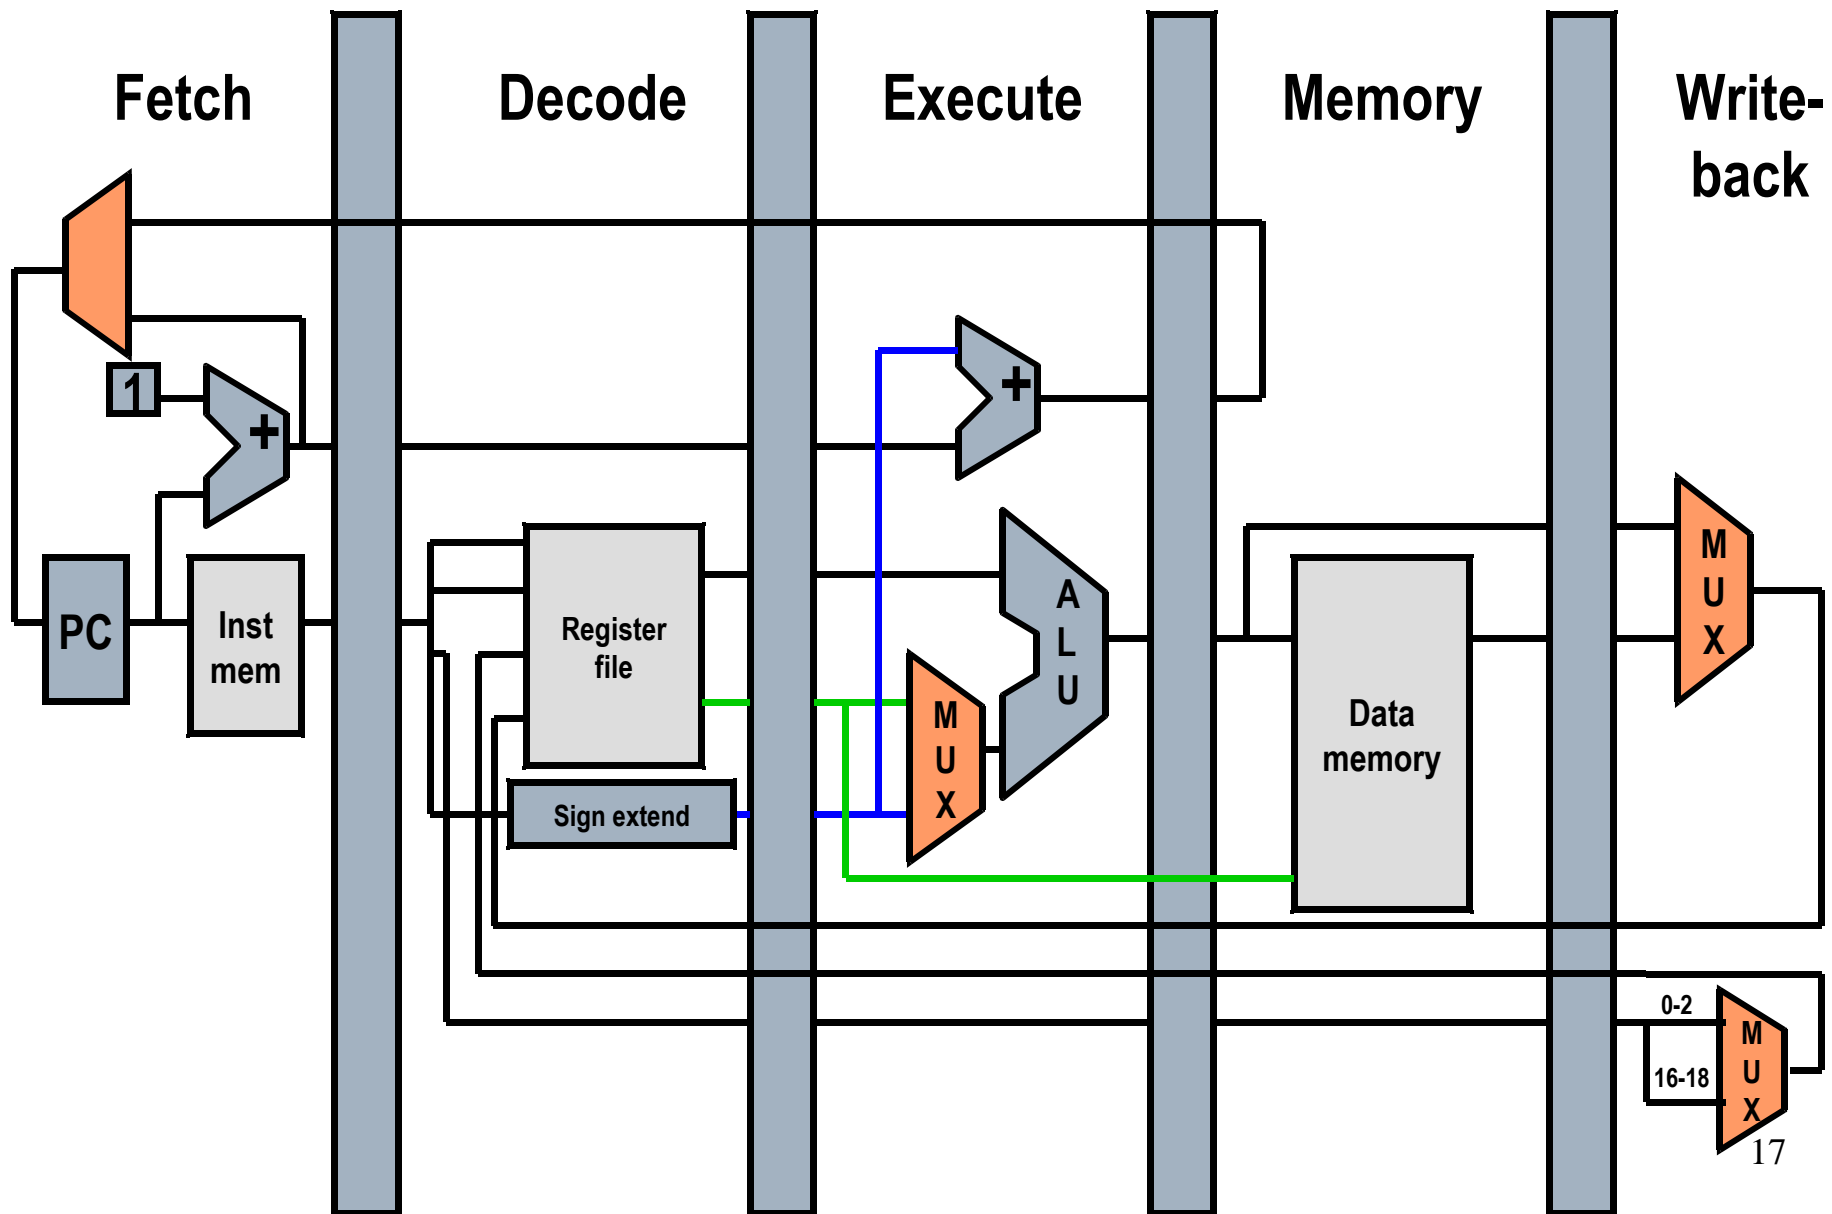
# Pipelined implementation of LC2Kx

---

- ❑ Break the execution of the instruction into cycles.
  - Similar to the multi-cycle datapath.
- ❑ Design a separate datapath **stage** for the execution performed during each cycle.
  - Build **pipeline registers** to communicate between the stages.



# Pipelined datapath

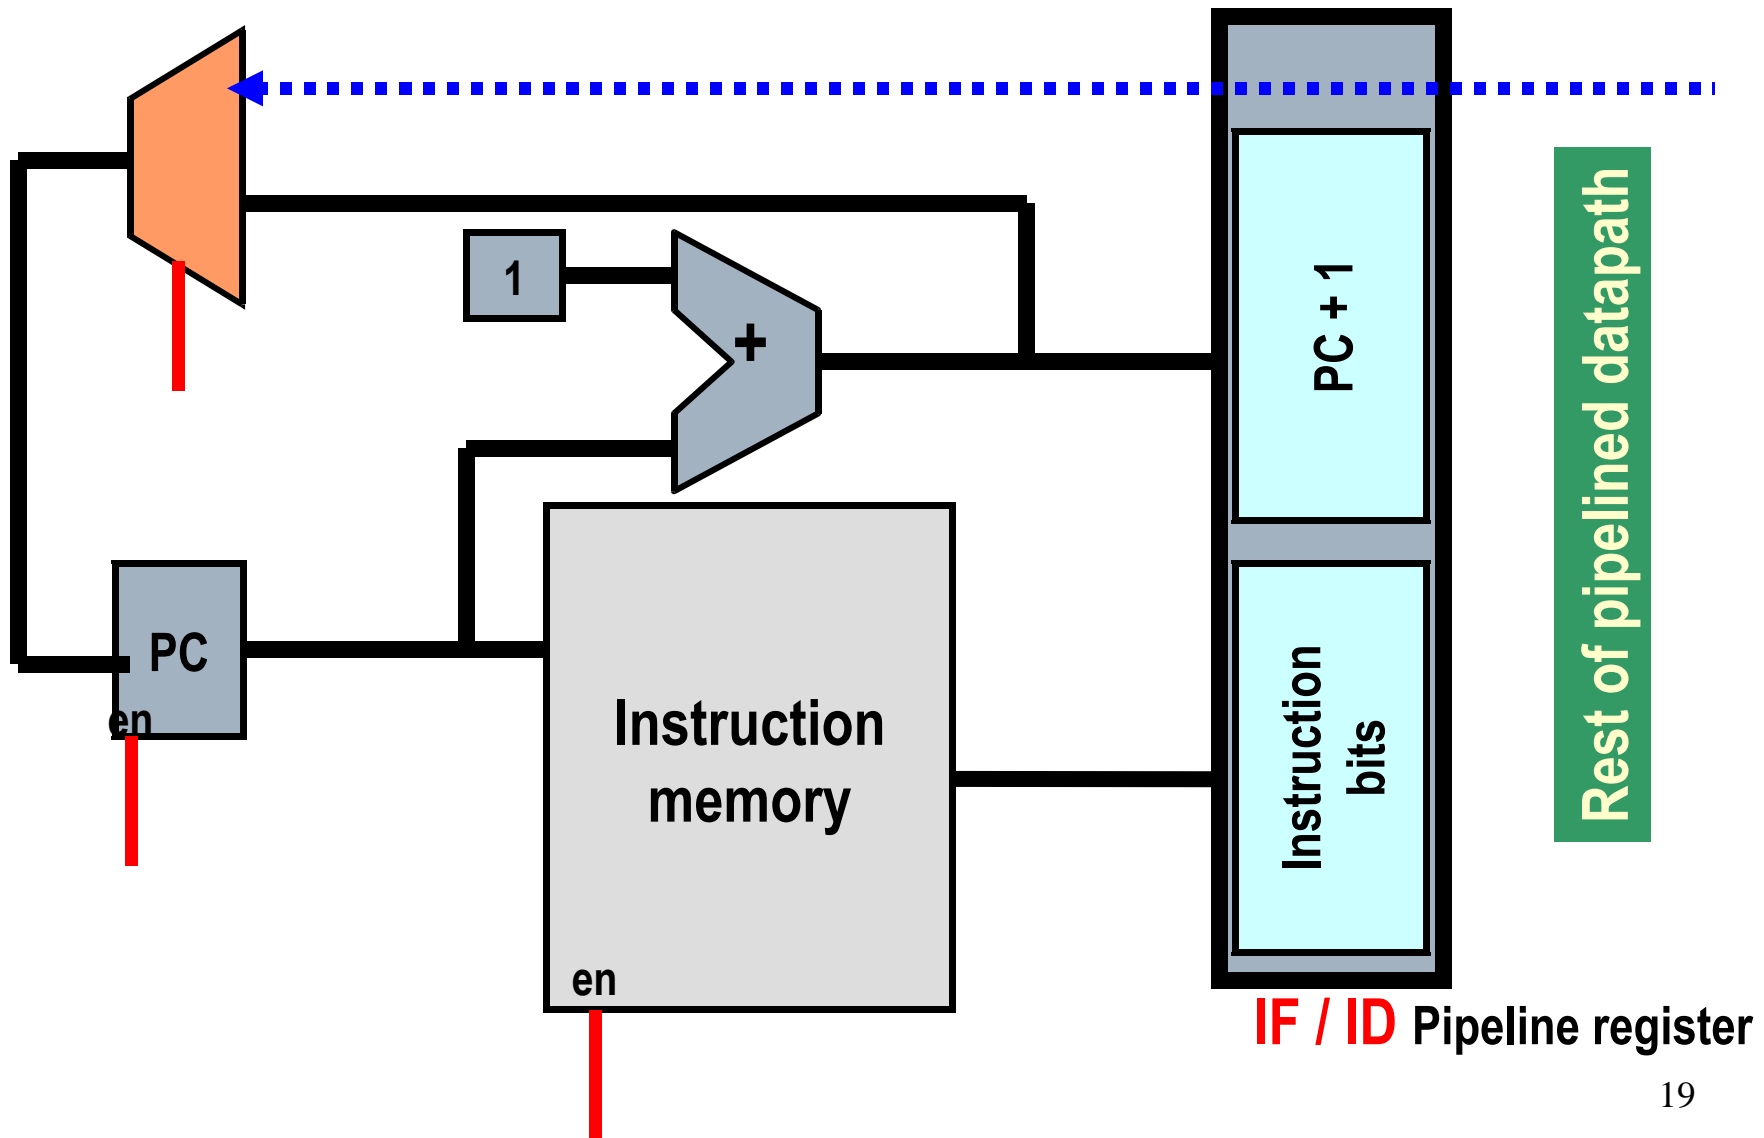


## Stage 1: fetch

---

- ❑ Design a datapath that can fetch an instruction from memory every cycle.
  - Use PC to index memory to read instruction.
  - Increment the PC (assume no branches for now).
  
- ❑ Write everything needed to complete execution to the **pipeline register (IF/ID)**
  - The next **stage** will read this pipeline register.
  - Note that pipeline register must be edge-triggered.

# Pipeline datapath – fetch stage

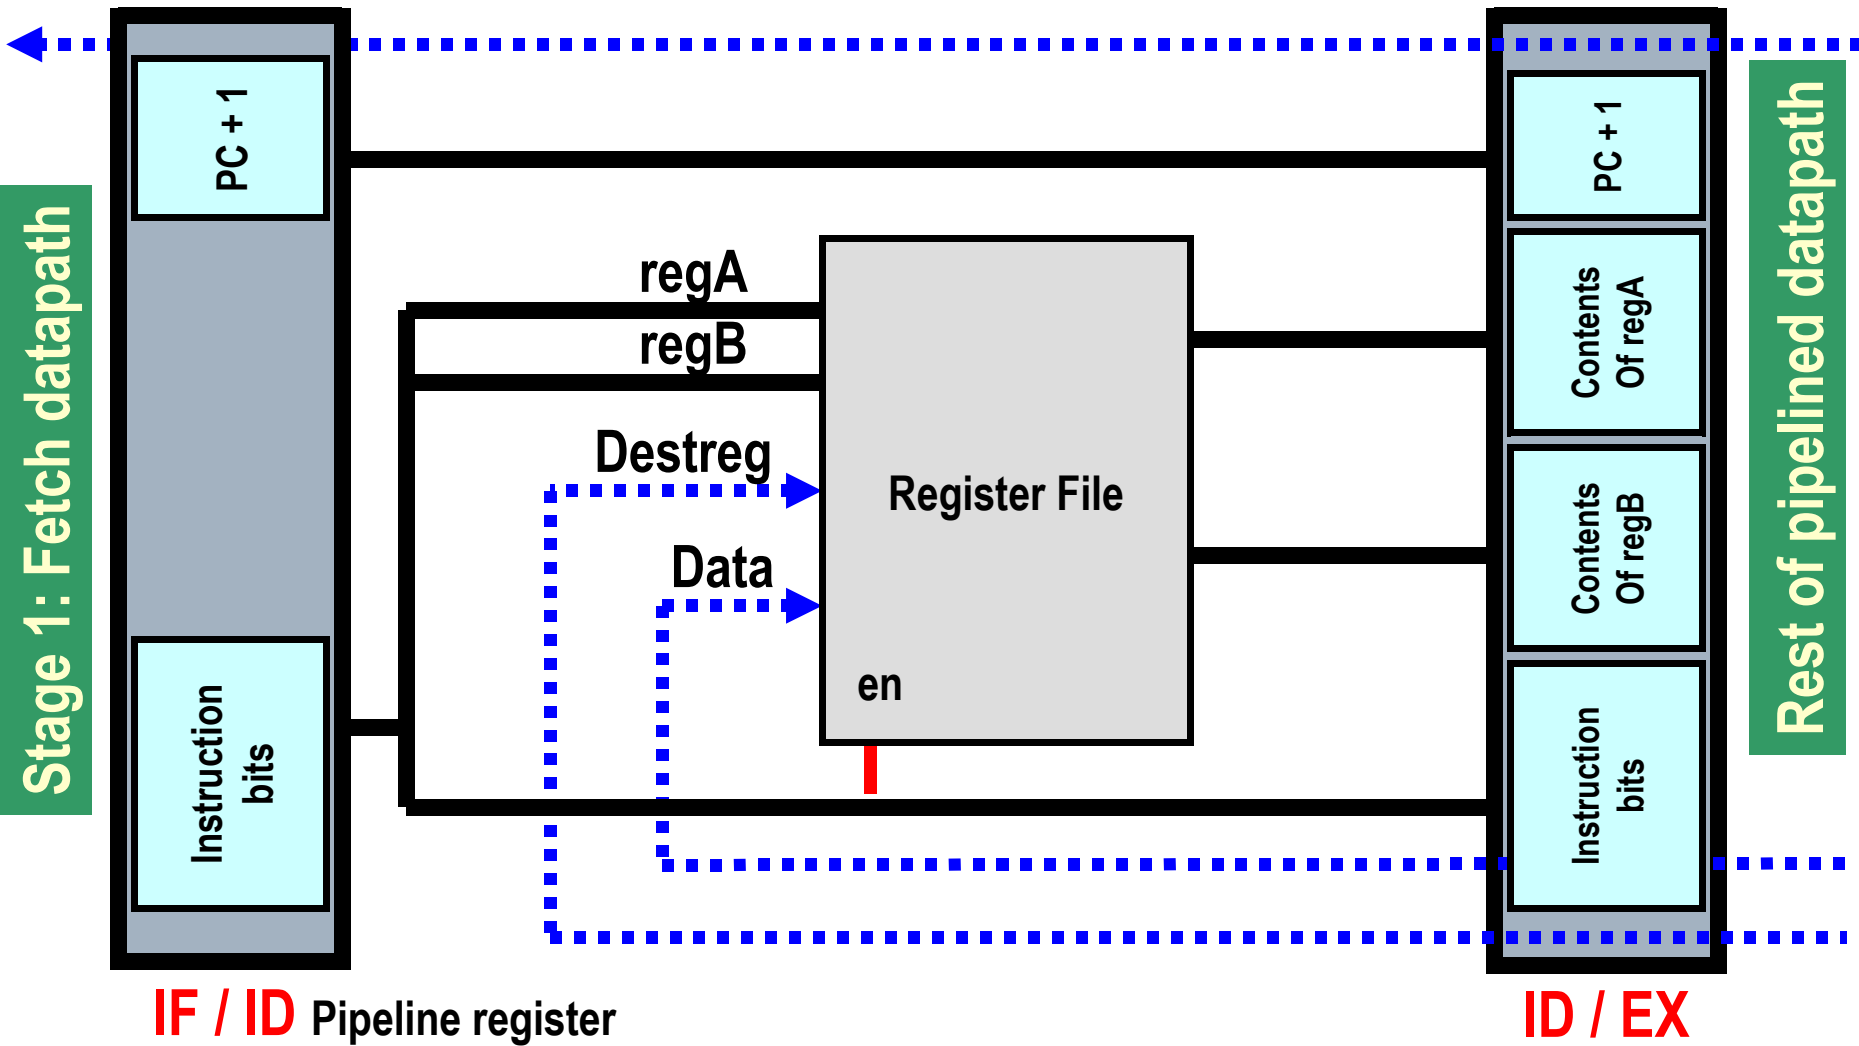


## Stage 2: decode

---

- ❑ Design a datapath that reads the IF/ID pipeline register, decodes instruction and reads register file (specified by regA and regB of instruction bits).
  - Decode is easy, just pass on the opcode and let later stages figure out their own control signals for the instruction.
  
- ❑ Write everything needed to complete execution to the **pipeline register (ID/EX)**.
  - Pass on the offset field and both destination register specifiers (or simply pass on the whole instruction!).
  - Including PC+1 even though decode didn't use it.

# Pipeline datapath – decode stage

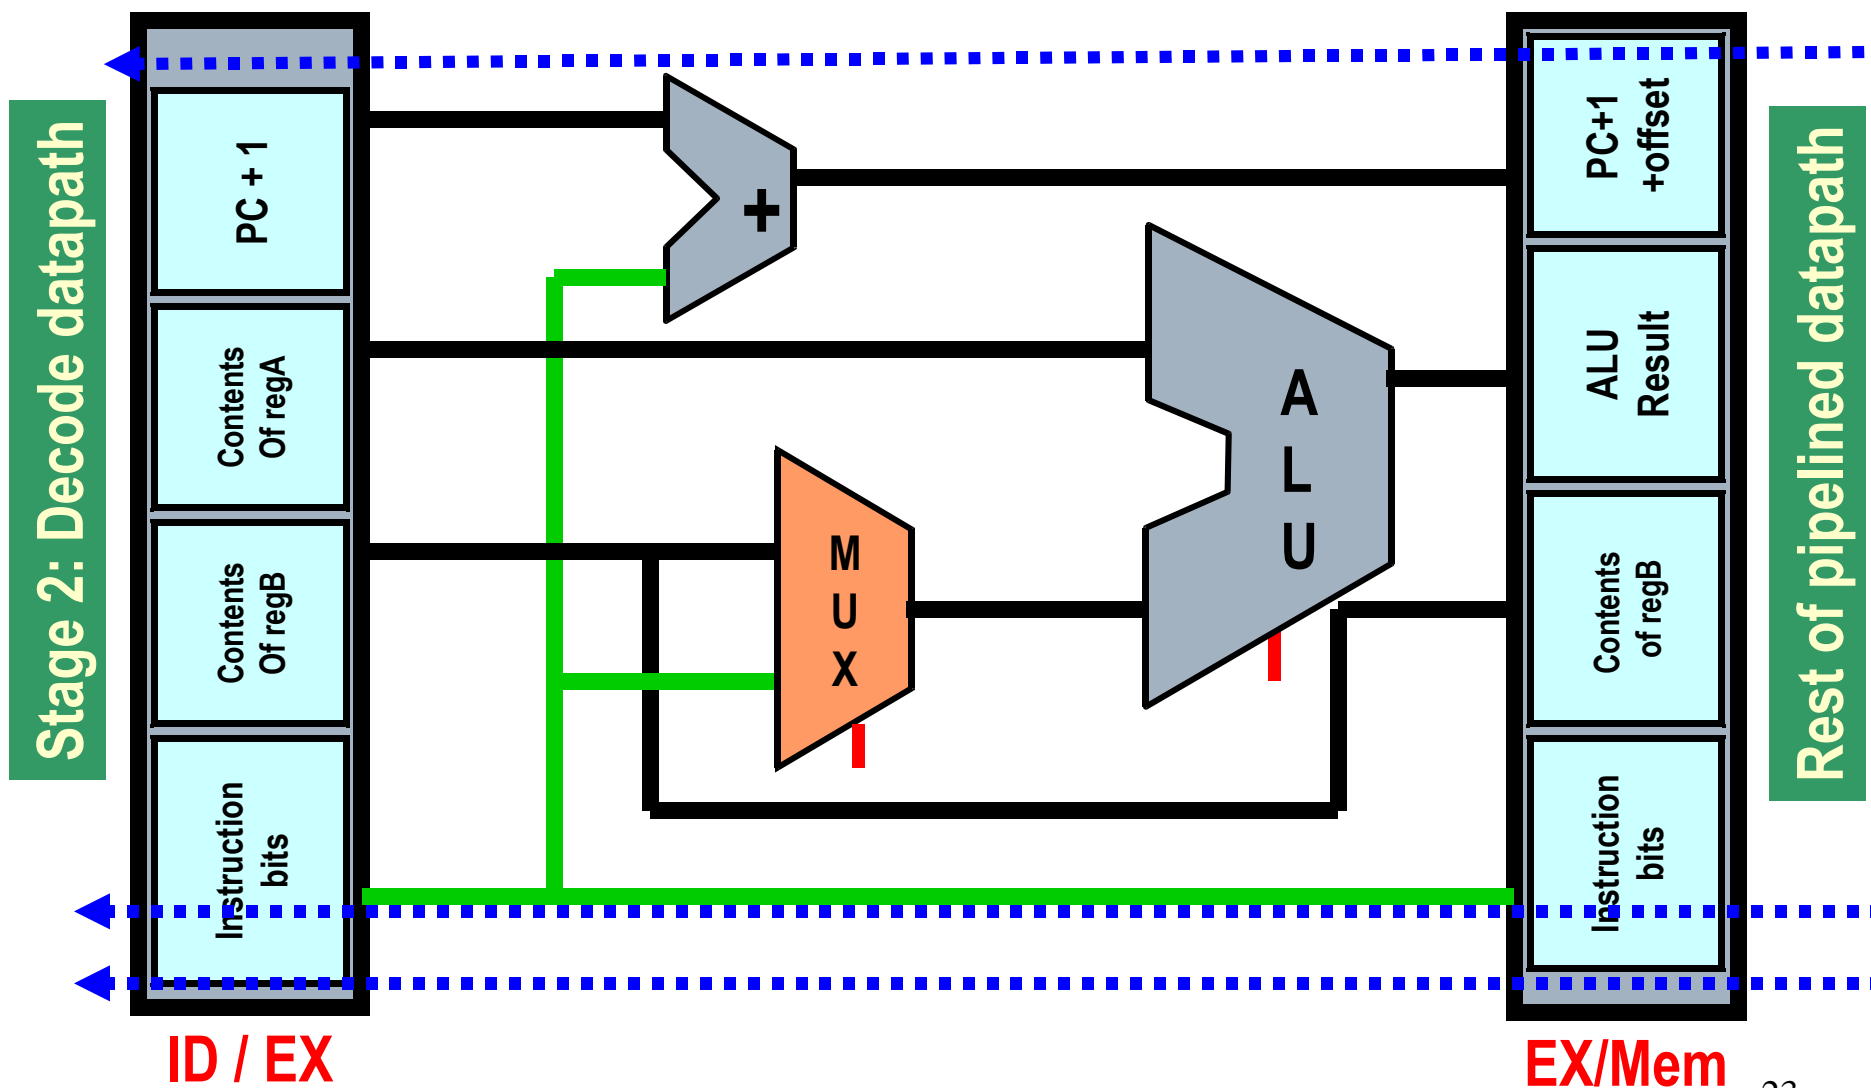


## Stage 3: execute

---

- ❑ Design a datapath that performs the proper ALU operation for the instruction specified and the values present in the ID/EX pipeline register.
  - The inputs are the contents of regA and either the contents of regB or the offset field on the instruction.
  - Also, calculate  $PC+1+offset$  in case this is a branch.
  
- ❑ Write everything needed to complete execution to the pipeline register (EX/Mem).
  - ALU result, contents of regB and  $PC+1+offset$ .
  - Instruction bits for opcode and destReg specifiers.
  - Result from comparison of regA and regB contents.

## Pipeline datapath – execute stage



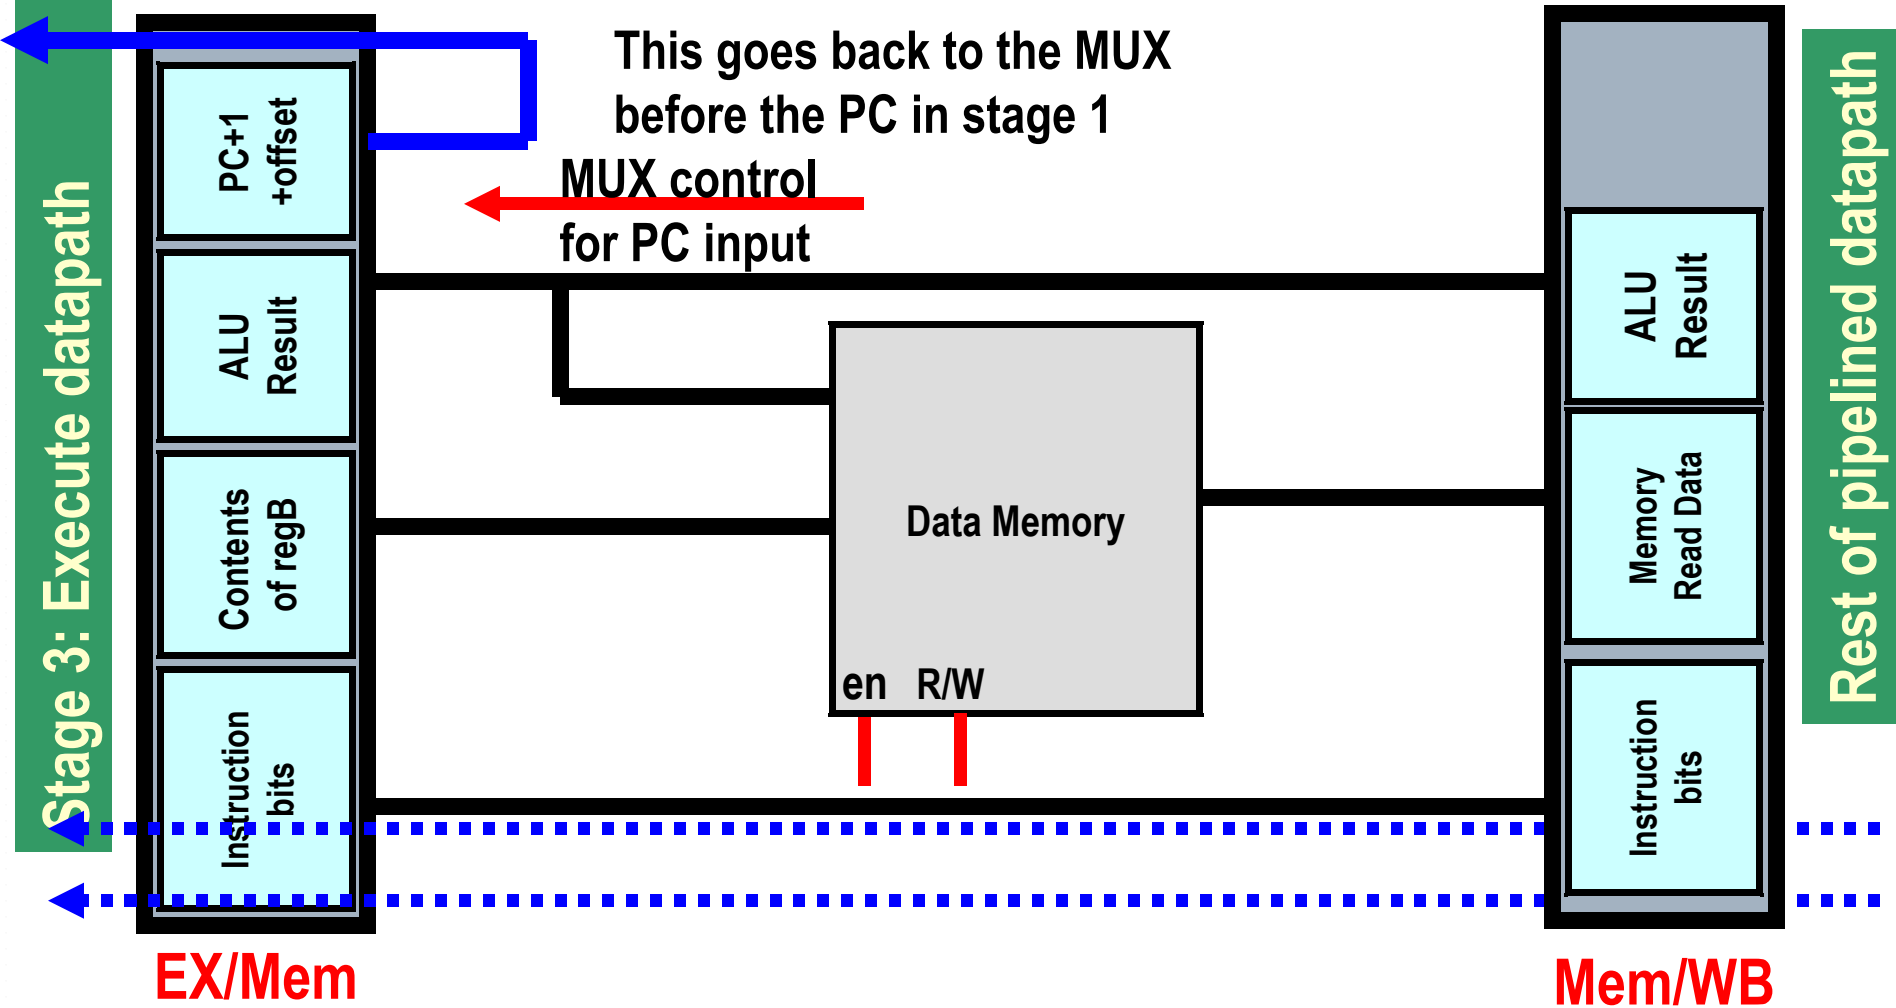
## Stage 4: memory operation

---

- ❑ Design a datapath that performs the proper memory operation for the instruction specified and the values present in the EX/Mem pipeline register.
  - ALU result contains address for **ld** and **st** instructions.
  - Opcode bits control memory R/W and enable signals.
  
- ❑ Write everything needed to complete execution to the **pipeline register (Mem/WB)**.
  - ALU result and MemData.
  - Instruction bits for opcode and destReg specifiers.



# Pipeline datapath – memory stage

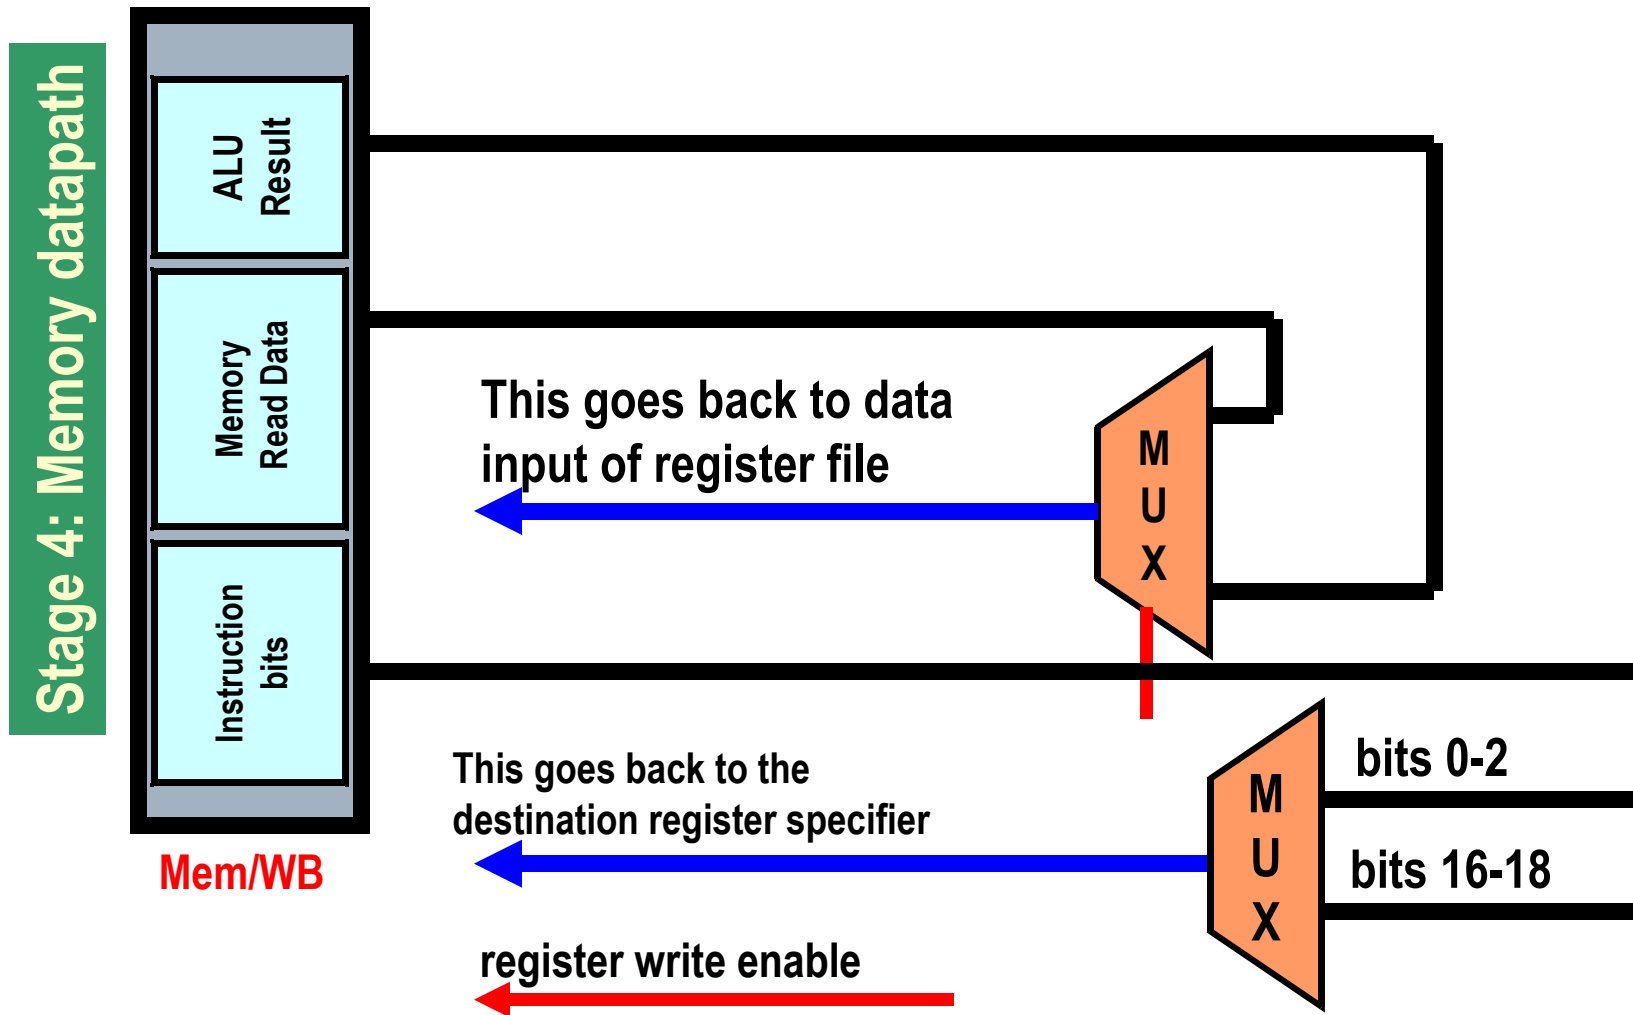


## Stage 5: write back

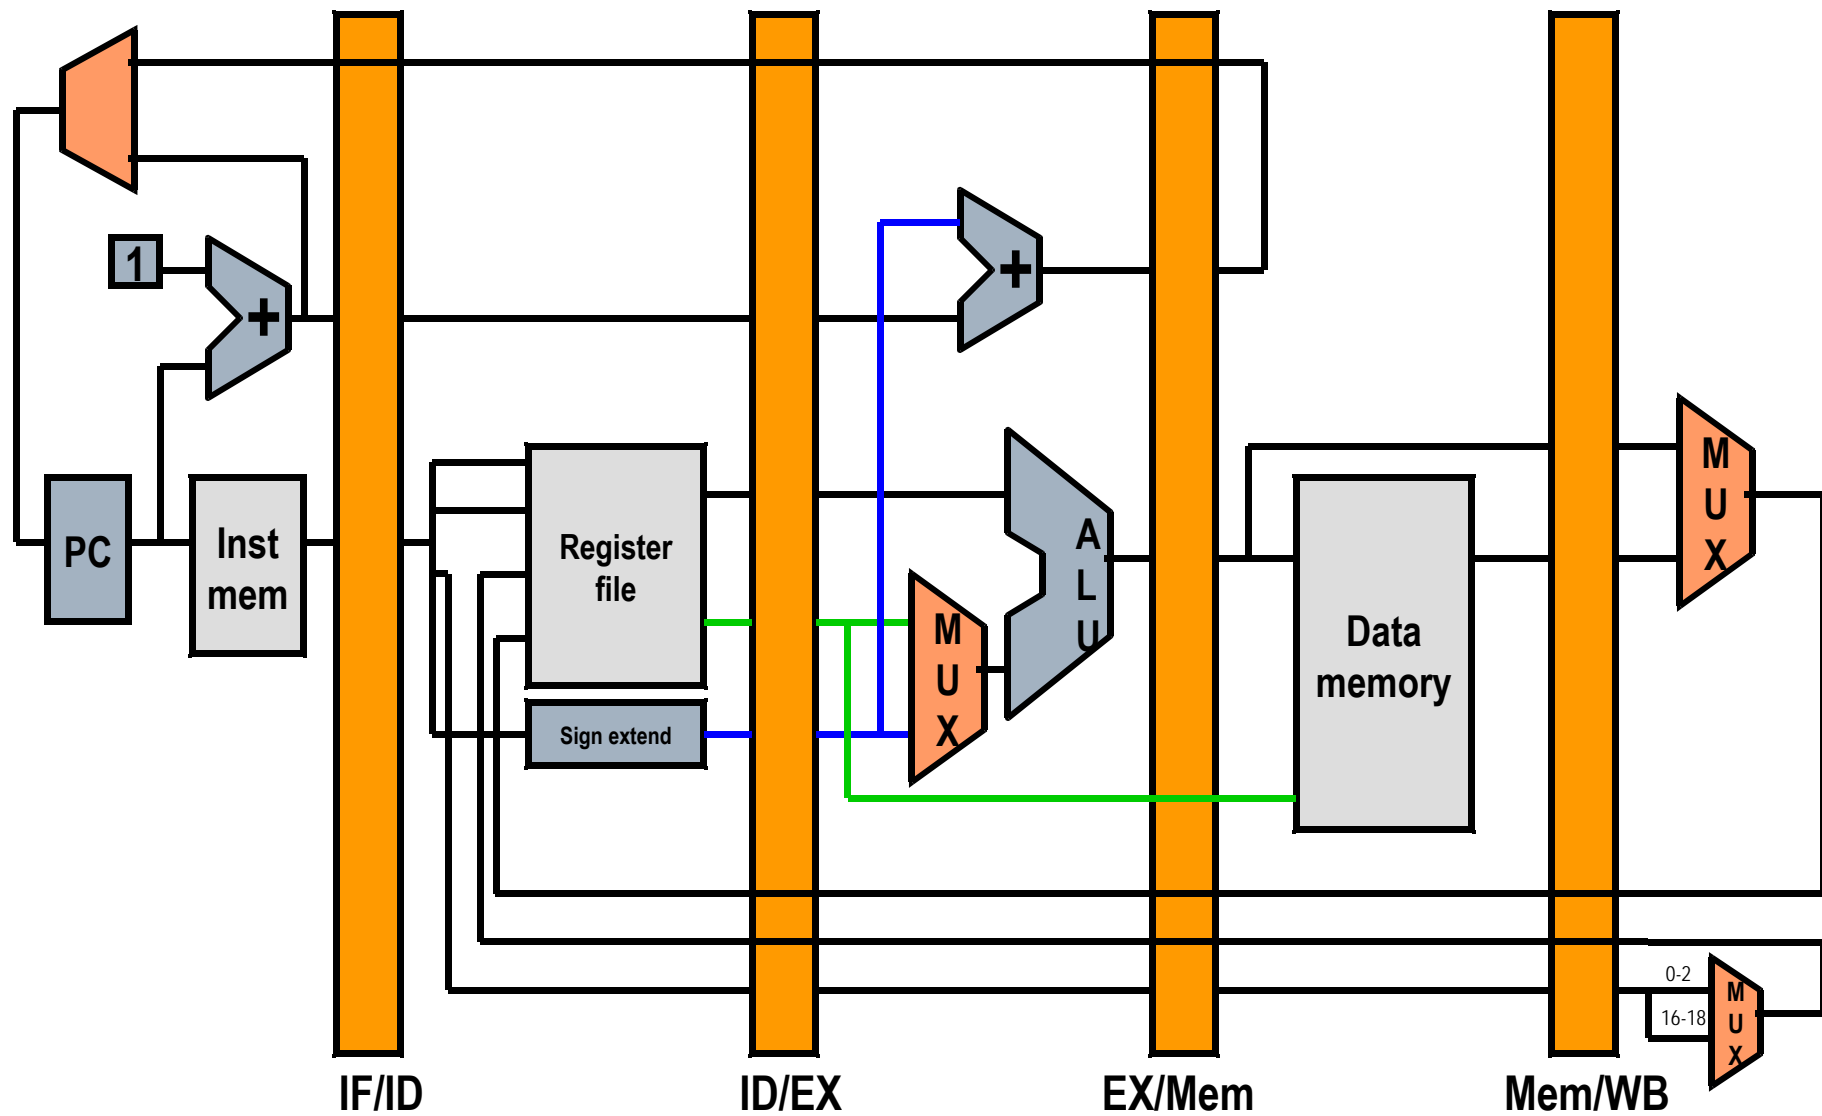
---

- ❑ Design a datapath that completes the execution of this instruction, writing to the register file if required.
  - Write MemData to destReg for ld instruction.
  - Write ALU result to destReg for add or nand instructions.
  - Opcode bits also control register write enable signal.

# Pipeline datapath – writeback stage



# Interaction among stages



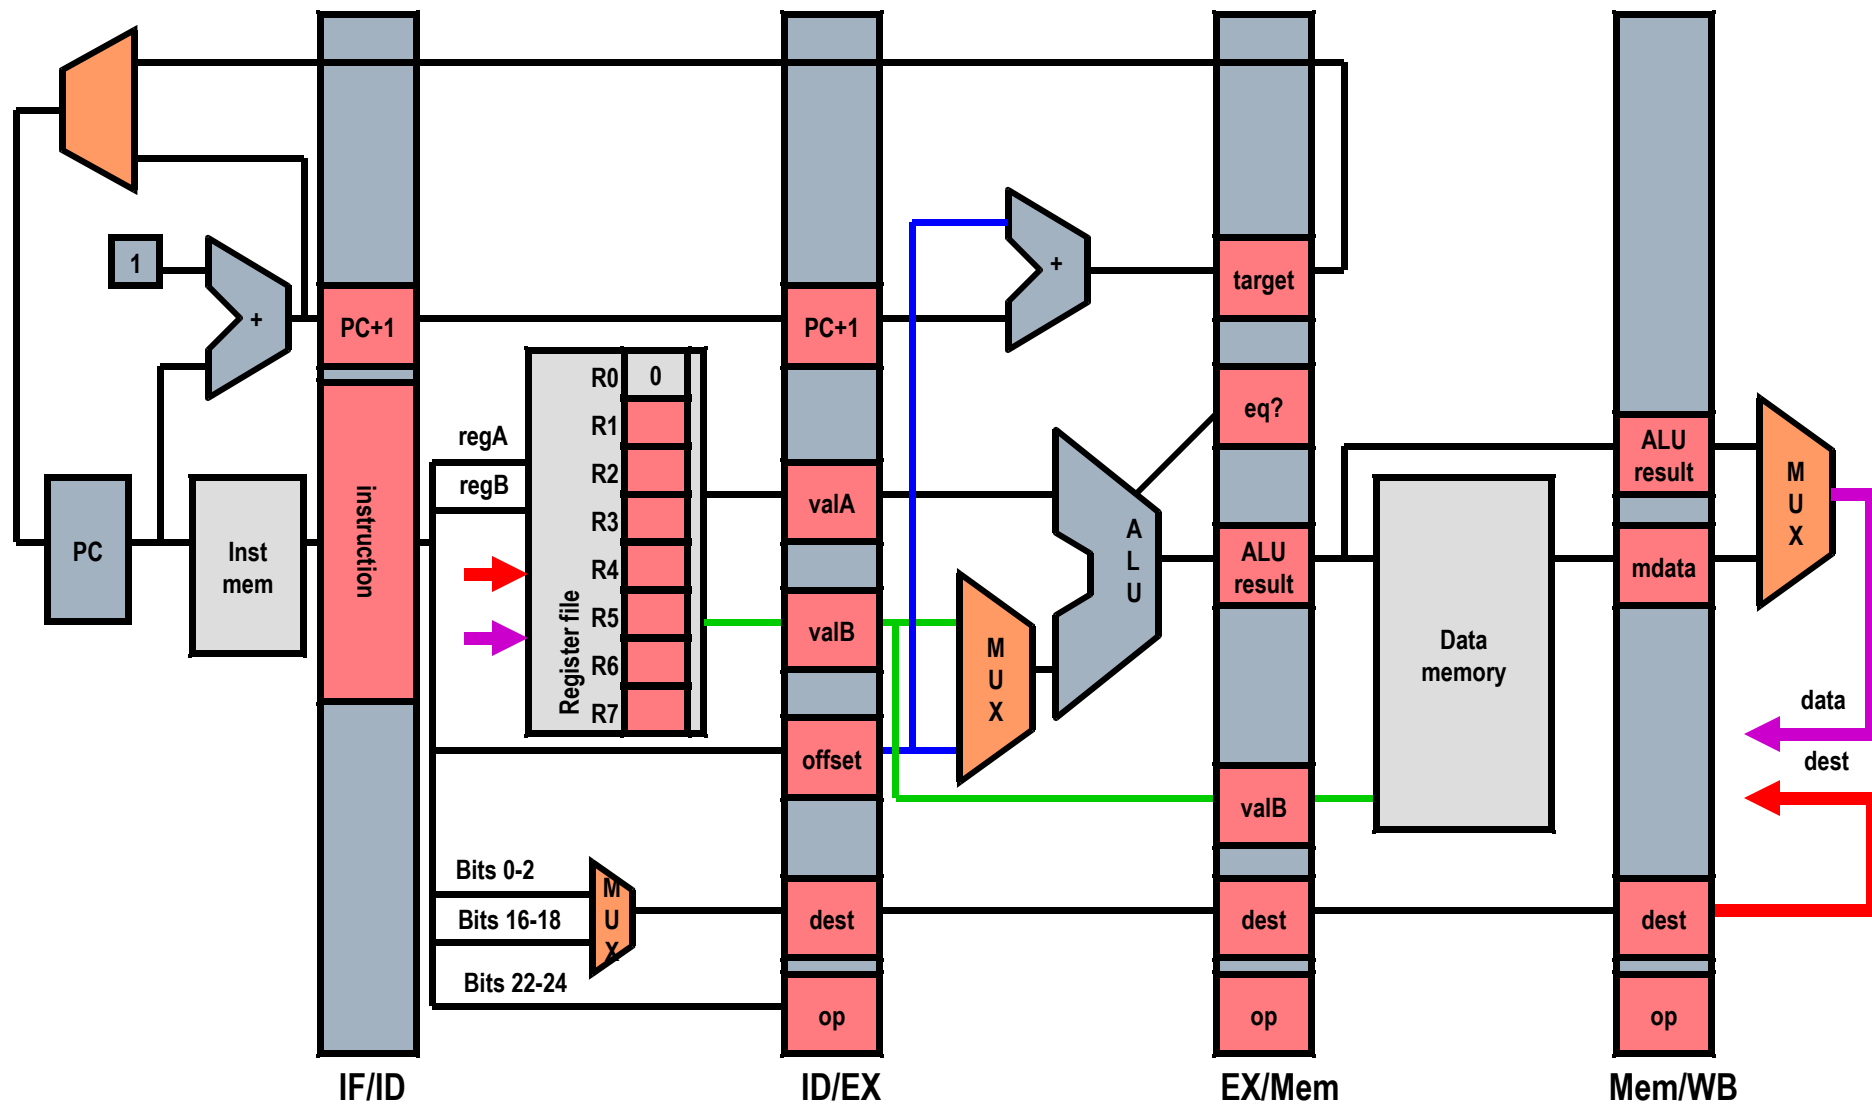
## Sample code (simple)

---

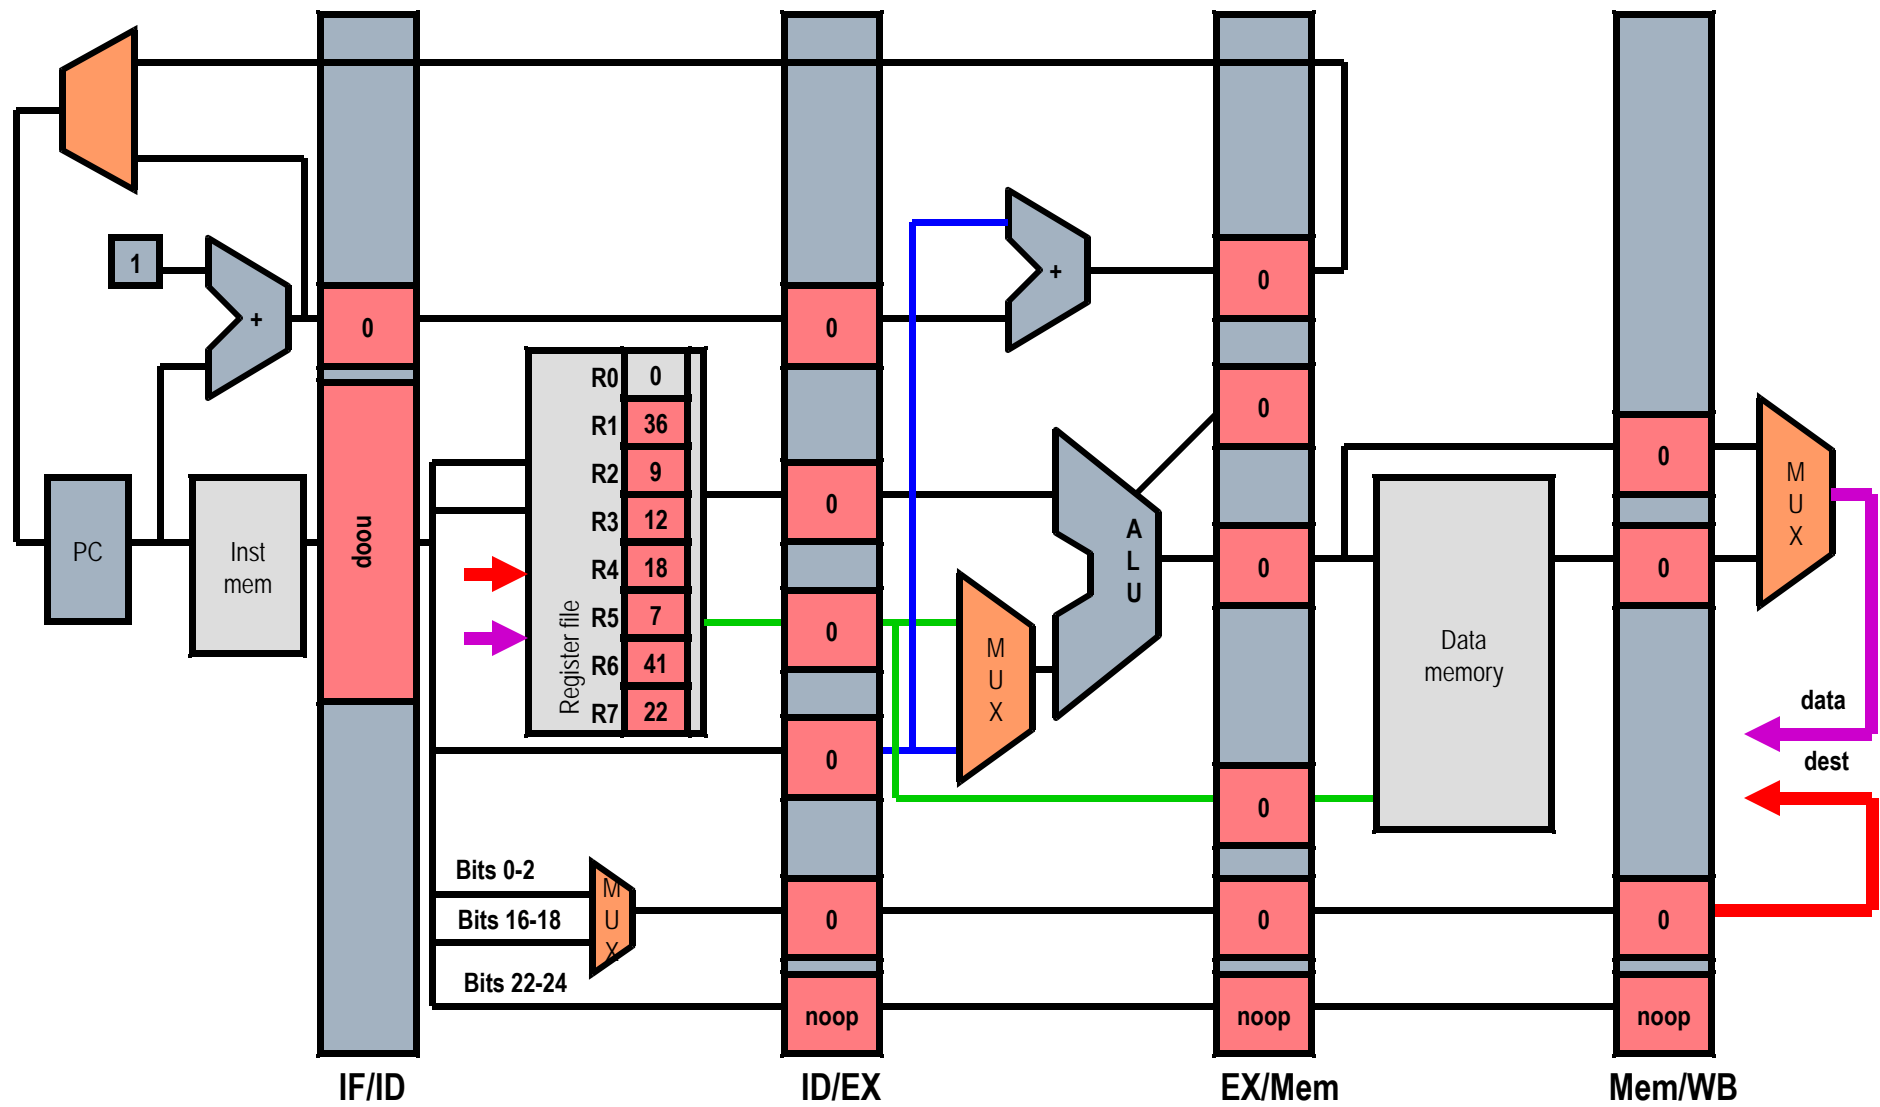
Let's run the following code on pipelined LC2K2x

- add     1   2   3     ; reg 3 = reg 1 + reg 2
- nand    4   5   6     ; reg 6 = ~(reg 4 & reg 5)
- lw      2   4   20    ; reg 4 = Mem[reg2+20]
- add     2   5   5     ; reg 5 = reg 2 + reg 5
- sw      3   7   10    ; Mem[reg3+10] = reg 7

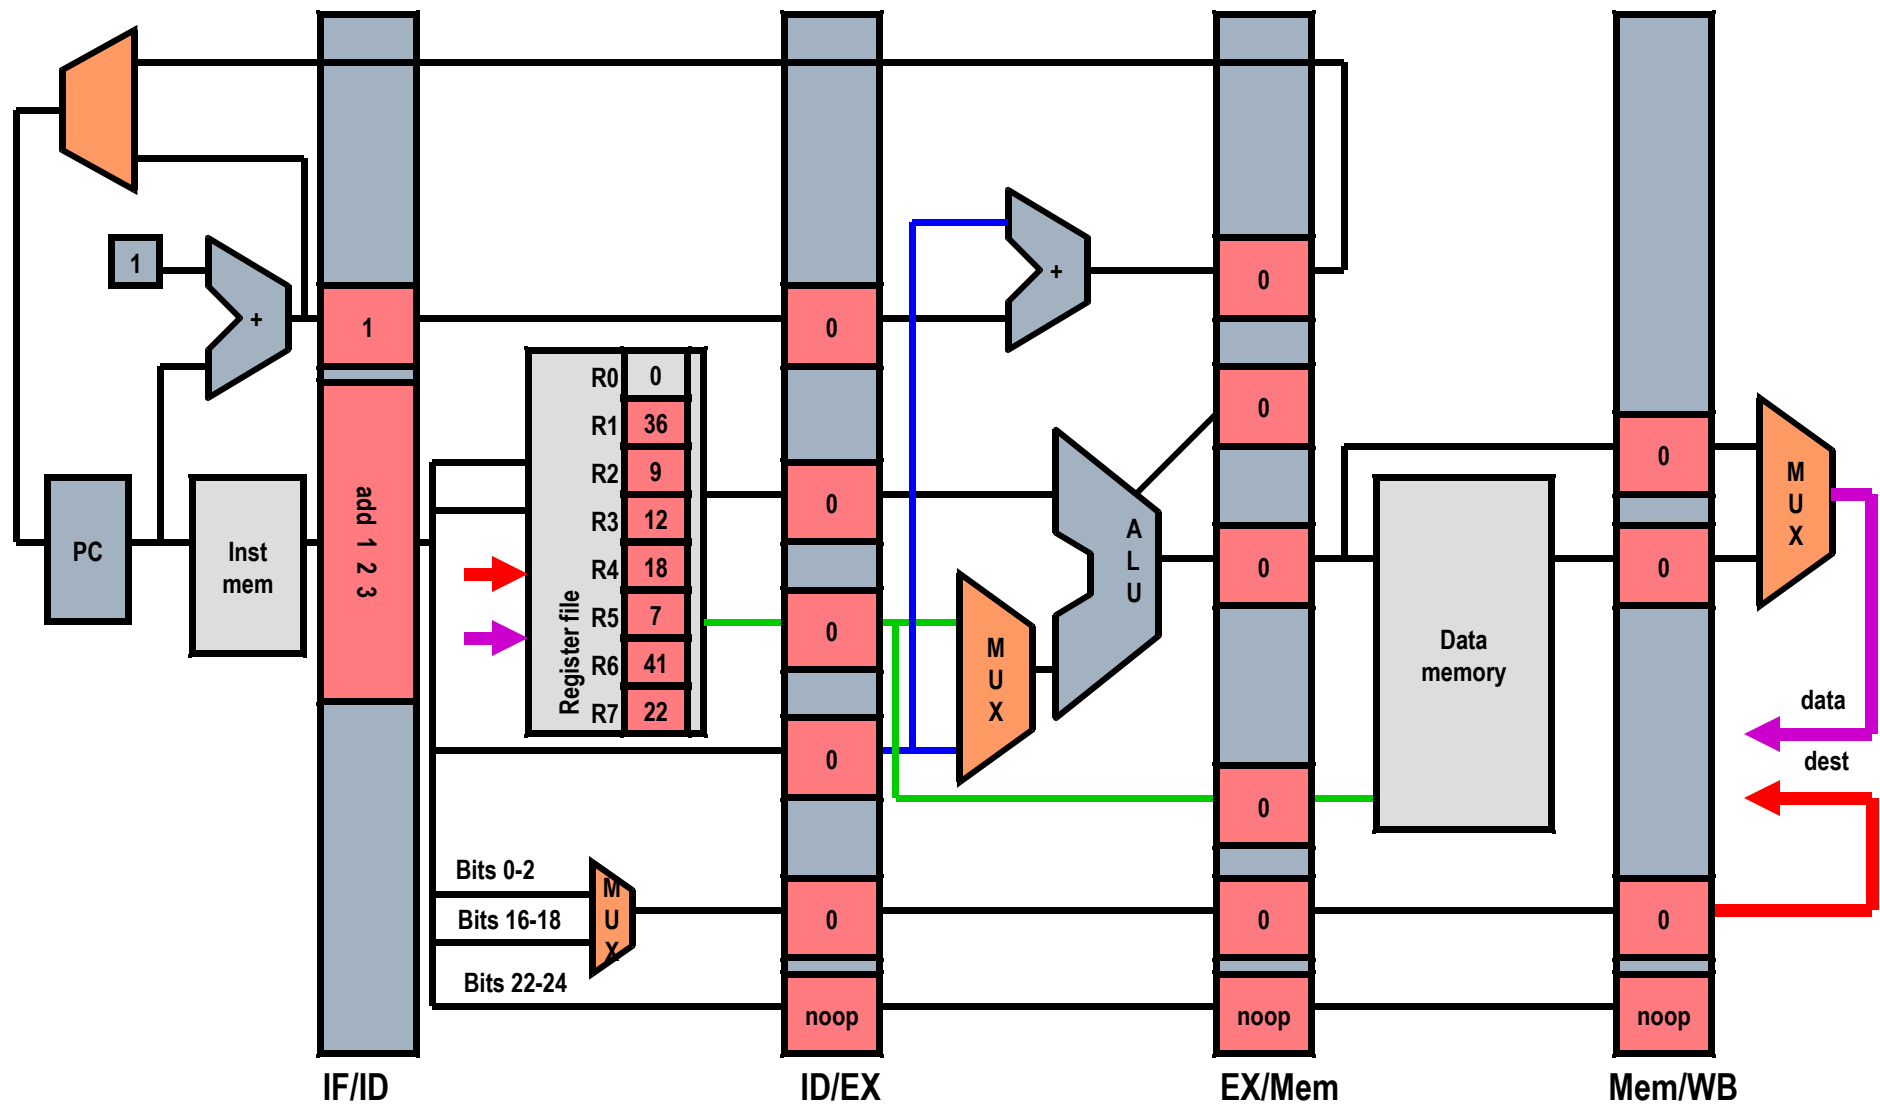
# Pipelined datapath



# Time 0 - initial state



# Time 1 - fetch: add 1 2 3

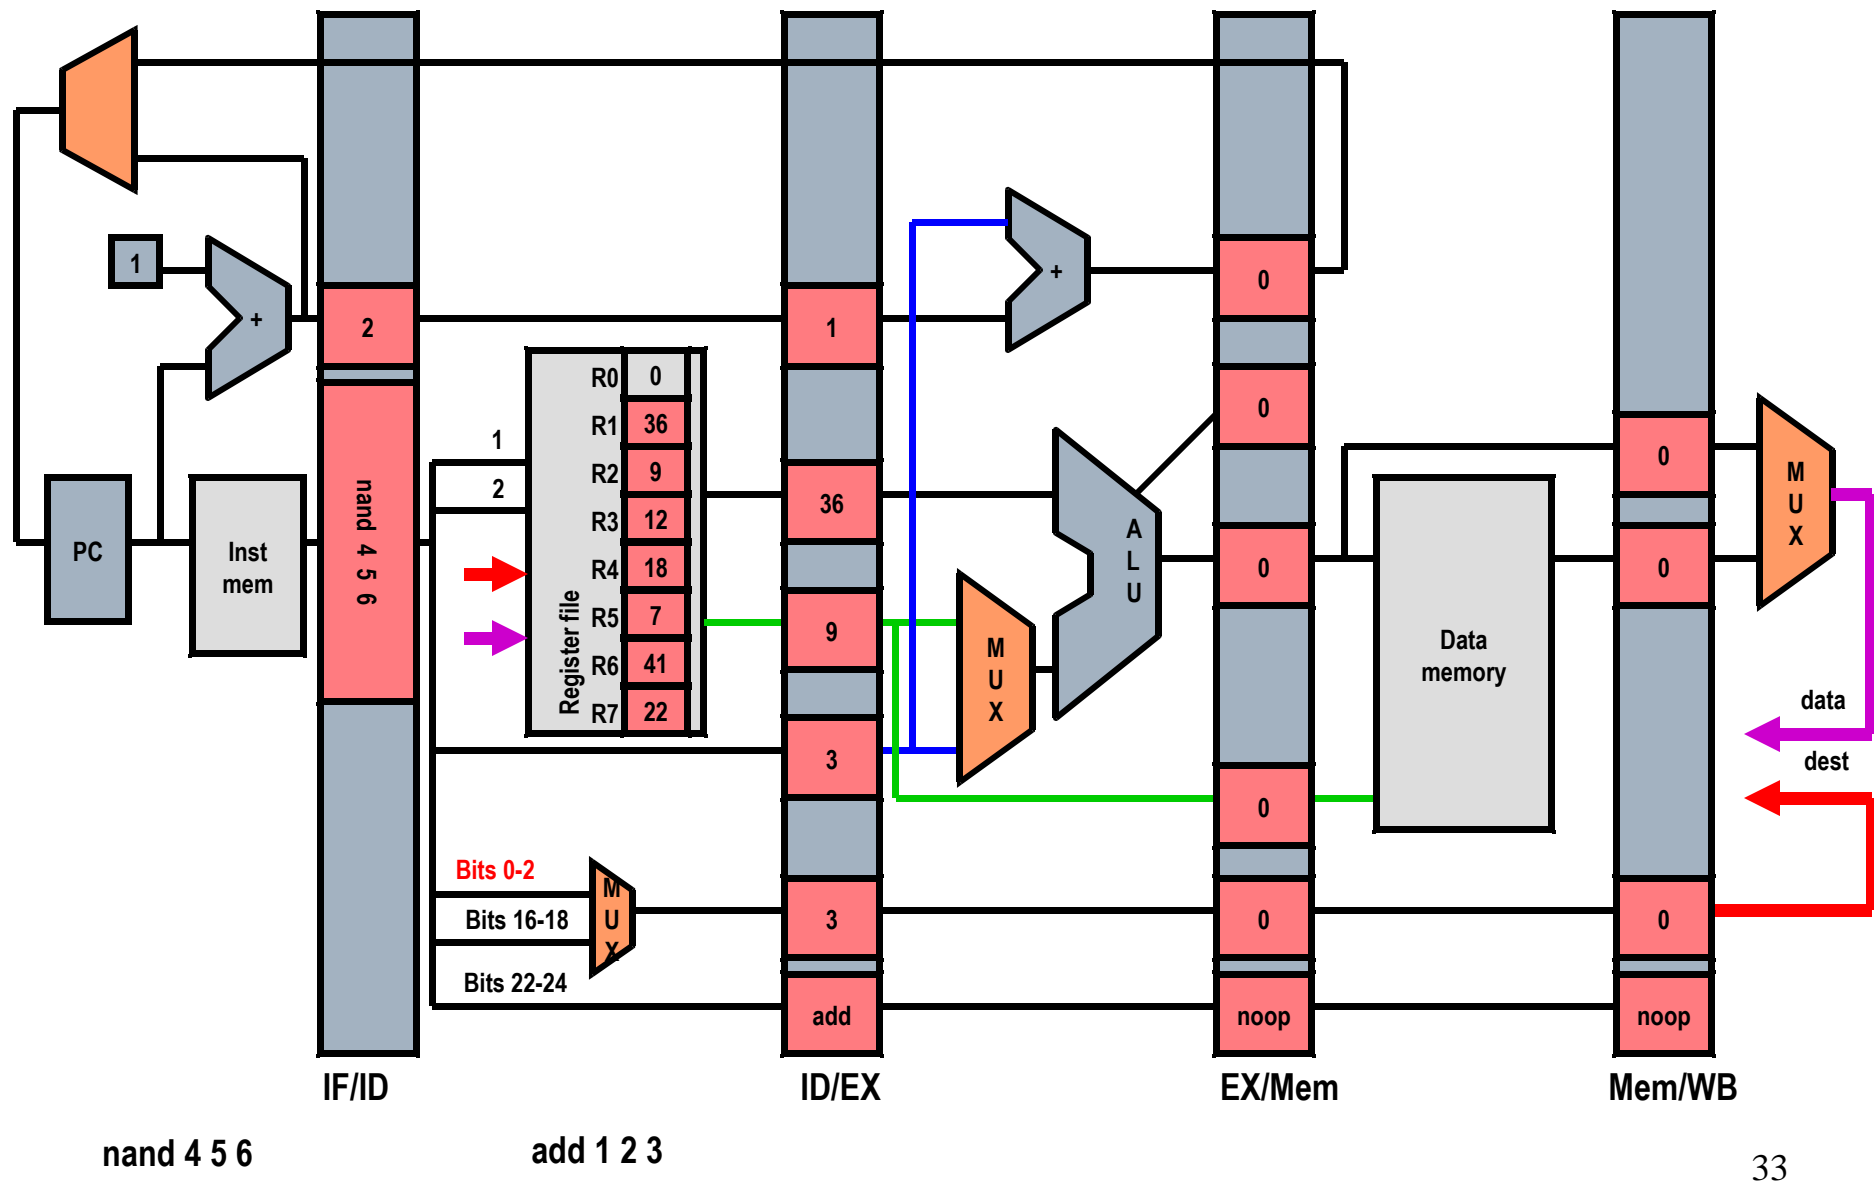


add 1 2 3

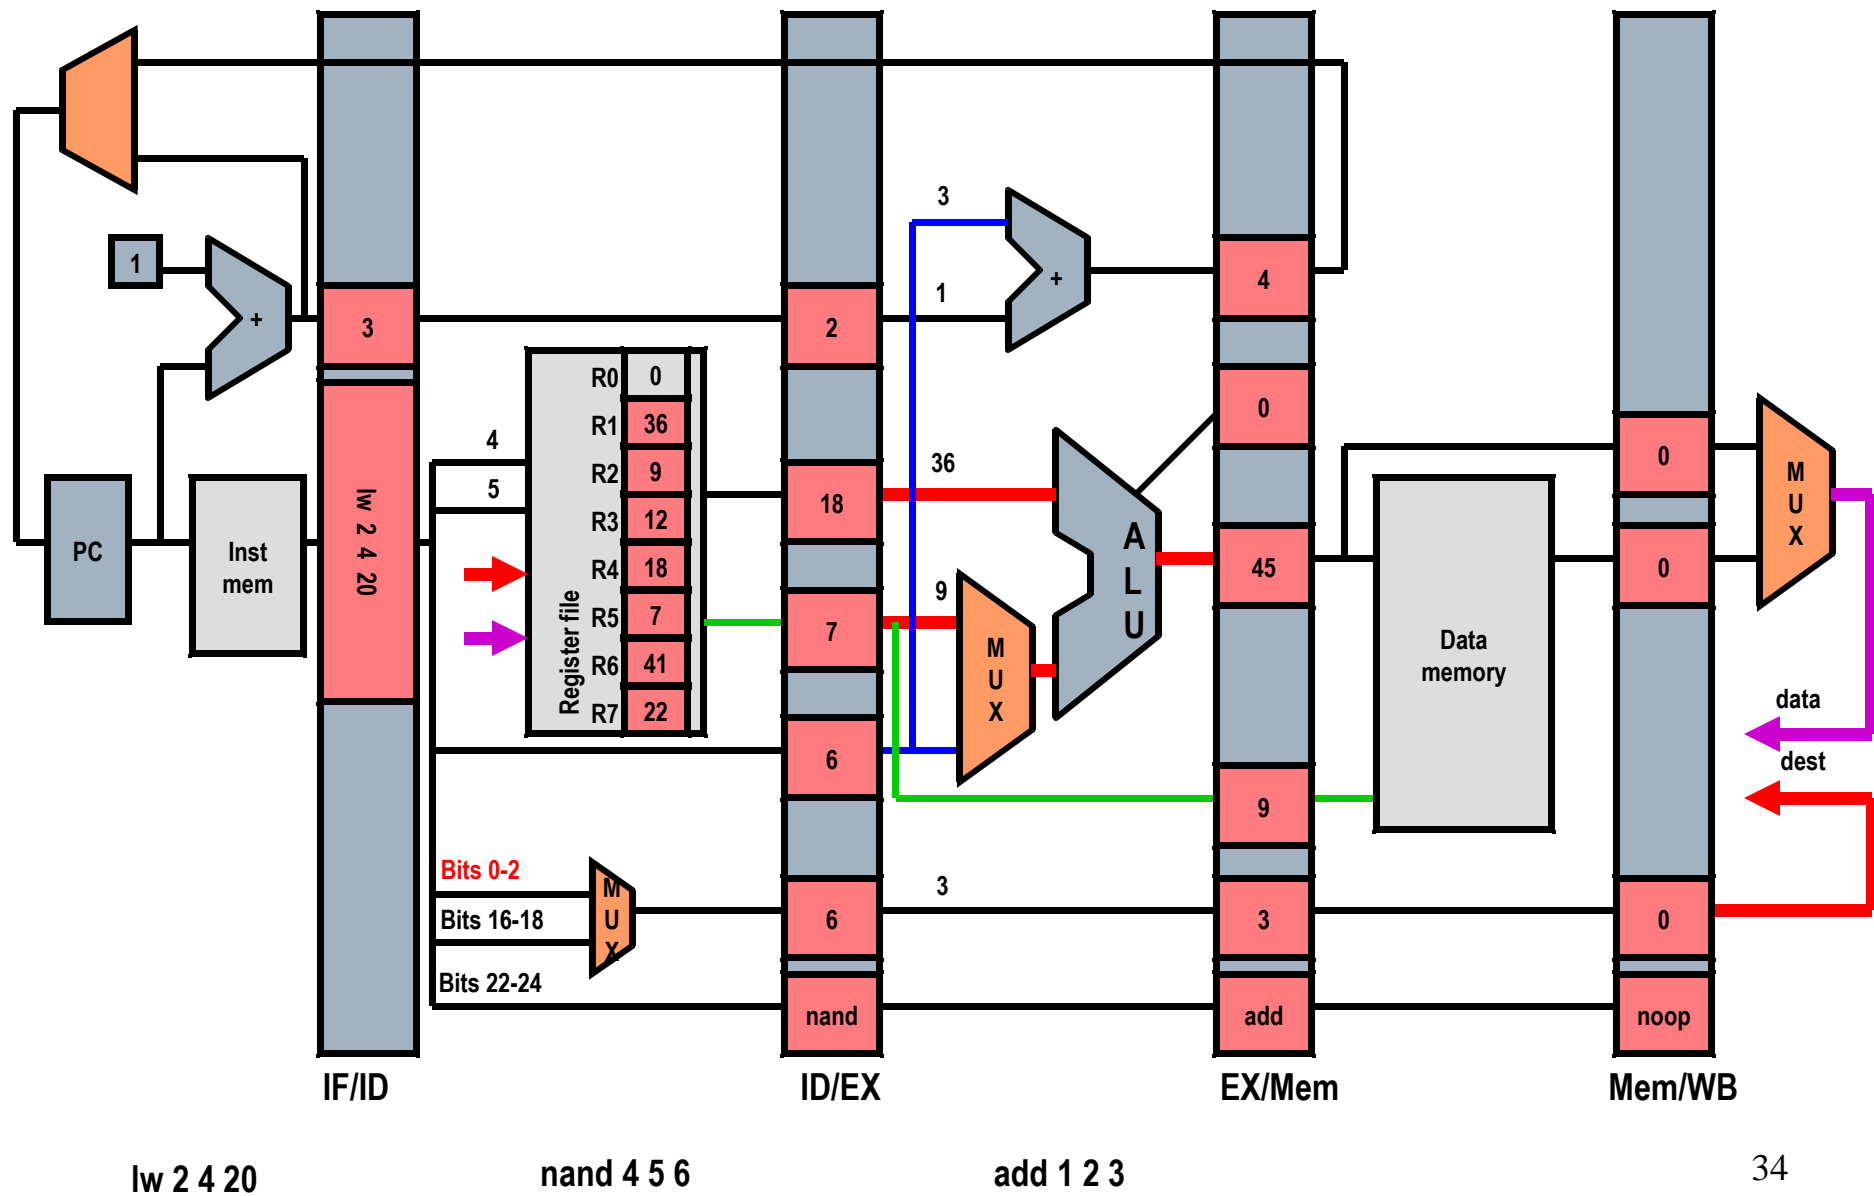
32



# Time 2 - fetch: nand 4 5 6

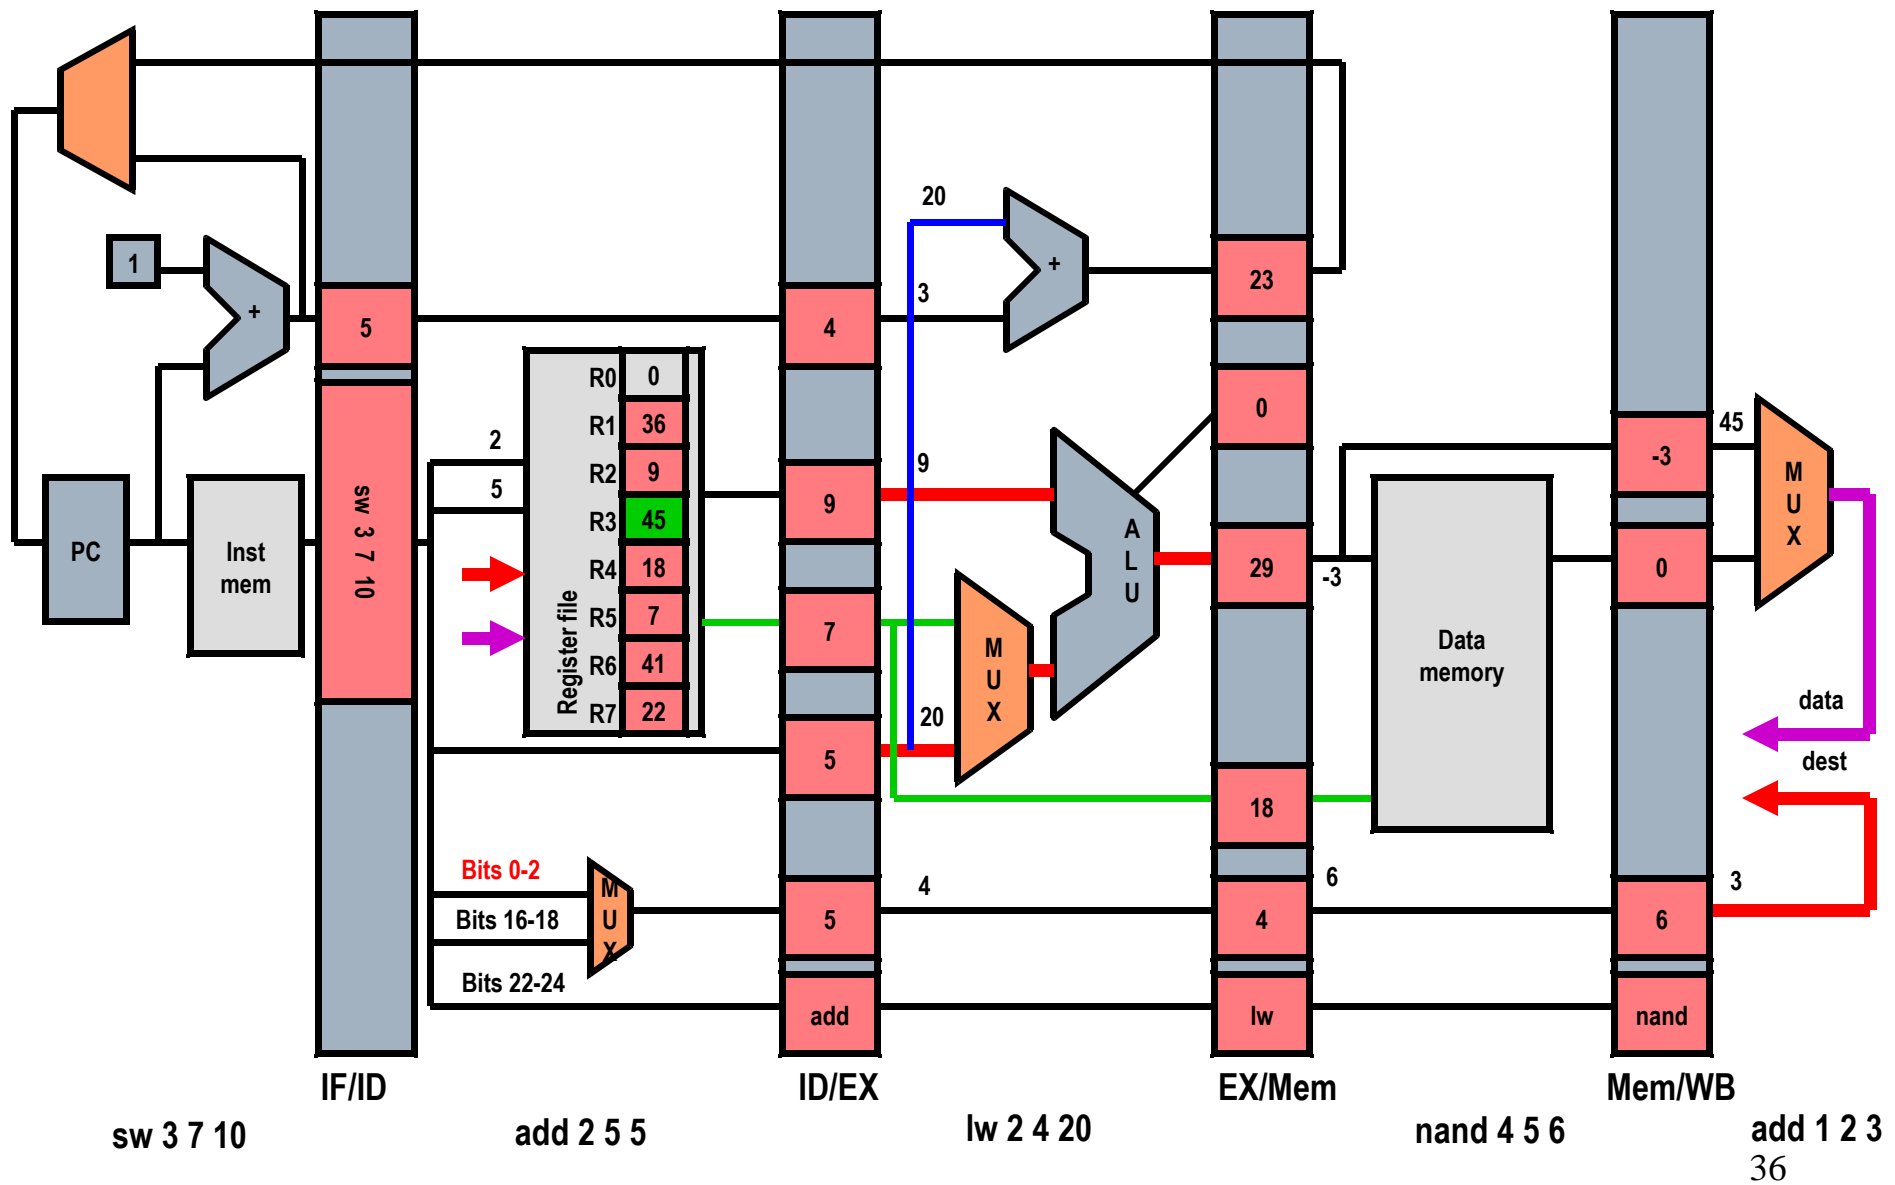


# Time 3 - fetch: lw 2 4 20

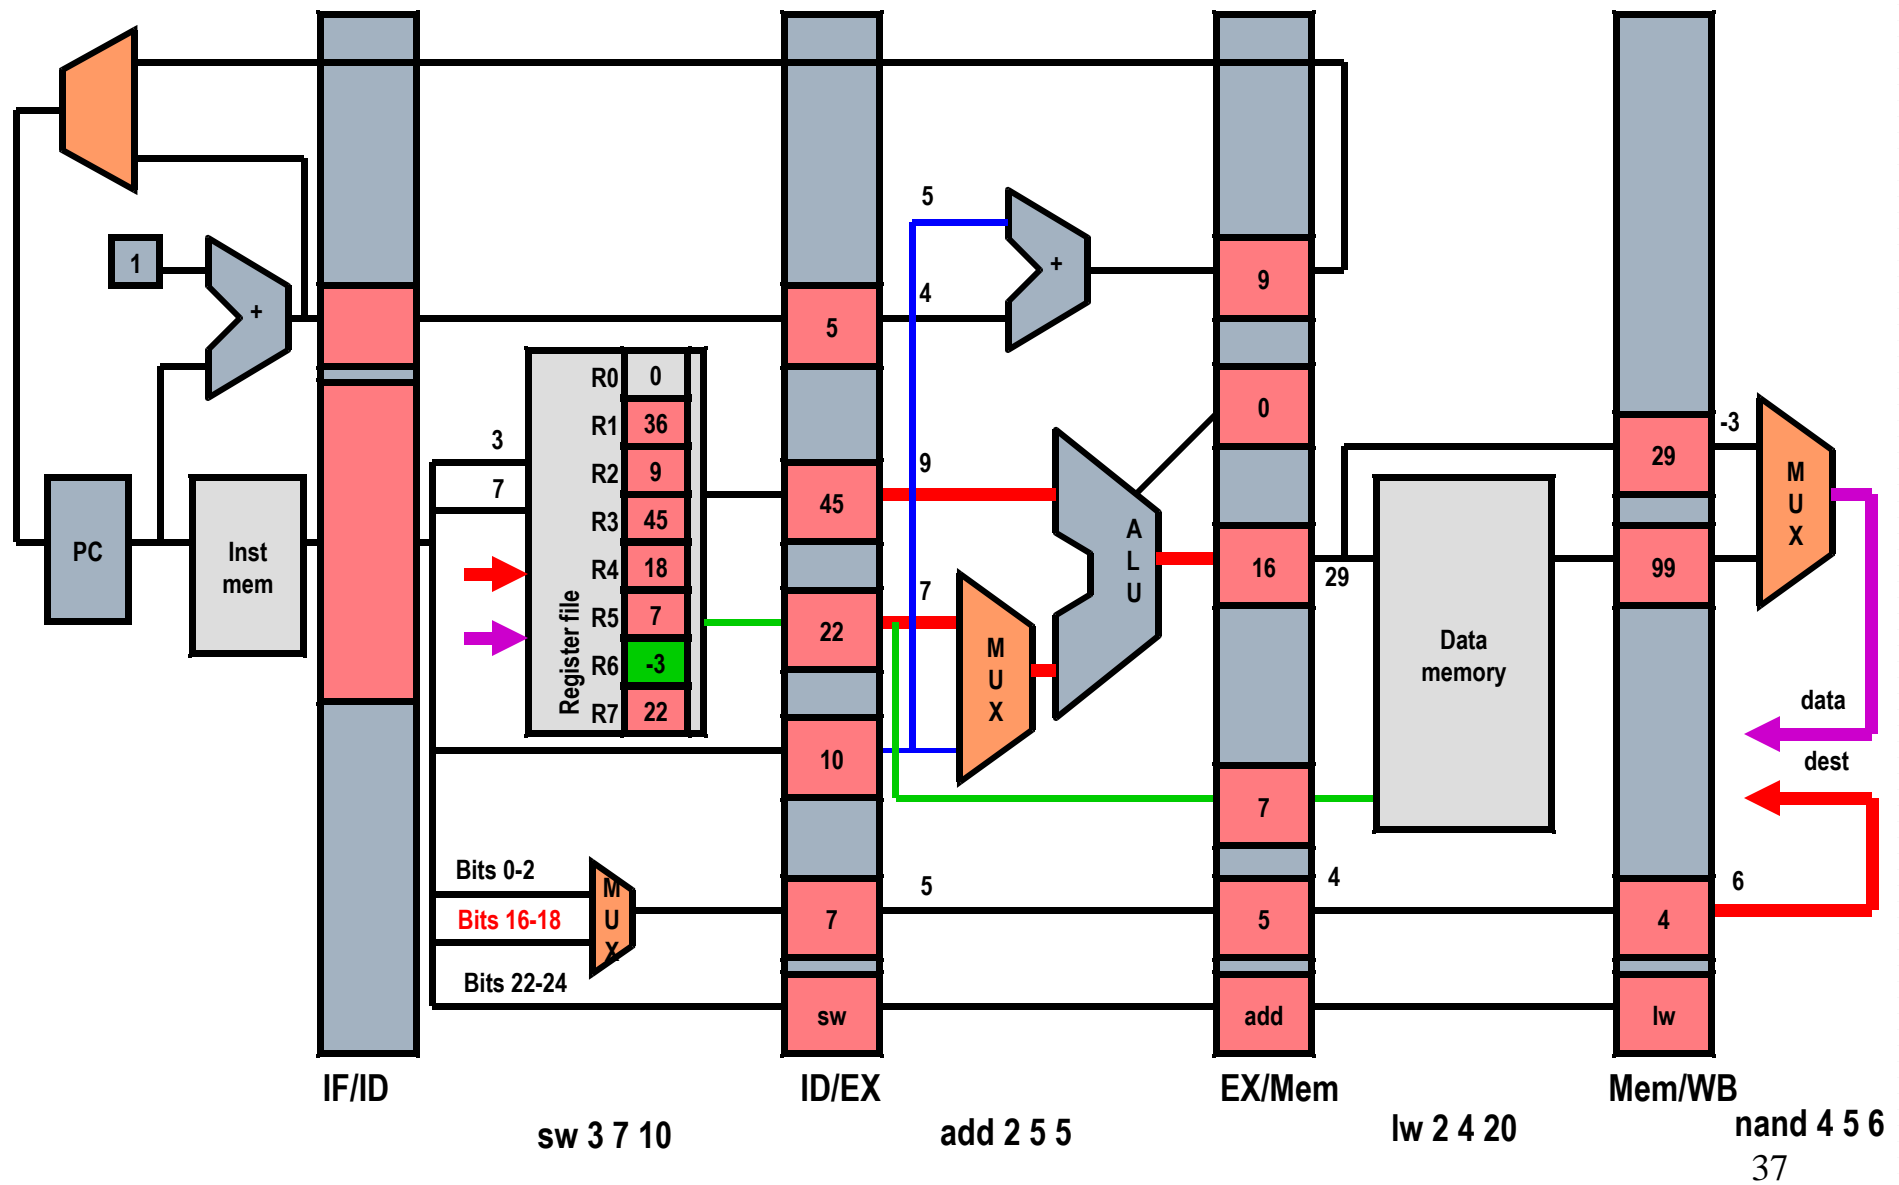




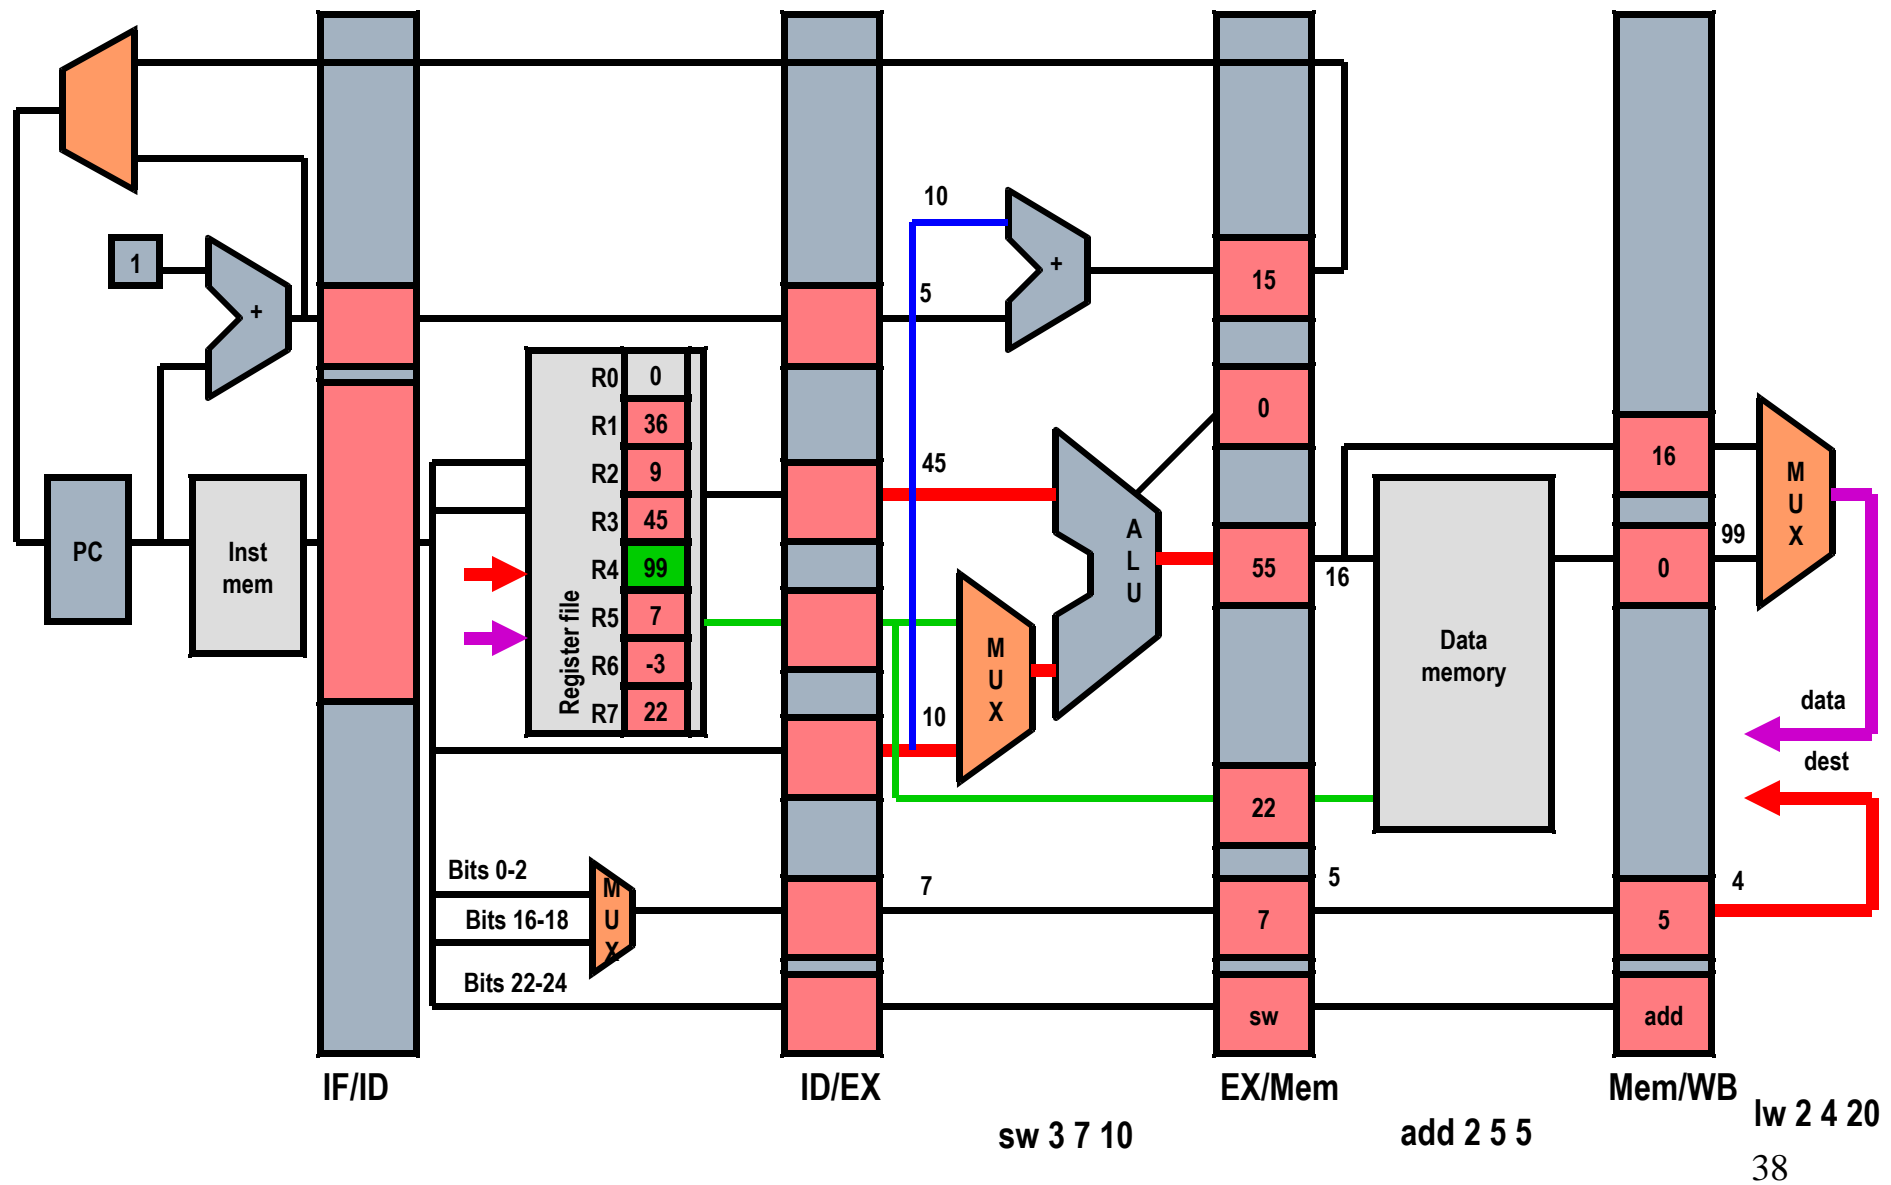
# Time 5 - fetch: sw 3 7 10



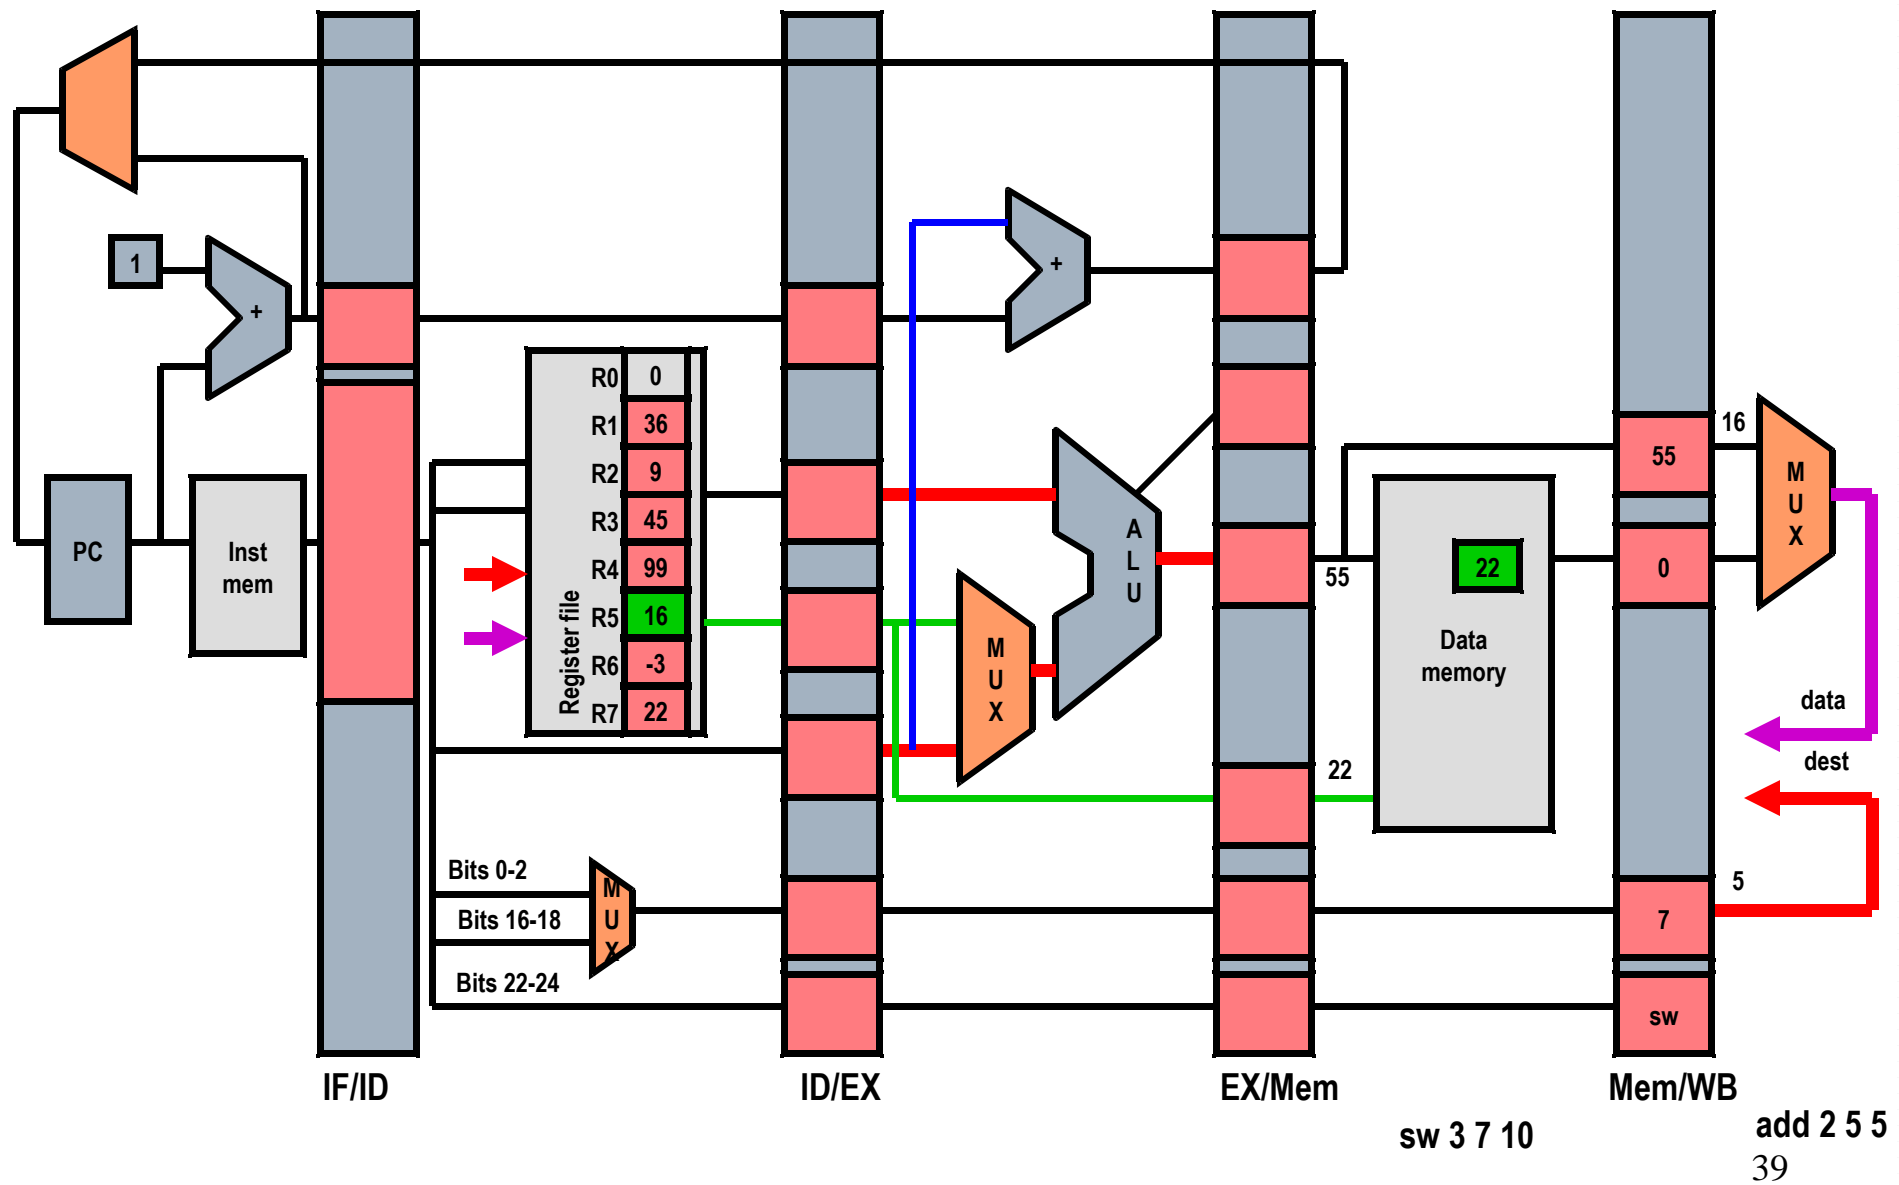
# Time 6 – no more instructions



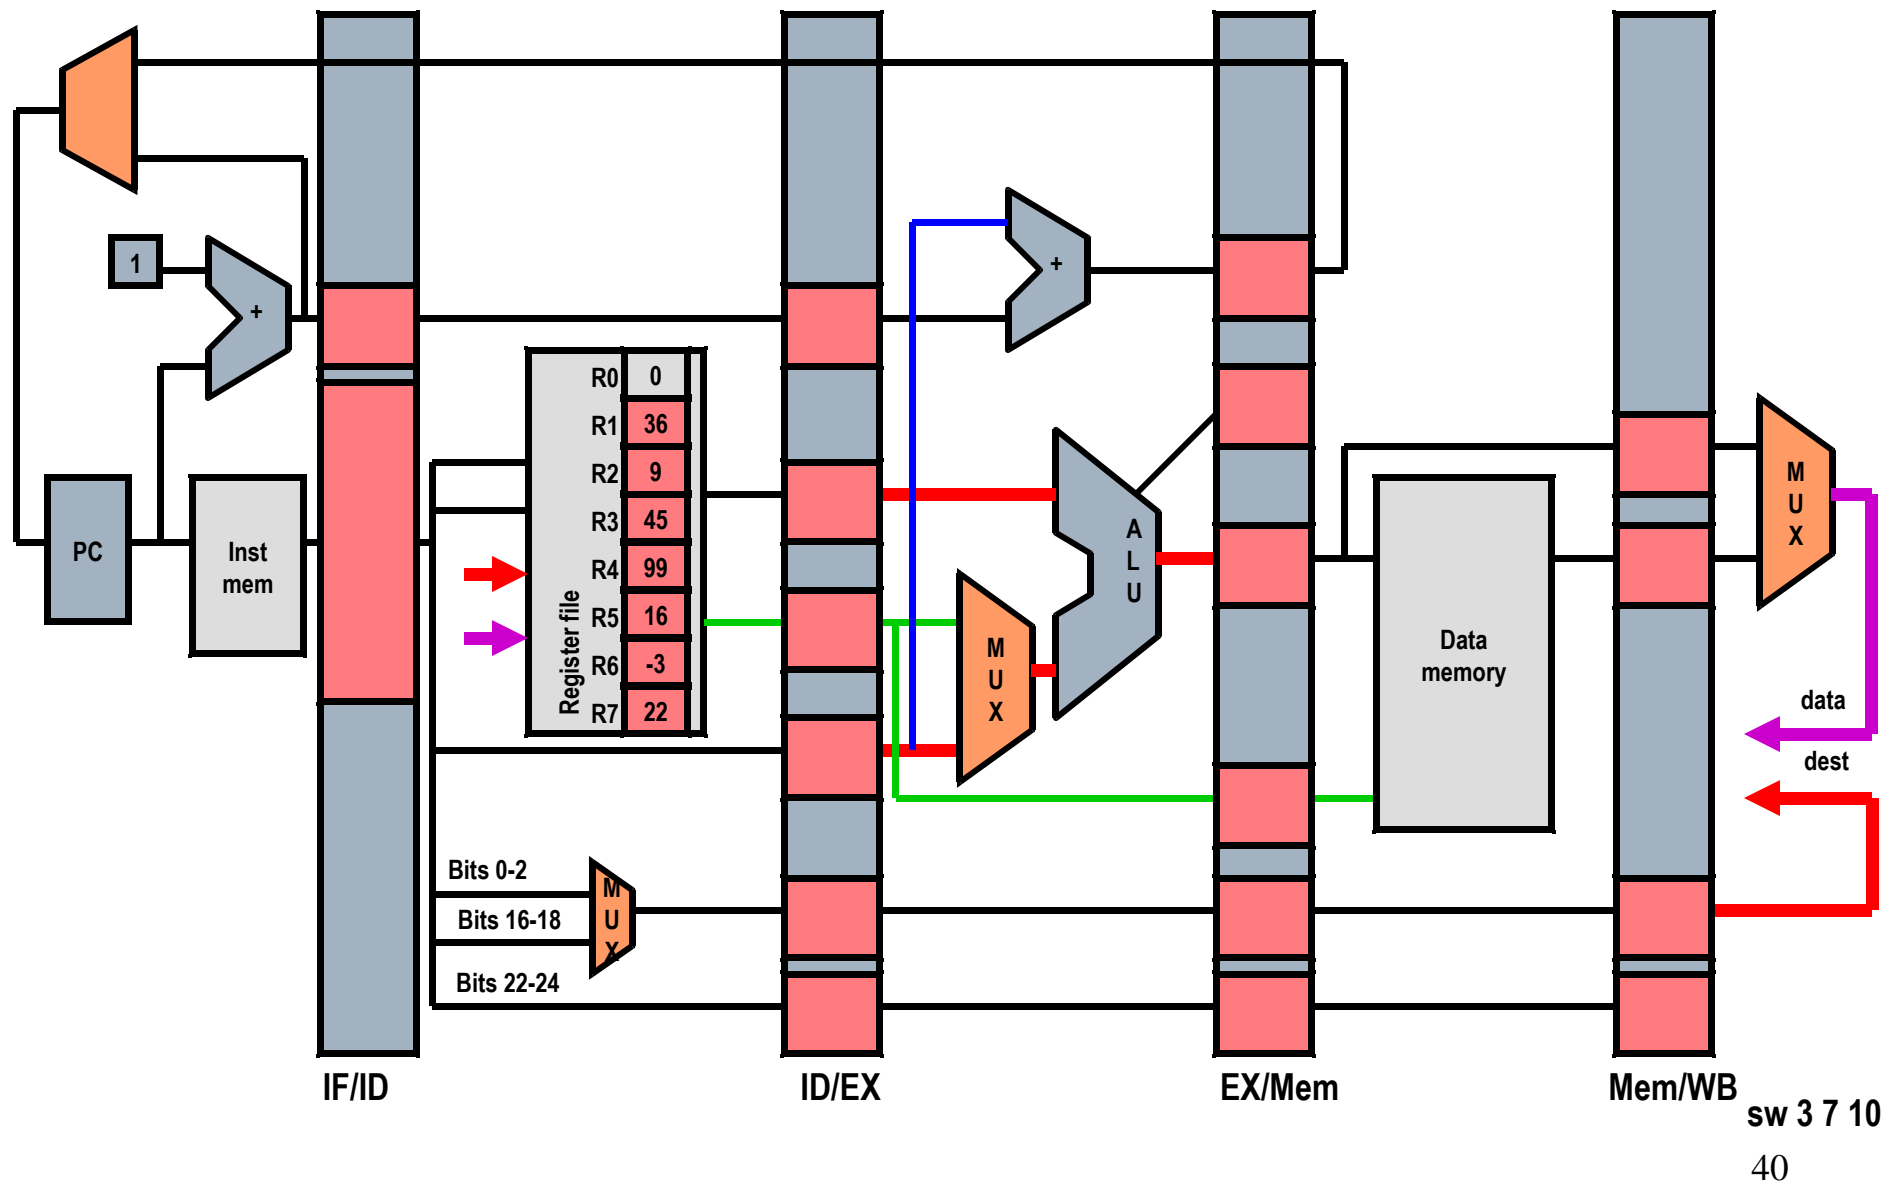
# Time 7 – no more instructions



# Time 8 – no more instructions




# Time 9 – no more instructions





## Time graphs (pipeline trace)



A vertical slice reports the entire activity of the pipeline at time 5.

# What can go wrong?

---

**Data hazards:** since register reads occur in stage 2 and register writes occur in stage 5, it is possible to read the wrong value if it is about to be written.

**Control hazards:** A branch instruction may change the PC, but not until stage 4.  
What do we fetch before that?

**Exceptions:** How do you handle exceptions in a pipelined processor with 5 instructions in flight?

Next Lecture: data hazards.

# Next time

---

- Hazards