

12. Pipeline Data Hazards

EECS 370 – Introduction to Computer Organization
Fall 2013

Prof. Valeria Bertacco, Robert Dick, and Satish Narayanasamy

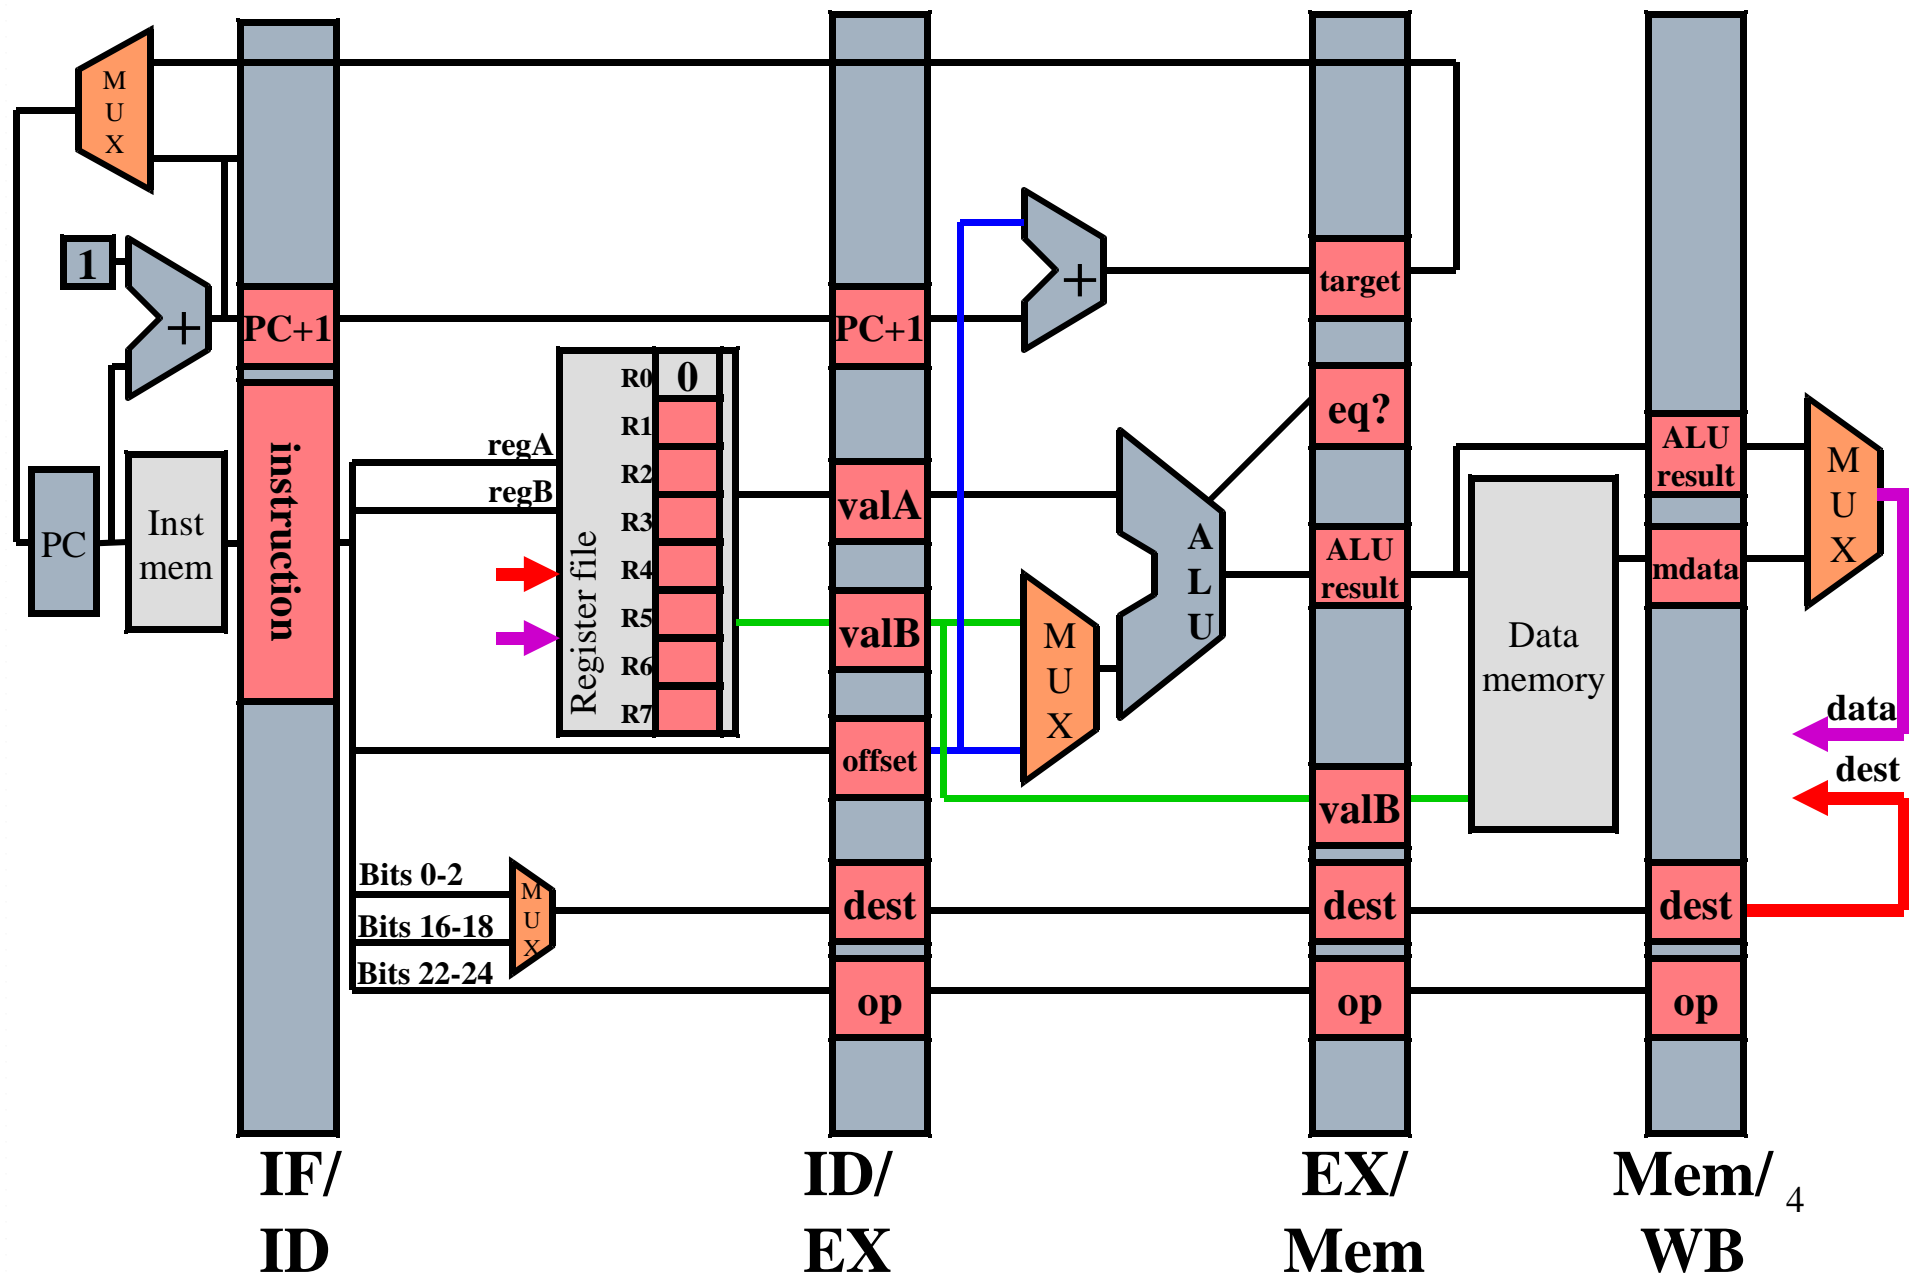
EECS Department
University of Michigan in Ann Arbor, USA

Announcements

- ❑ Project 2s autograder was buggy. Fixed now. Was passing invalid simulators. If your final submission was before 23 October, resubmit.
- ❑ See email for common problems that autograder previously missed.
- ❑ 24 October: Homework 4 due.
- ❑ 6:00 PM 28 October Project 2s and competition due.
- ❑ 25 October: Rest of Project 2 due.
- ❑ 5 November: Homework 5 due.
- ❑ 12 November: Project 3 due.

Pipelining - What can go wrong?


- ❑ **Data hazards**: since register reads occur in stage 2 and register writes occur in stage 5 it is possible to read the wrong value if is about to be written.
- ❑ **Control hazards**: A branch instruction may change the PC, but not until stage 4. What do we fetch before that?
- ❑ **Exceptions**: How do you handle exceptions in a pipelined processor with 5 instructions in flight?
- ❑ **Today - Data hazards**
 - What are they?
 - How do you detect them?
 - How do you deal with them?



Pipeline function for ADD

- ❑ Fetch: read instruction from memory
- ❑ Decode: read source operands from reg
- ❑ Execute: calculate sum
- ❑ Memory: pass results to next stage
- ❑ Writeback: write sum into register file


Data Hazards



If not careful, nand will read the wrong value of R3

Data Hazards

add 1 2 3
nand 3 4 5



Assume Register File gives the right value of R3 when read/written during same cycle. This is consistent with the book and the MIPS architecture, but not Project 3.

Class Problem 1

Which data hazards do you see?

```
add 1 2 3
nand 3 4 5
add 6 3 7
lw 3 6 10
sw 6 2 12
```

What about here?

```
add 1 2 3
beq 3 4 1
add 3 5 6
add 3 6 7
```


Three approaches to handling data hazards

- ❑ Avoid
 - Make sure there are no hazards in the code
- ❑ Detect and stall
 - If hazards exist, stall the processor until they go away.
- ❑ Detect and forward
 - If hazards exist, augment the pipeline to get the values where they are needed with reduced stalls

Handling data hazards I: Avoid all hazards

- ❑ Assume the programmer (or the compiler) knows about the processor implementation.
 - Make sure no hazards exist.
 - Put noops between any dependent instructions.

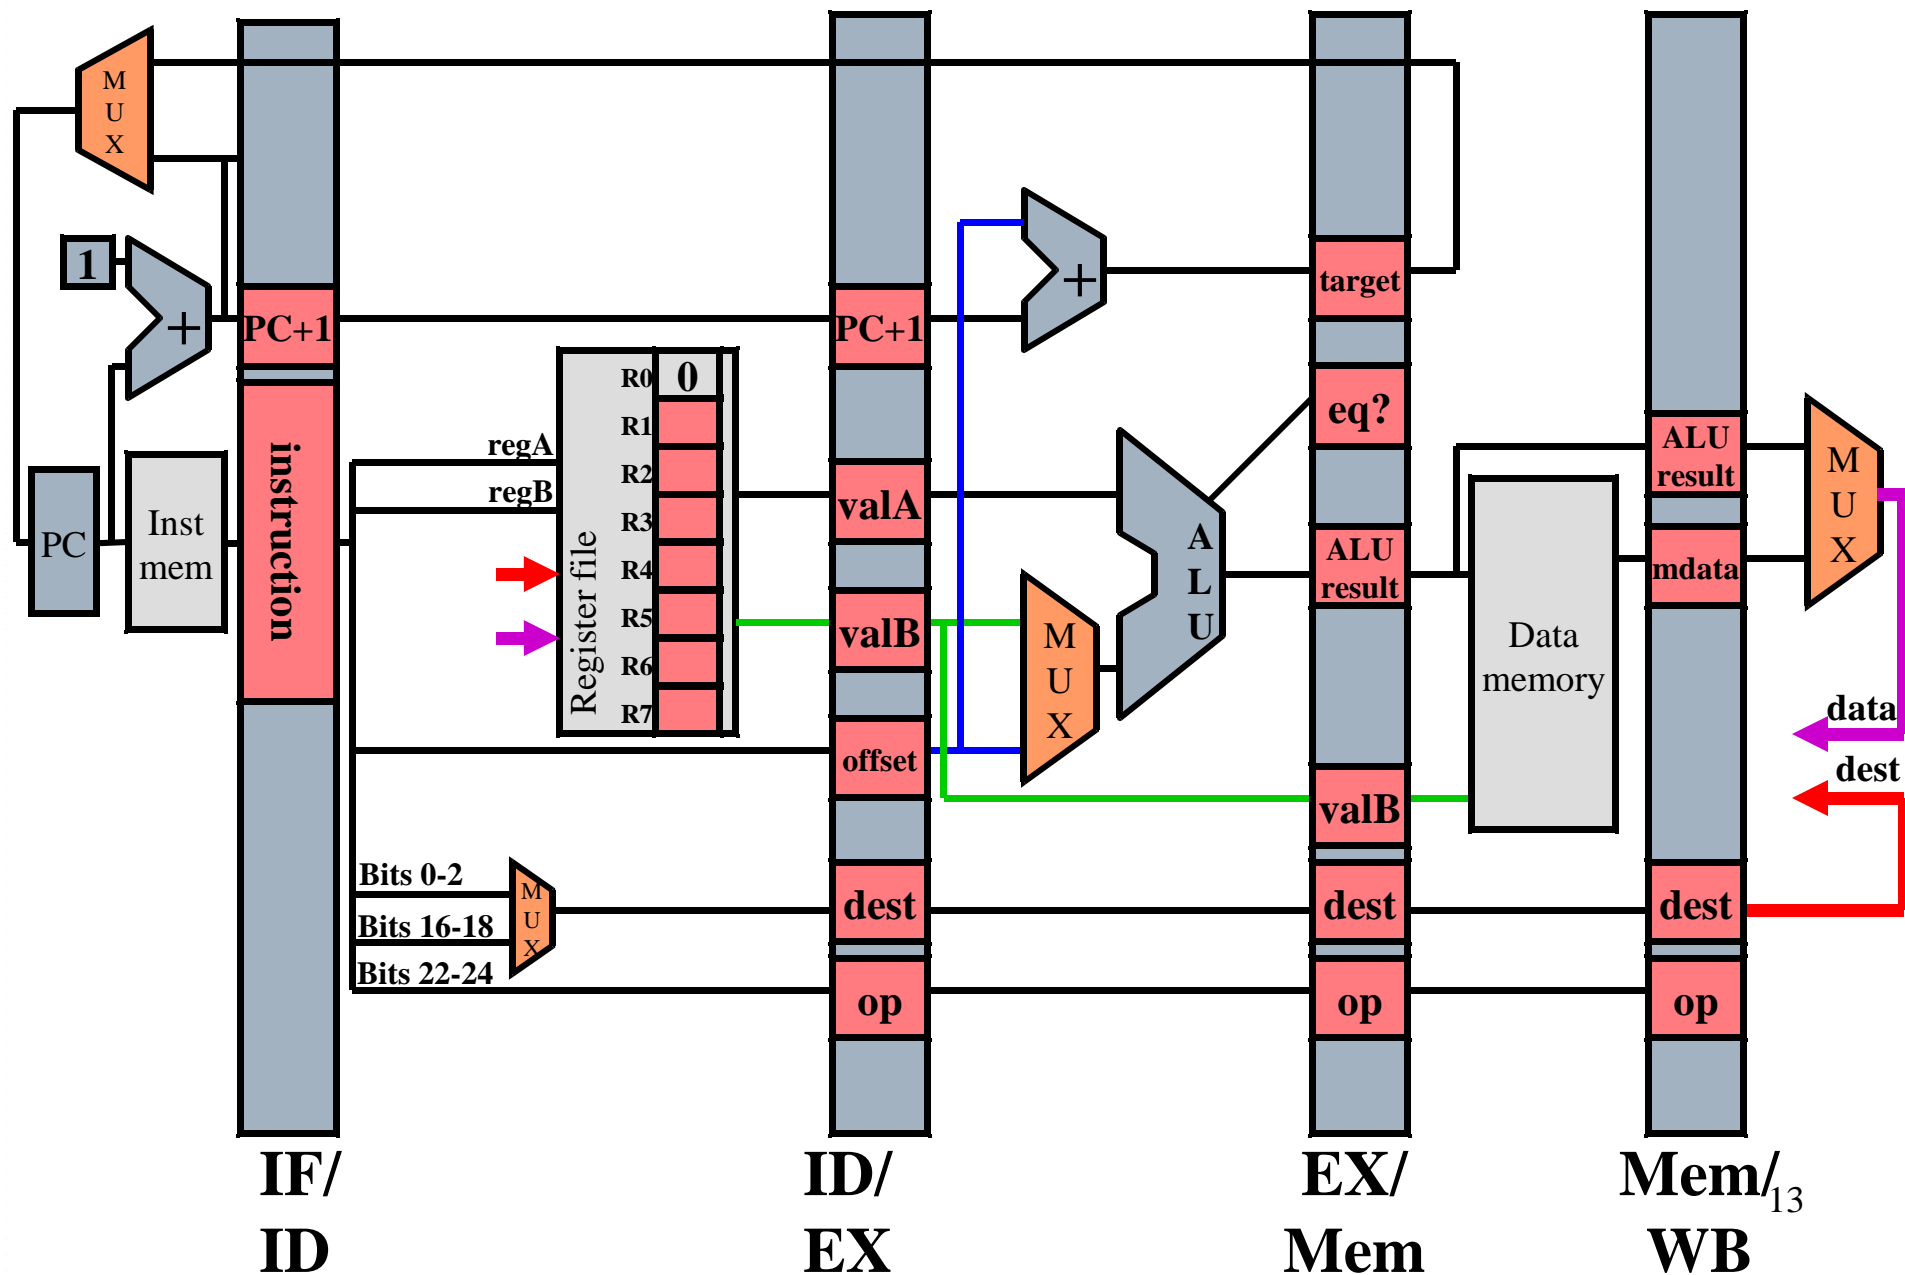
add 1 2 3 ← write R3 in cycle 5
noop
noop
nand 3 4 5 ← read R3 in cycle 5

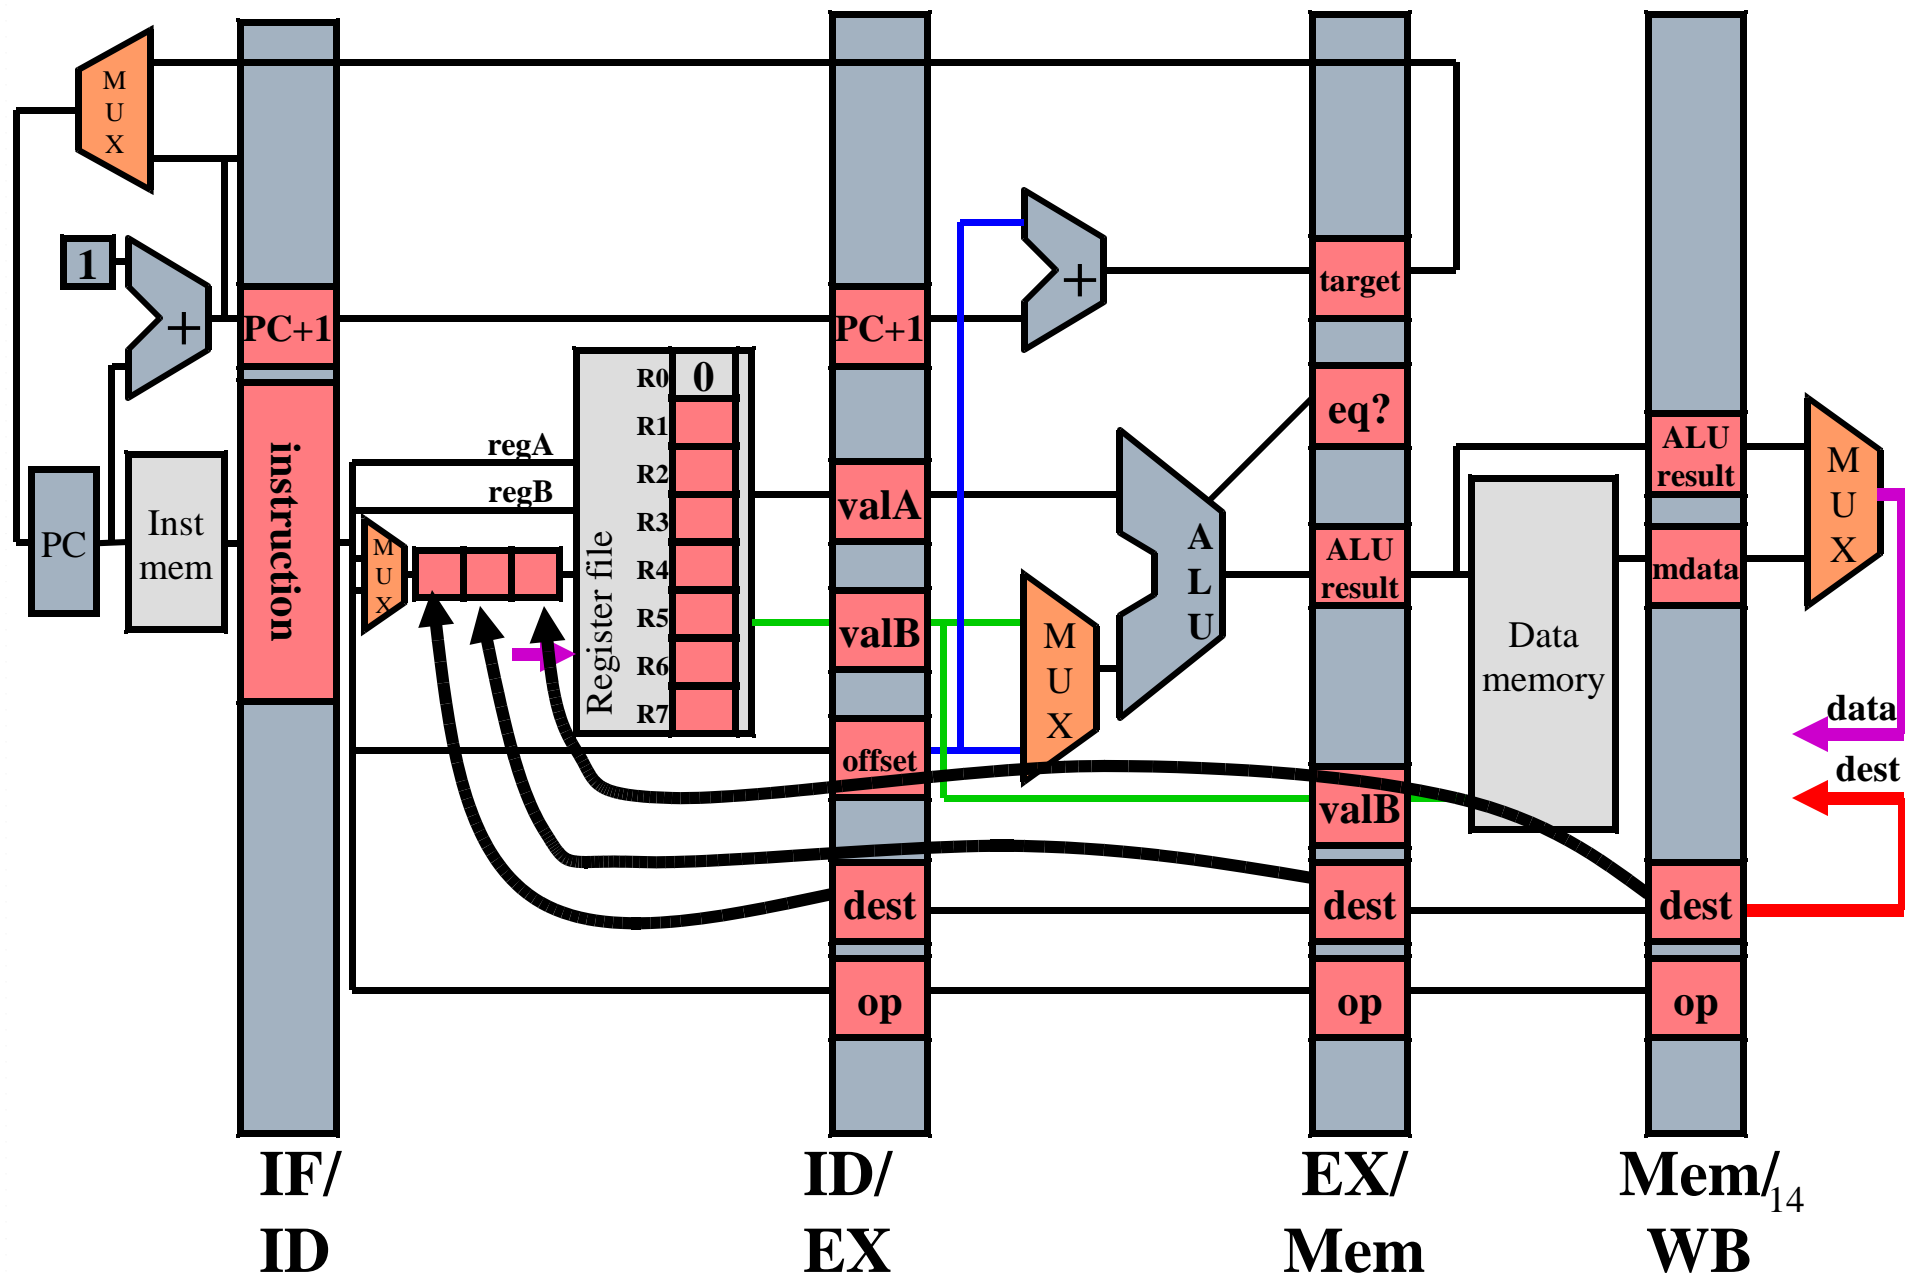
Problems with this solution

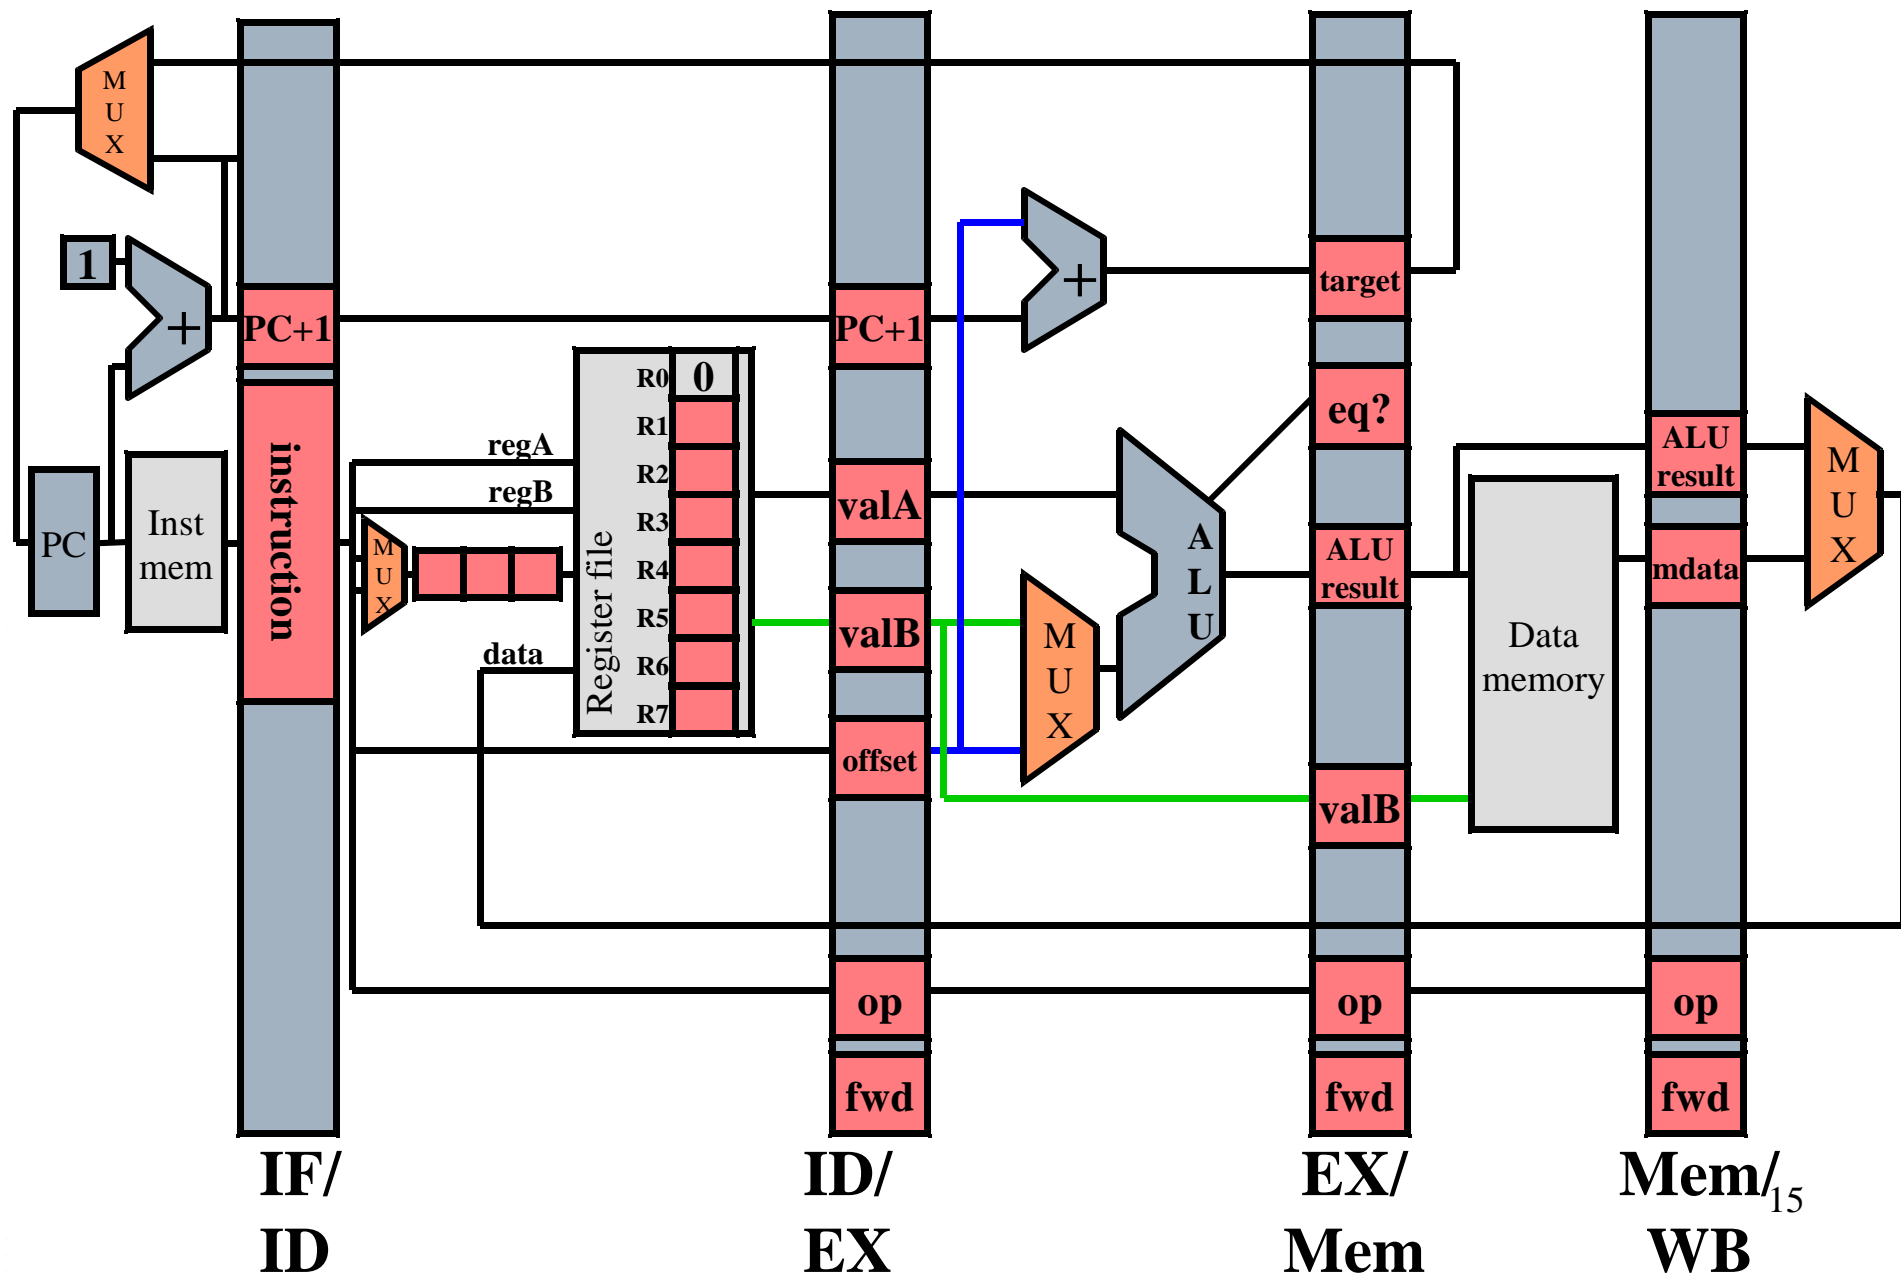
- ❑ Old programs (legacy code) may not run correctly on new implementations
 - Longer pipelines need more noops
- ❑ Programs get larger as noops are included
 - Especially a problem for machines that try to execute more than one instruction every cycle
 - Intel EPIC: Often 25% - 40% of instructions are noops
- ❑ Program execution is slower
 - CPI is 1, but some instructions are noops


Handling data hazards II: Detect and stall until ready


- ❑ Detect:
 - Compare regA with previous DestRegs
 - 3 bit operand fields
 - Compare regB with previous DestRegs
 - 3 bit operand fields
- ❑ Stall:
 - Keep current instructions in fetch and decode
 - Pass a noop to execute











Handling data hazards II: Detect and stall until ready

❑ Detect:

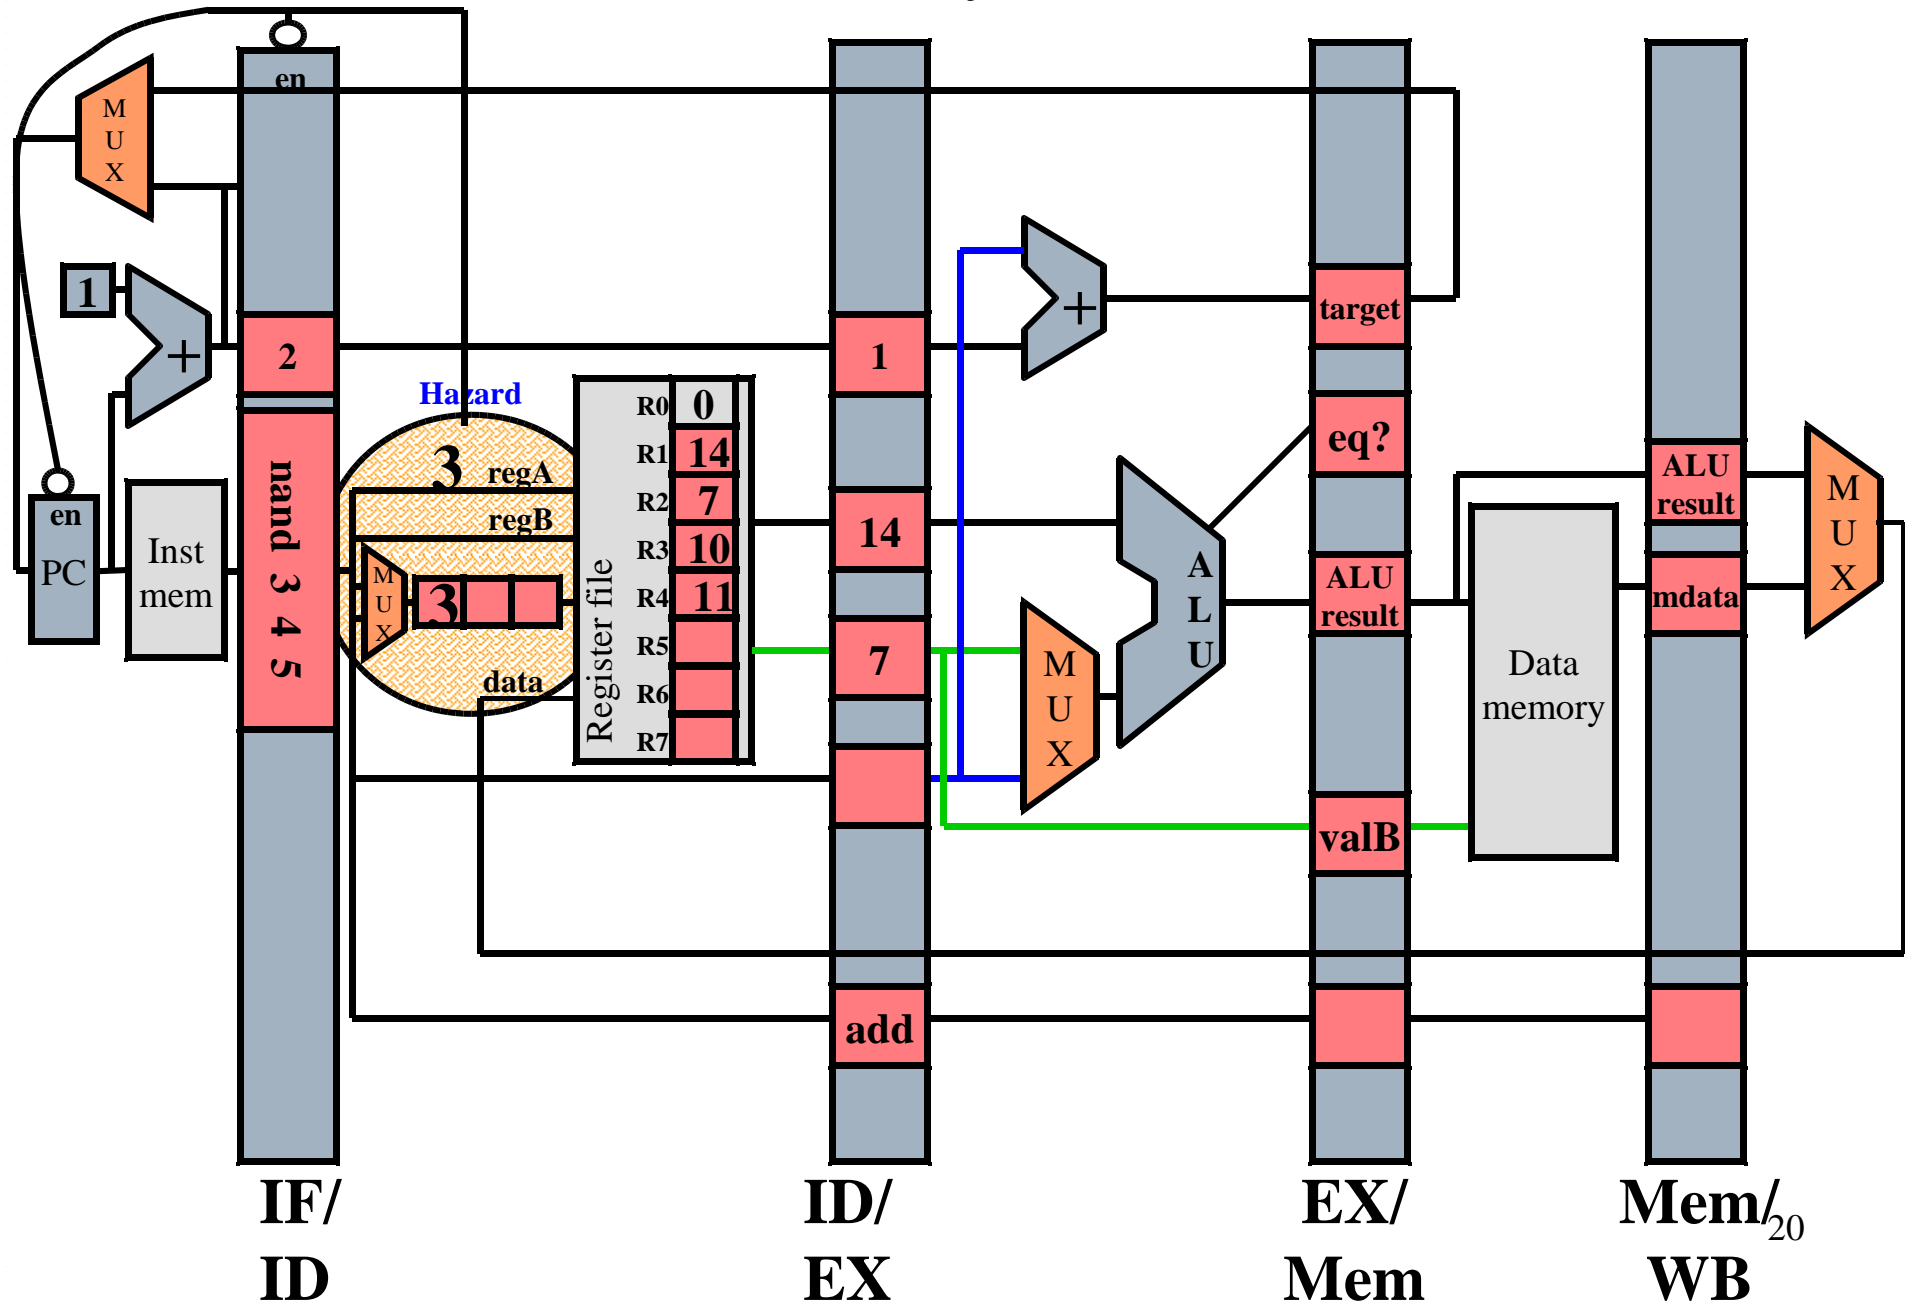
- Compare regA with previous DestReg
 - 3 bit operand fields
- Compare regB with previous DestReg
 - 3 bit operand fields

❑ Stall:

Keep current instructions in fetch and decode

Pass a noop to execute

First half of cycle 3



Handling data hazards II:

Detect and stall until ready

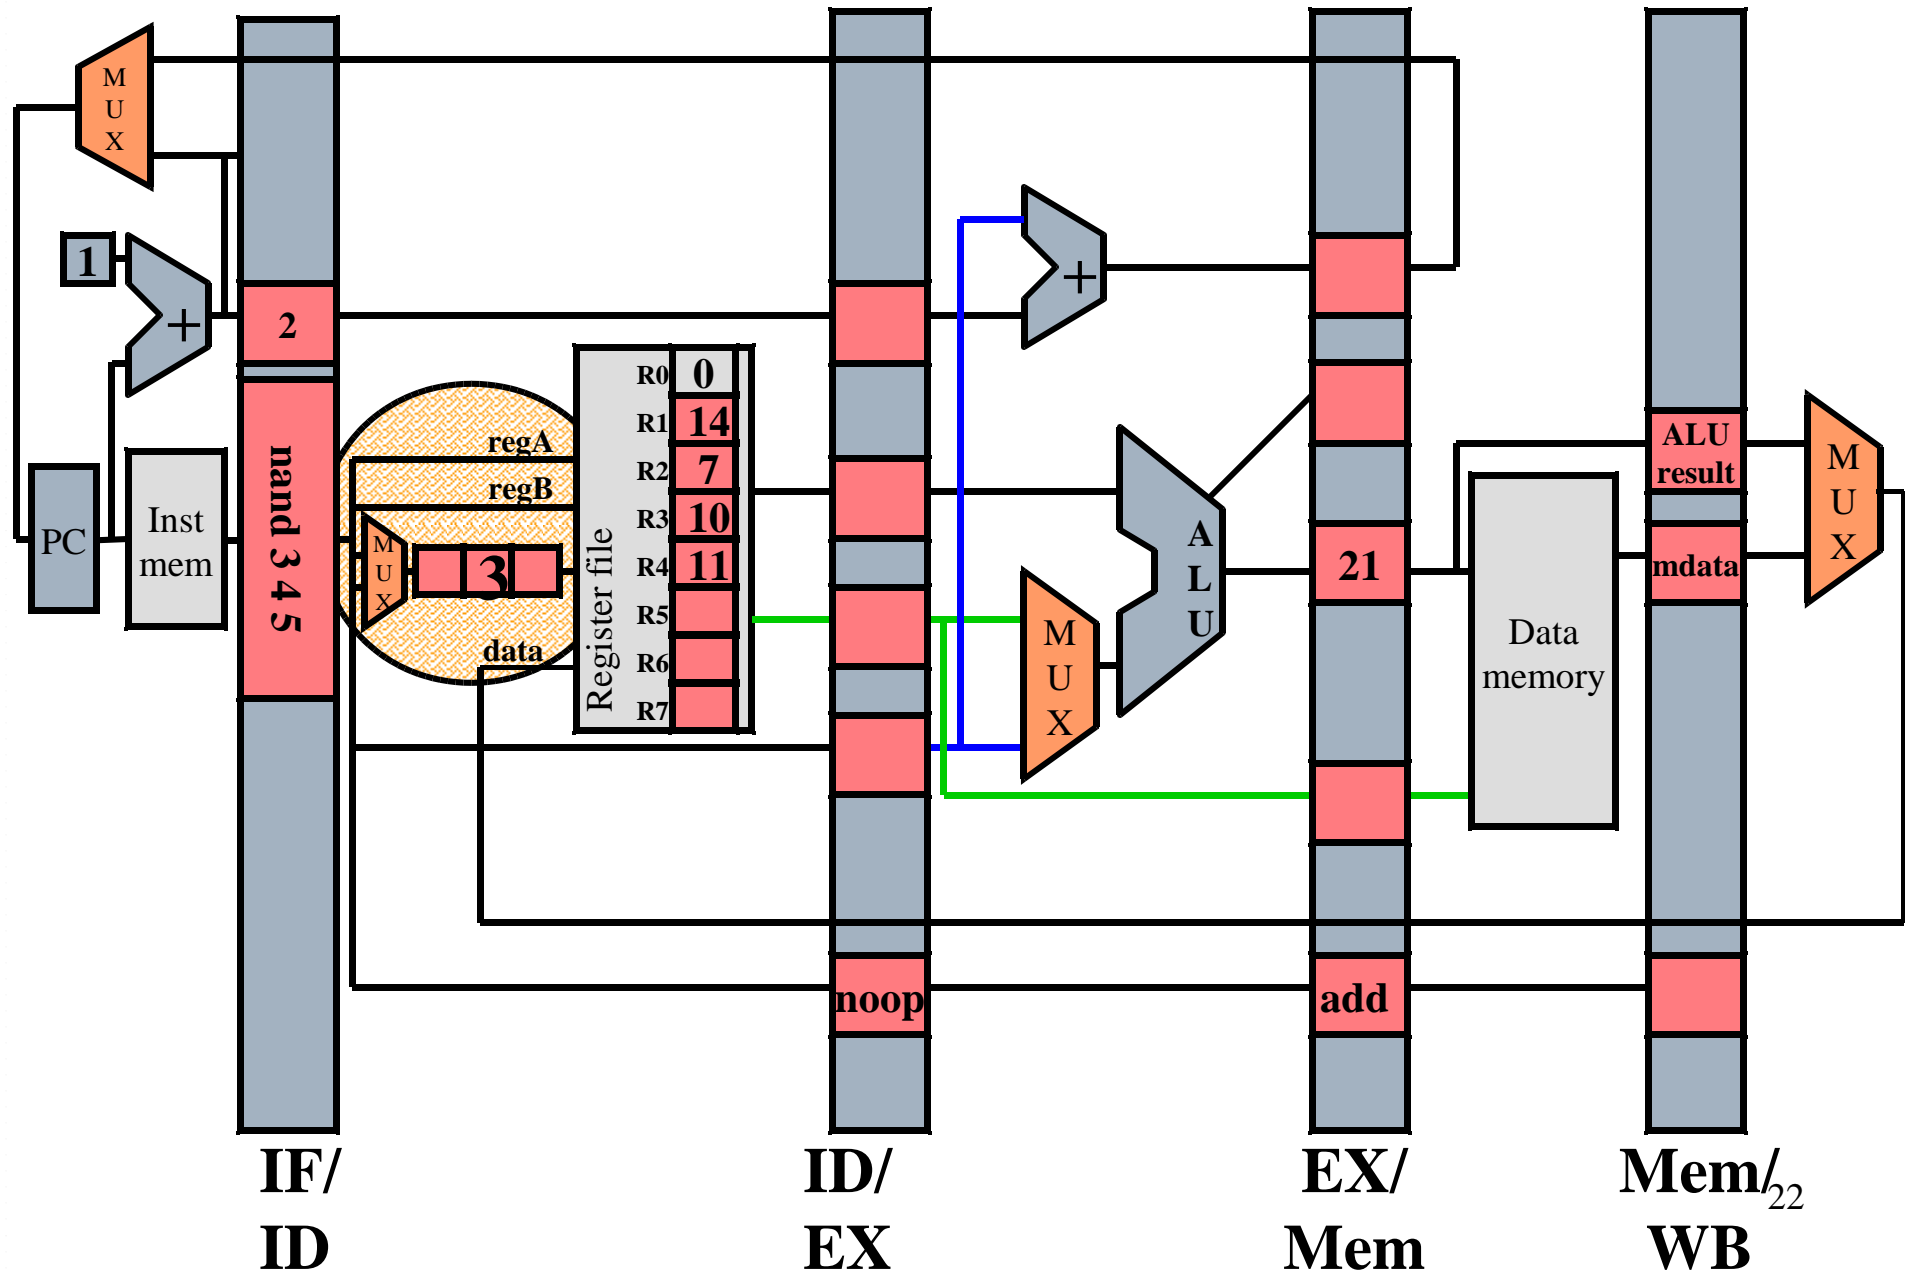
❑ Detect:

- Compare regA with previous DestReg
 - 3 bit operand fields
- Compare regB with previous DestReg
 - 3 bit operand fields

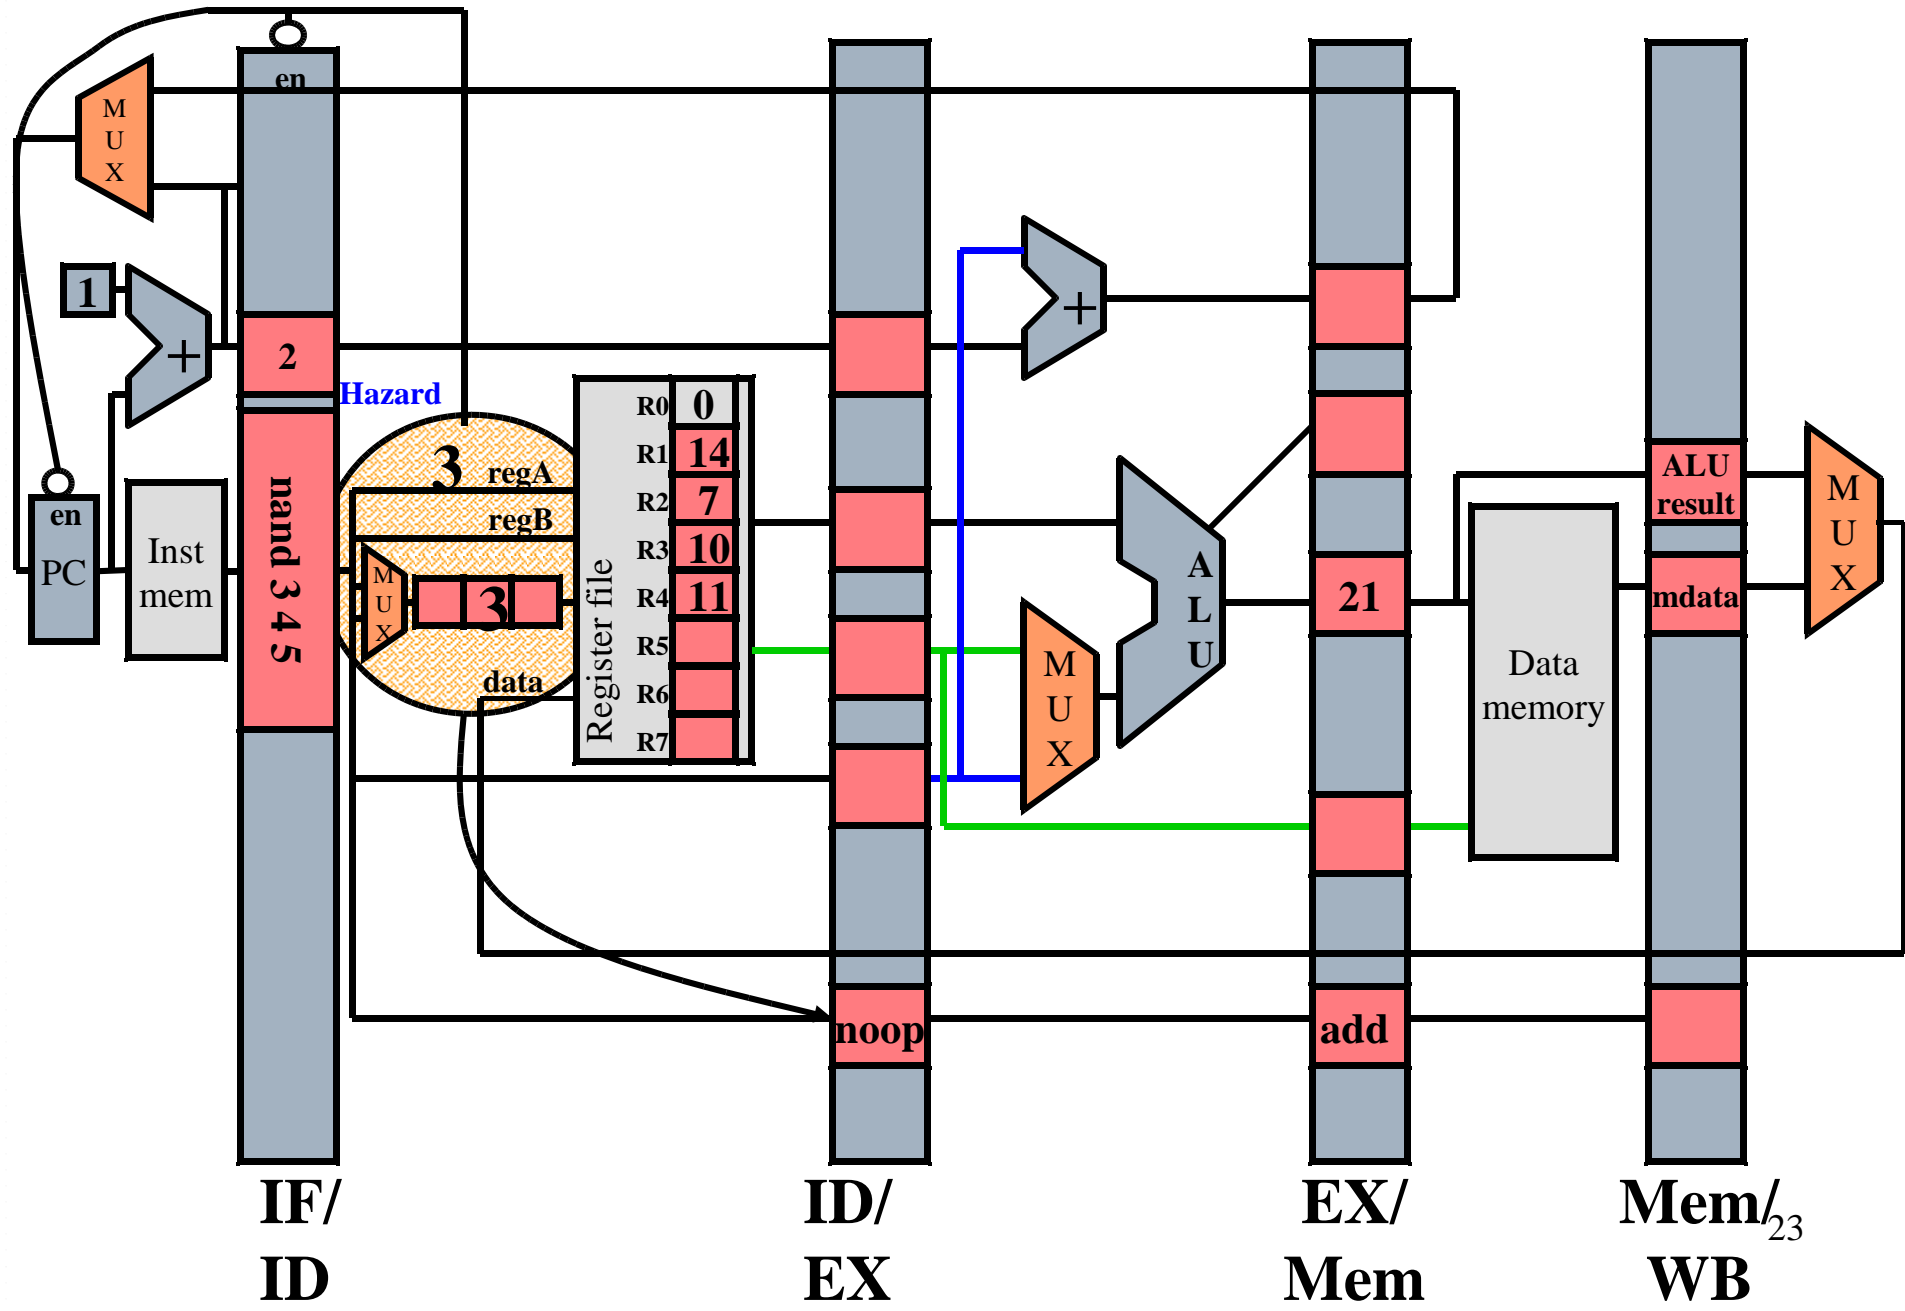
❑ Stall:

- Keep current instructions in fetch and decode
- Pass a noop to execute

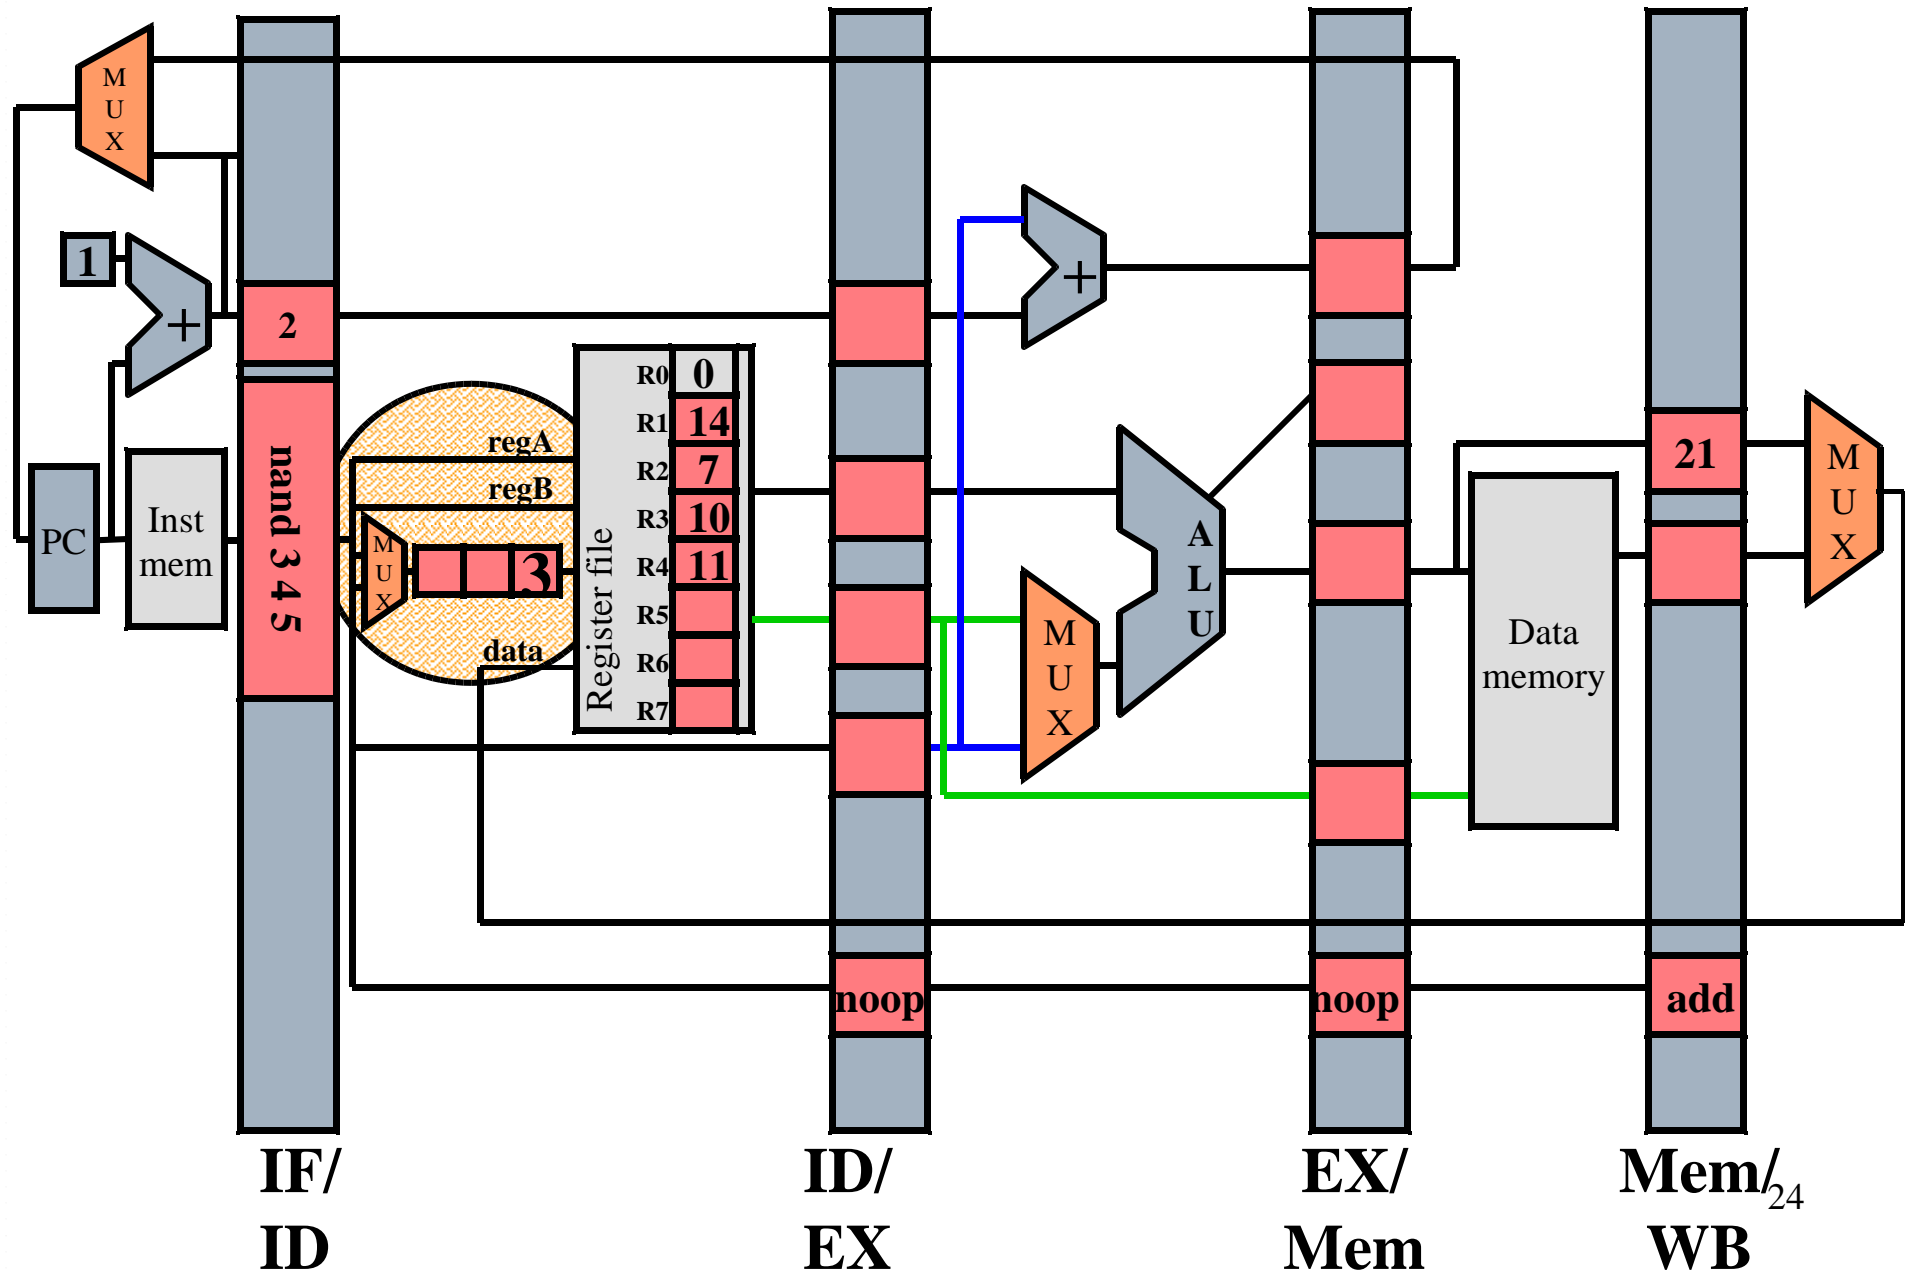
End of cycle 3



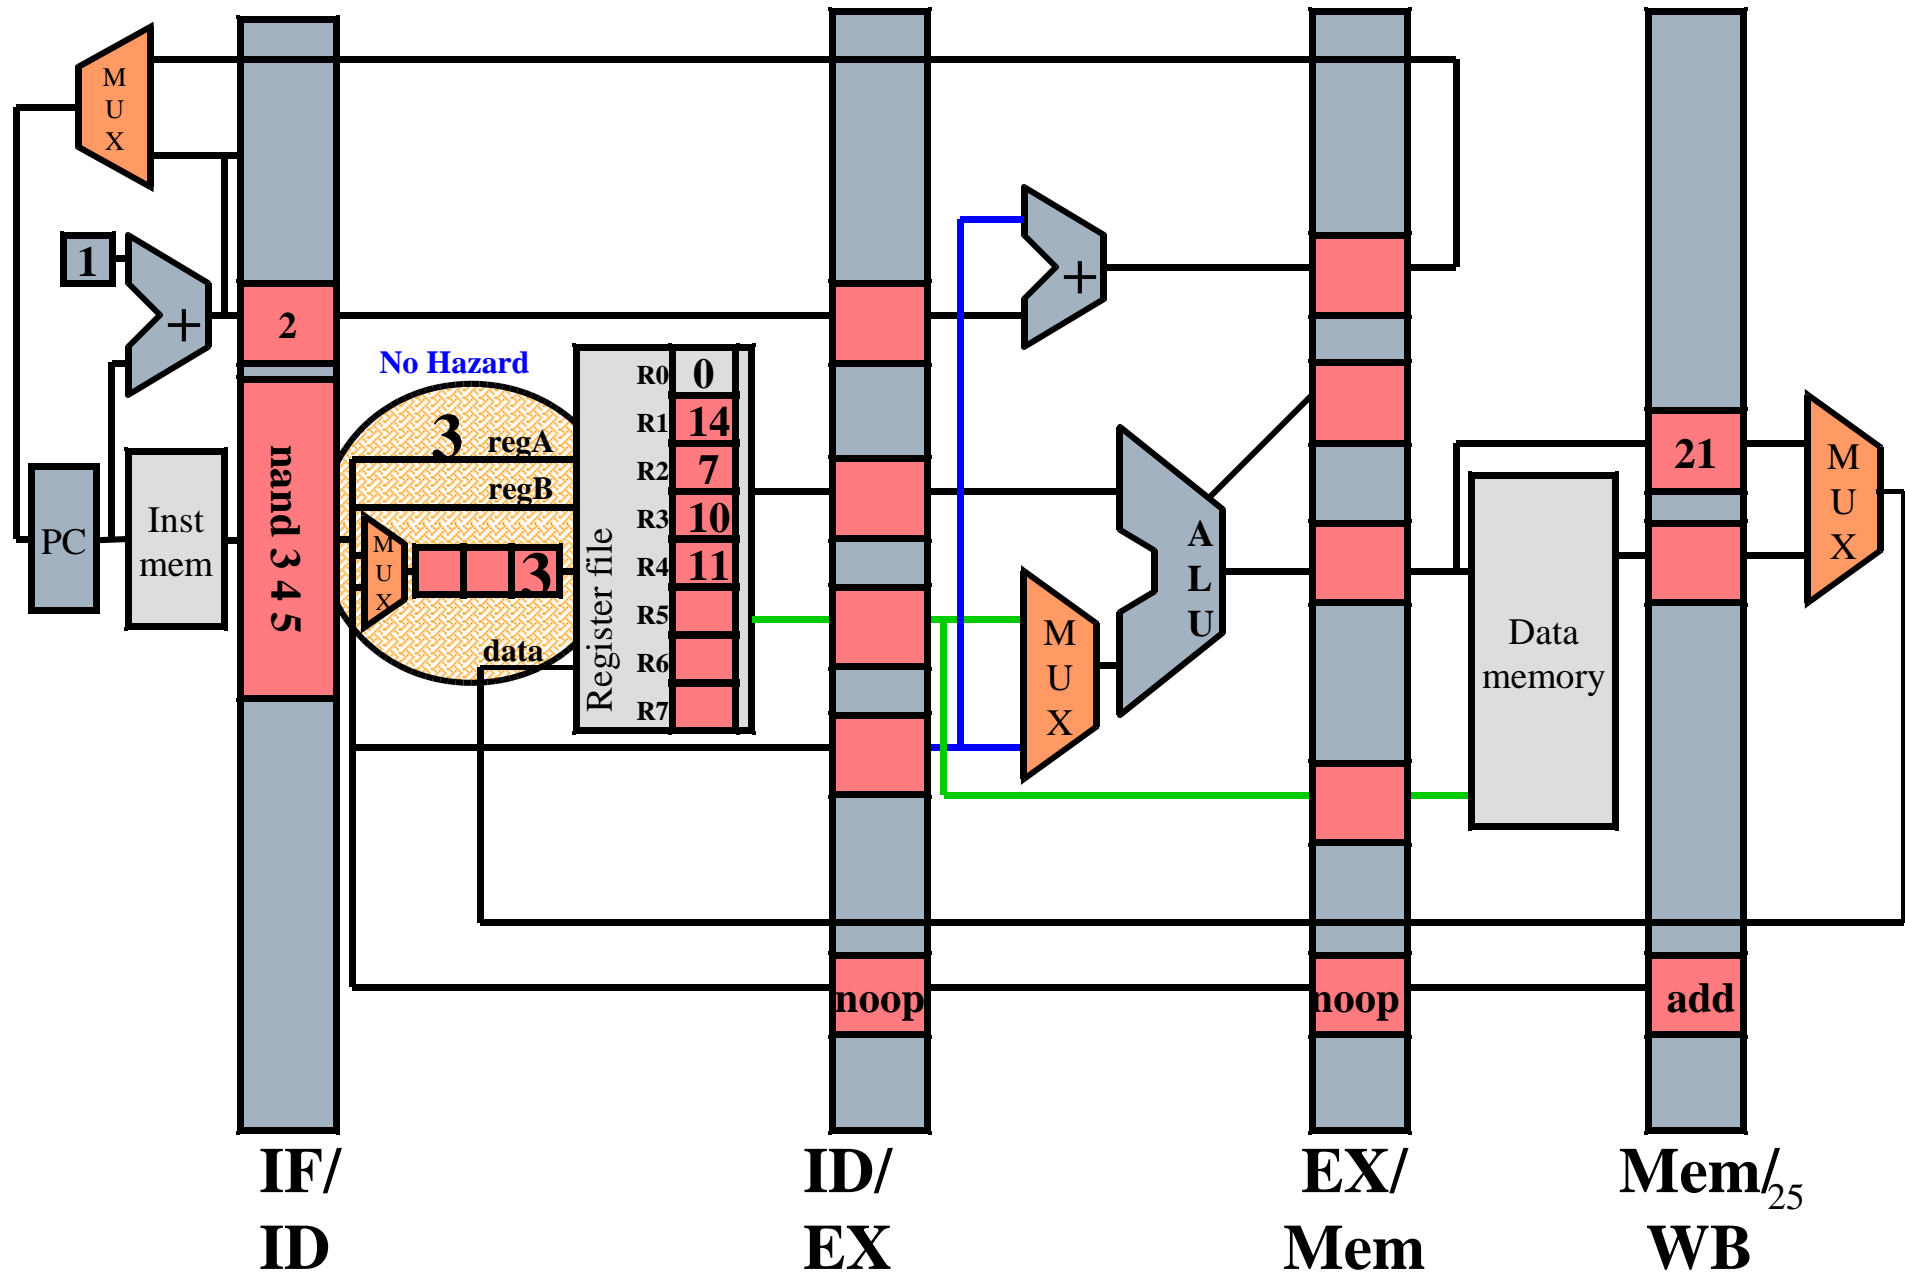
First half of cycle 4



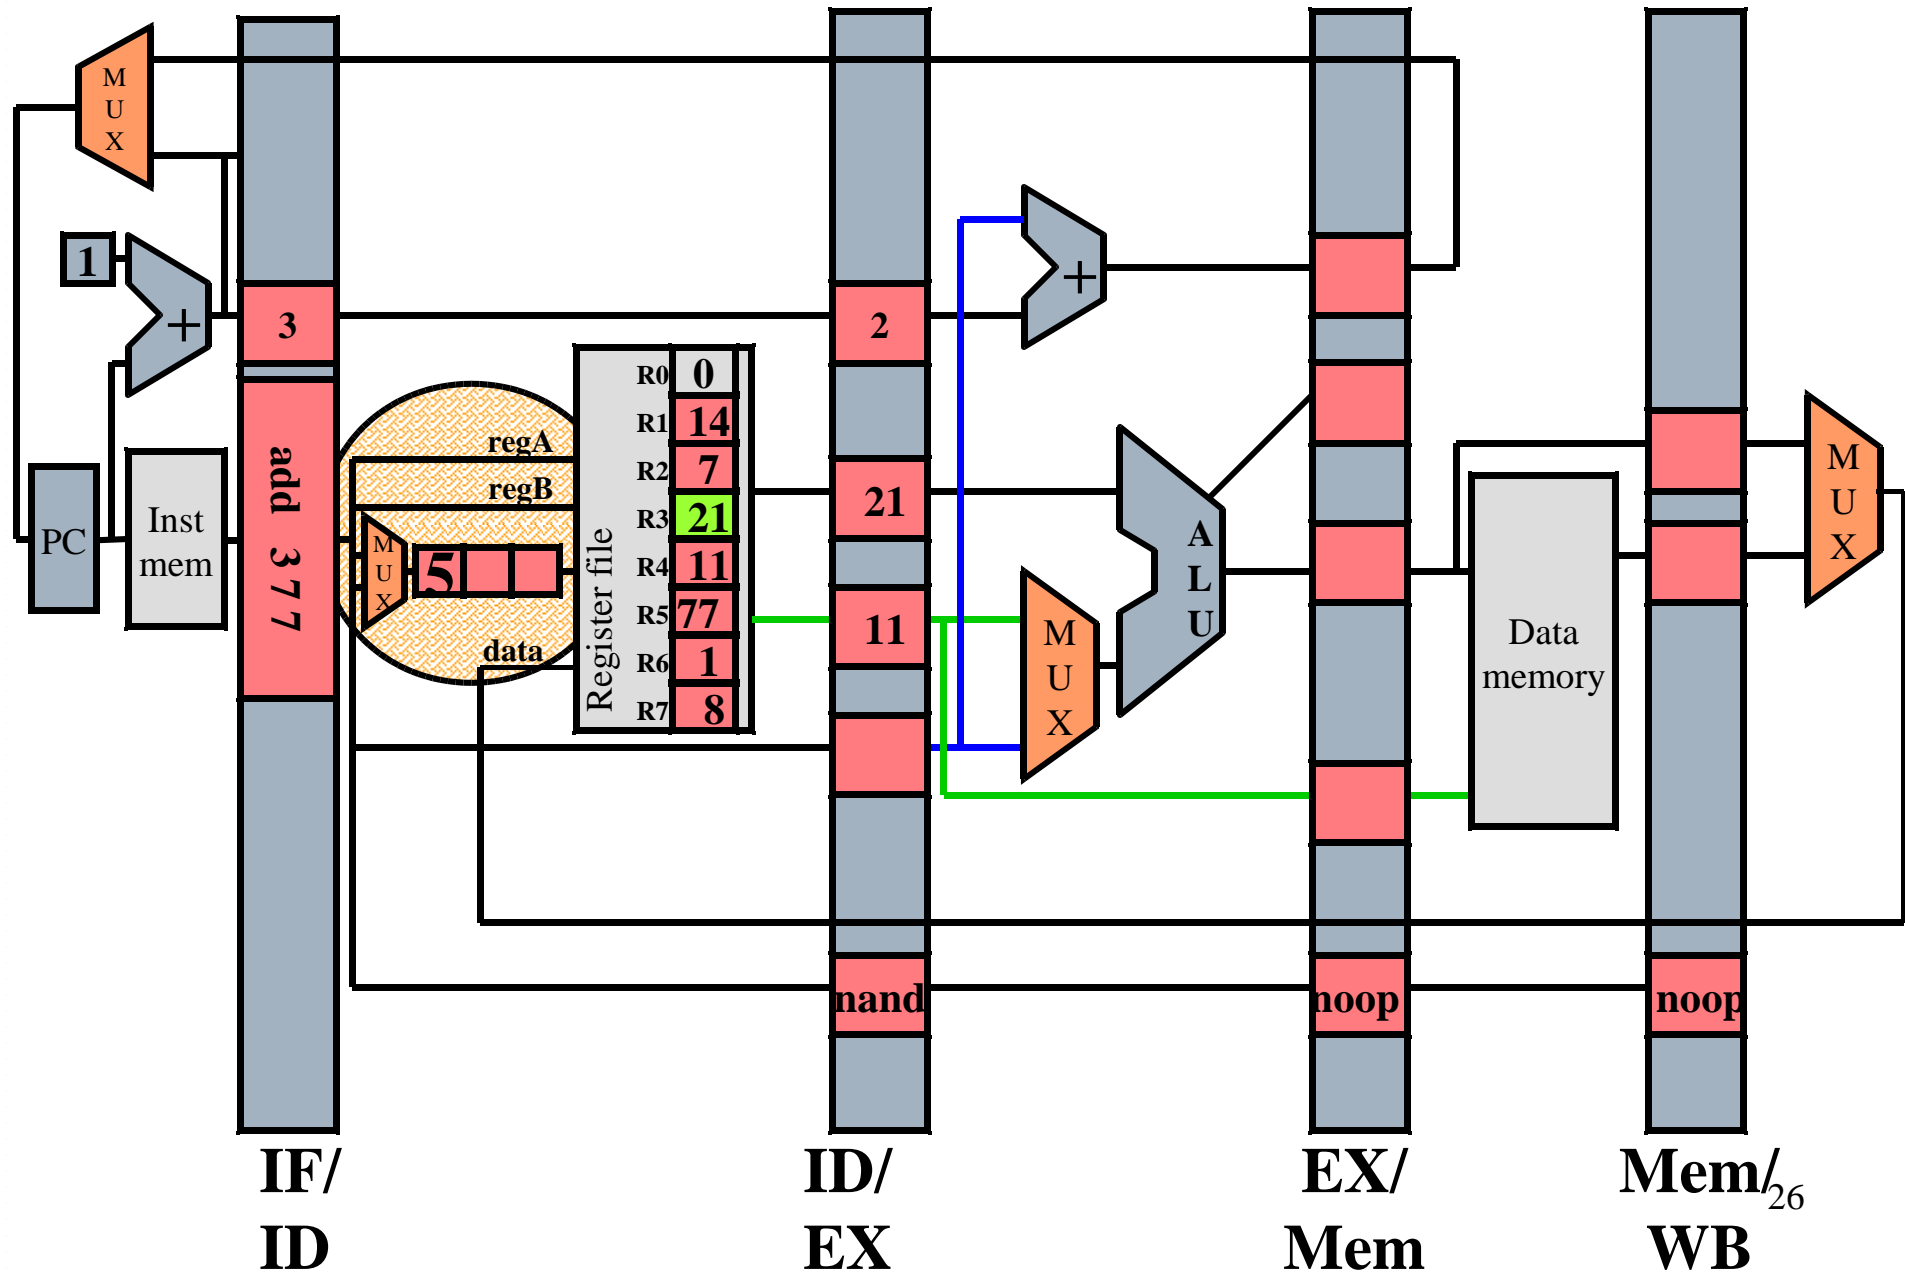
End of cycle 4



First half of cycle 5



End of cycle 5



Time Graph

Time:	1	2	3	4	5	6	7	8	9	10	11	12	13
add 1 2 3	IF	ID	EX	ME	WB								
nand 3 4 5		IF	no op	no op	ID	EX	ME	WB					
add 6 3 7					IF	ID	EX	ME	WB				
lw 3 6 10						IF	ID	EX	ME	WB			
sw 6 2 12							IF	no op	no op	ID	EX	ME	WB

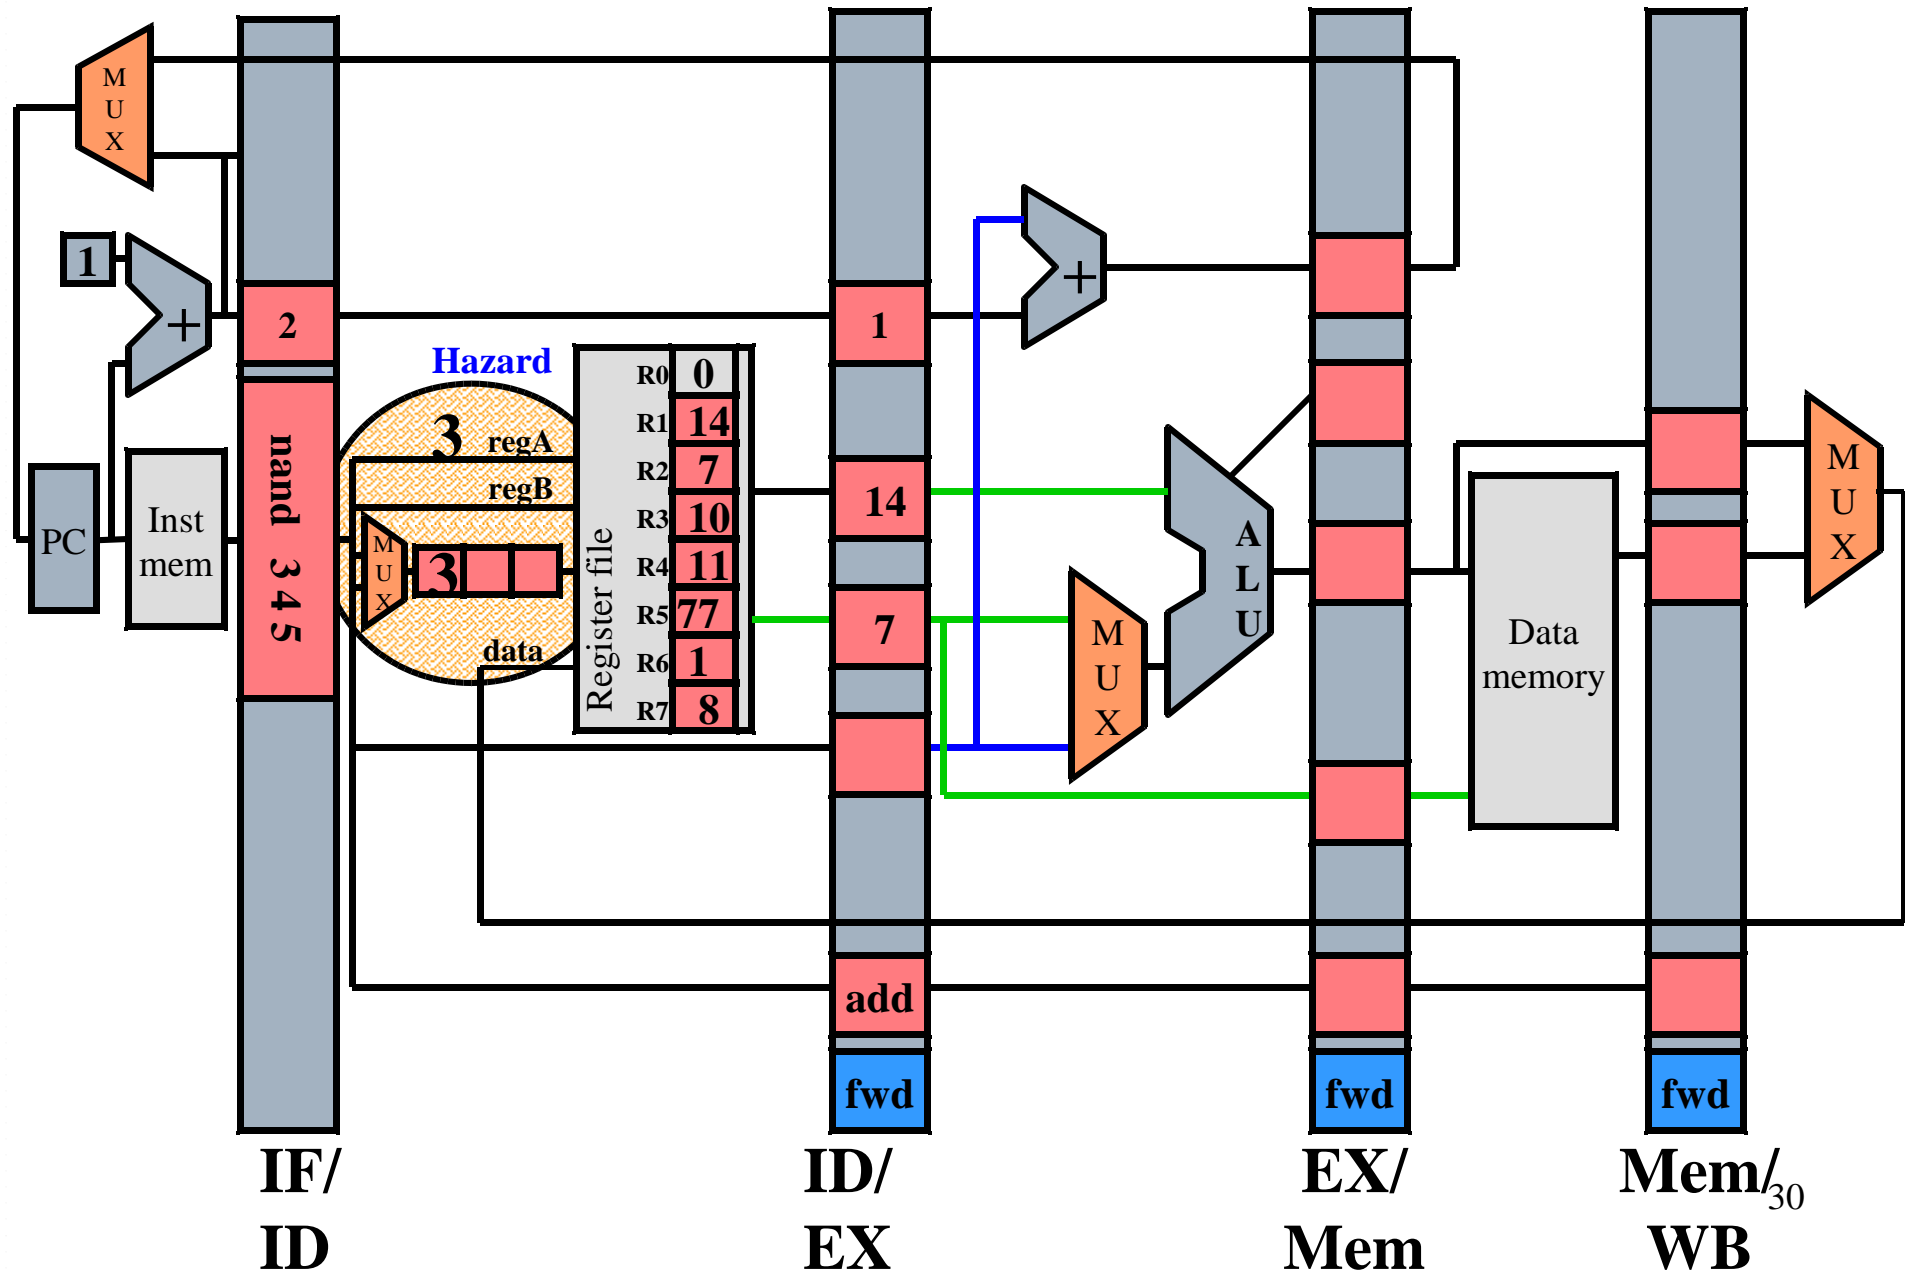
Problems with detect and stall

- ❑ CPI increases every time a hazard is detected!
- ❑ Is that necessary? Not always!
 - Re-route the result of the add to the nand
 - nand no longer needs to read R3 from reg file
 - It can get the data later (when it is ready)
 - This lets us complete the decode this cycle
 - But we need more control to remember that the data that we aren't getting from the reg file at this time will be found elsewhere in the pipeline at a later cycle.

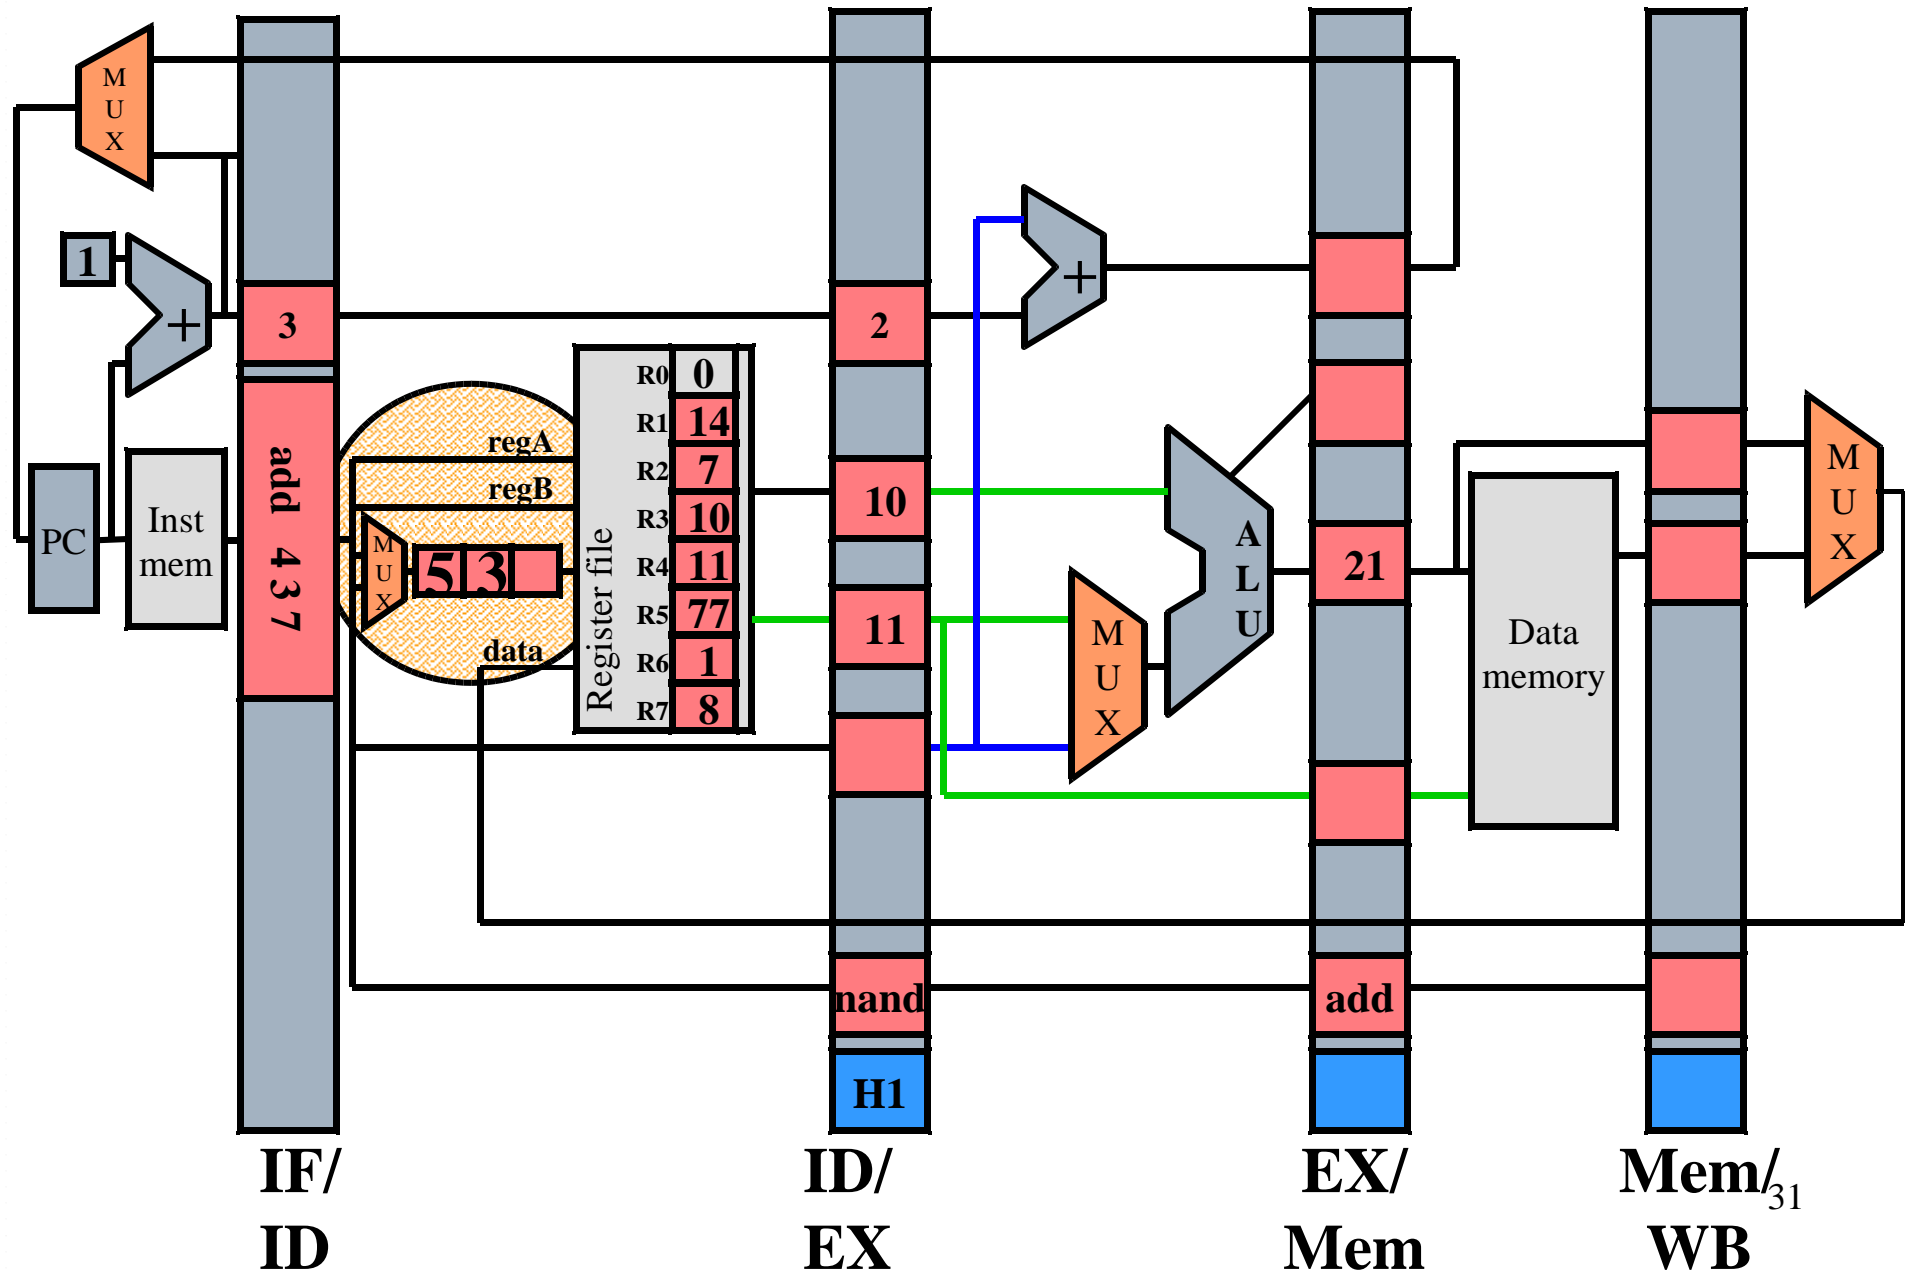
Handling data hazards III: Detect and forward

- ❑ Detect: same as detect and stall
 - Except that all 4 hazards are treated differently
 - I.e., you can't logical-OR the 4 hazard signals
- ❑ Forward:
 - New **bypass datapaths** route computed data to where it is needed
 - New MUX and control to pick the right data
- ❑ **Beware:** Stalling may still be required even in the presence of forwarding

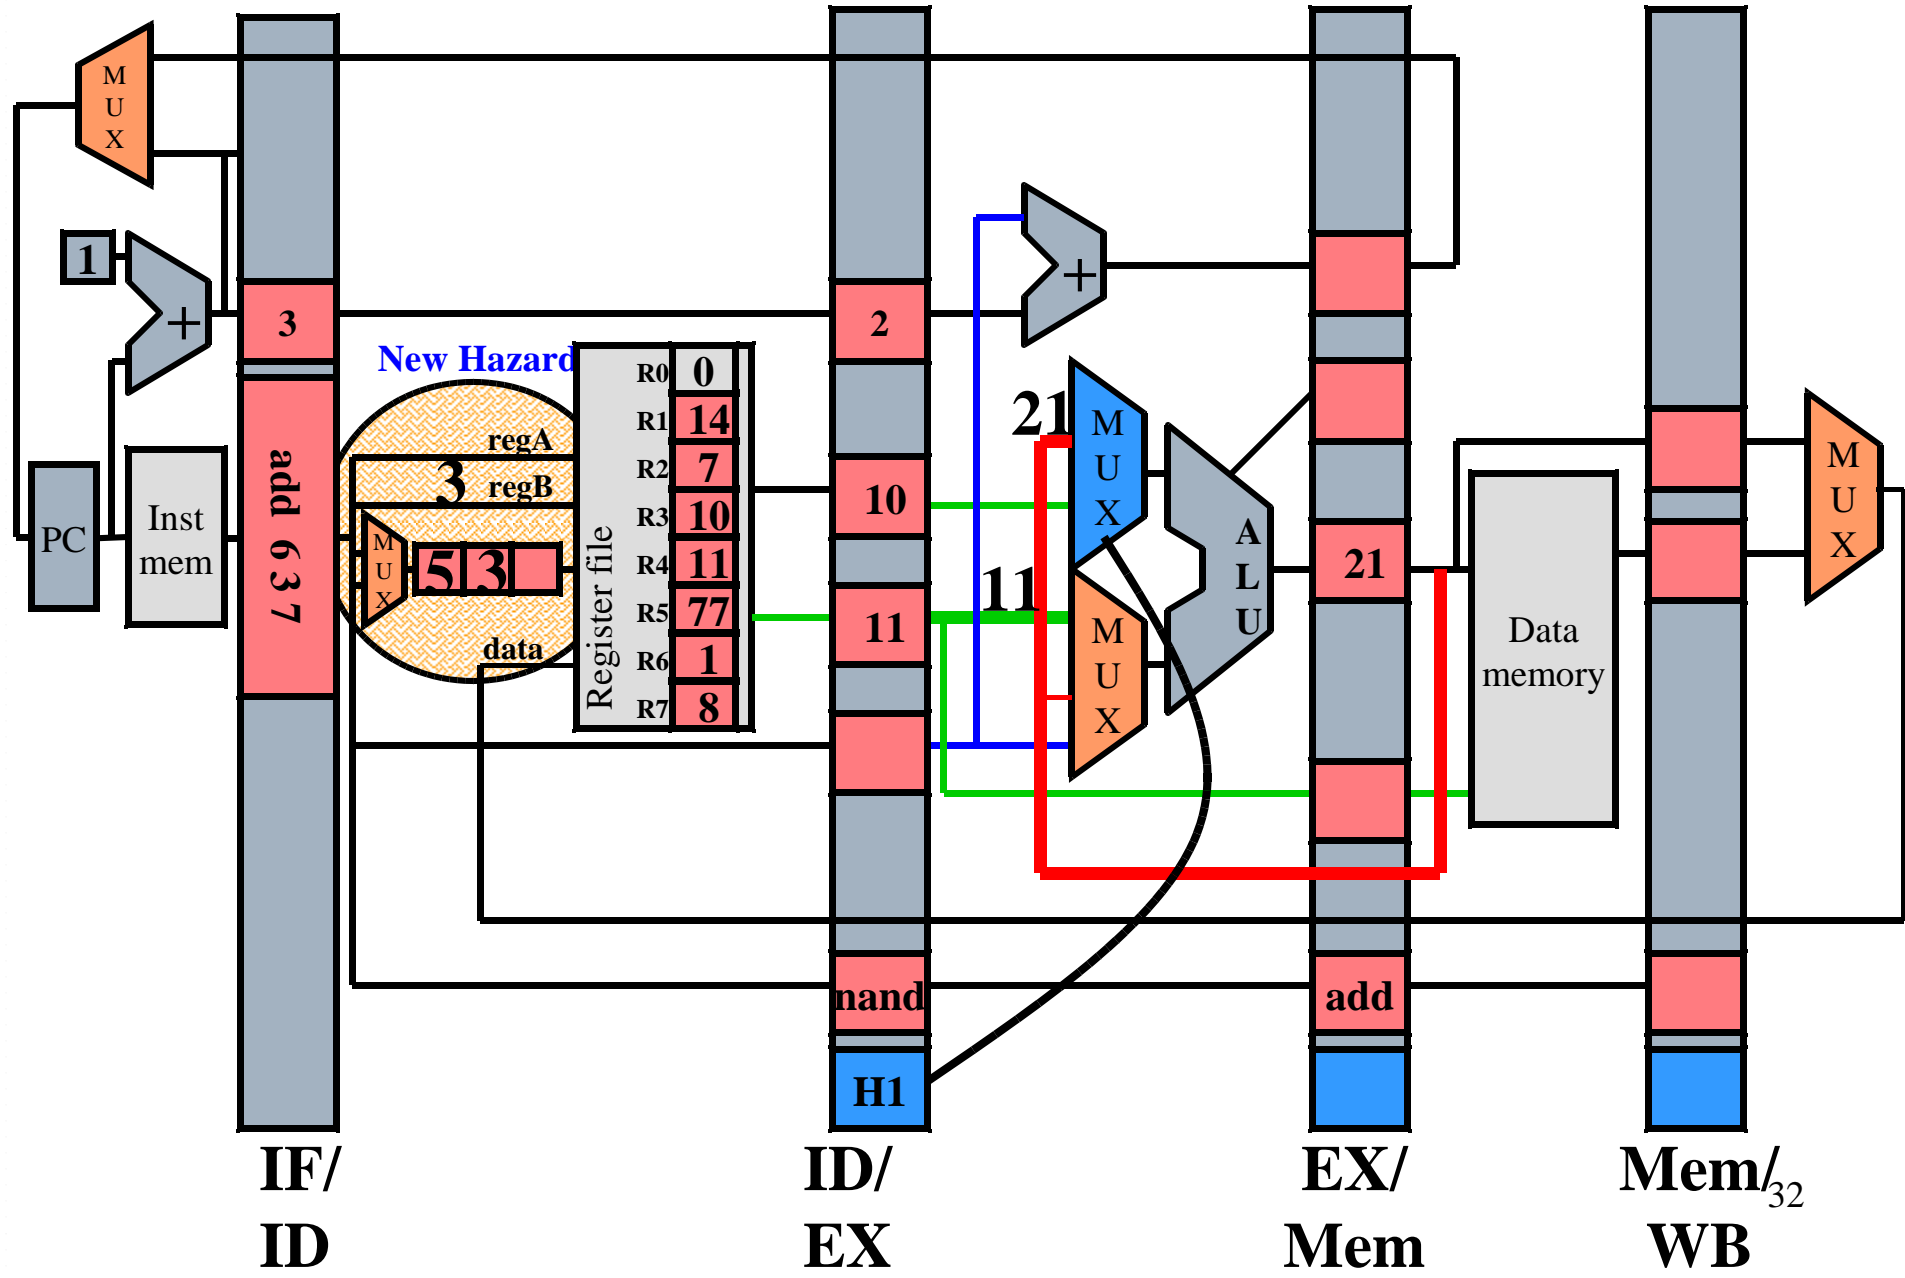
First half of cycle 3



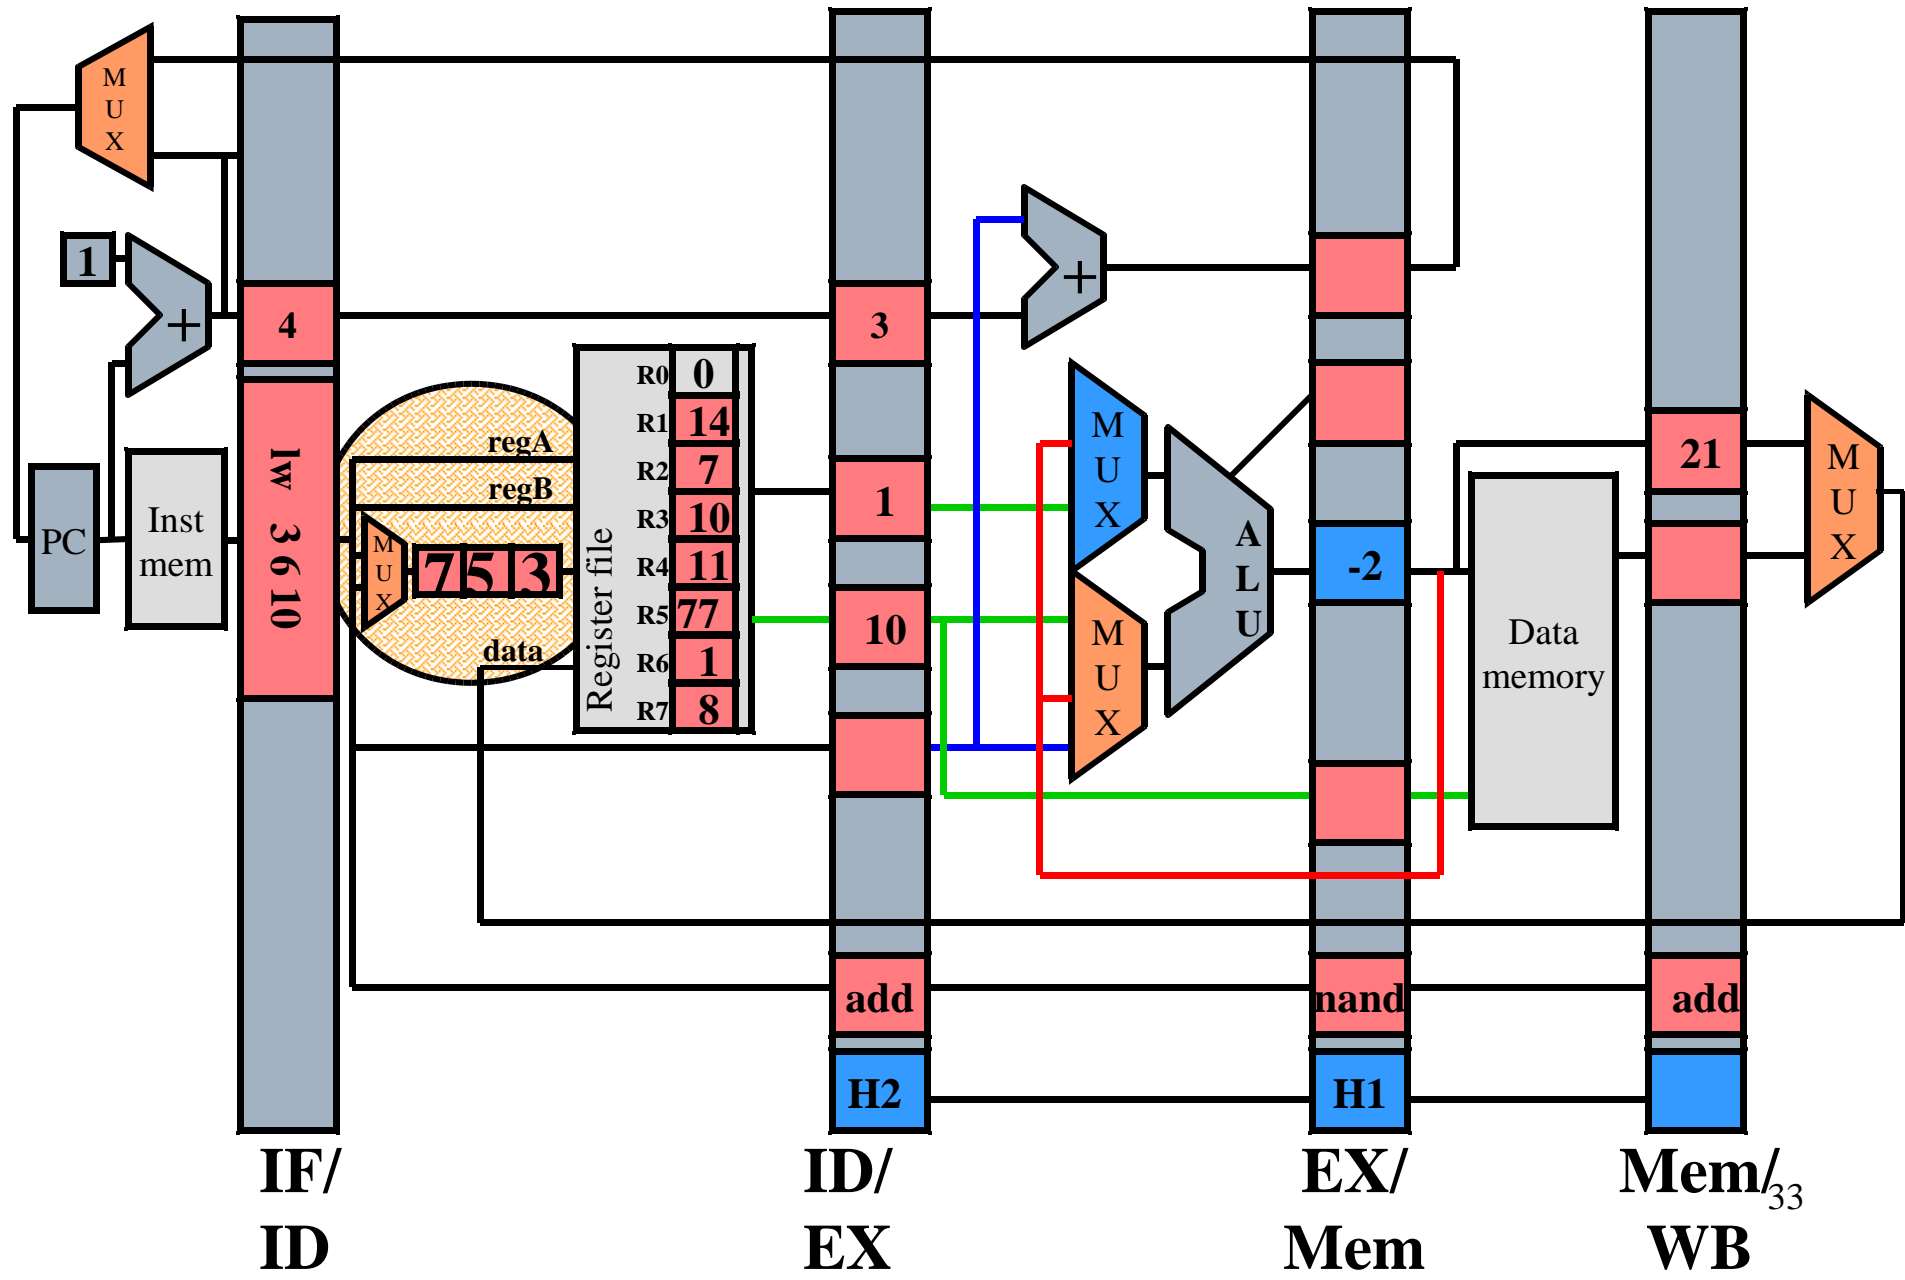
End of cycle 3



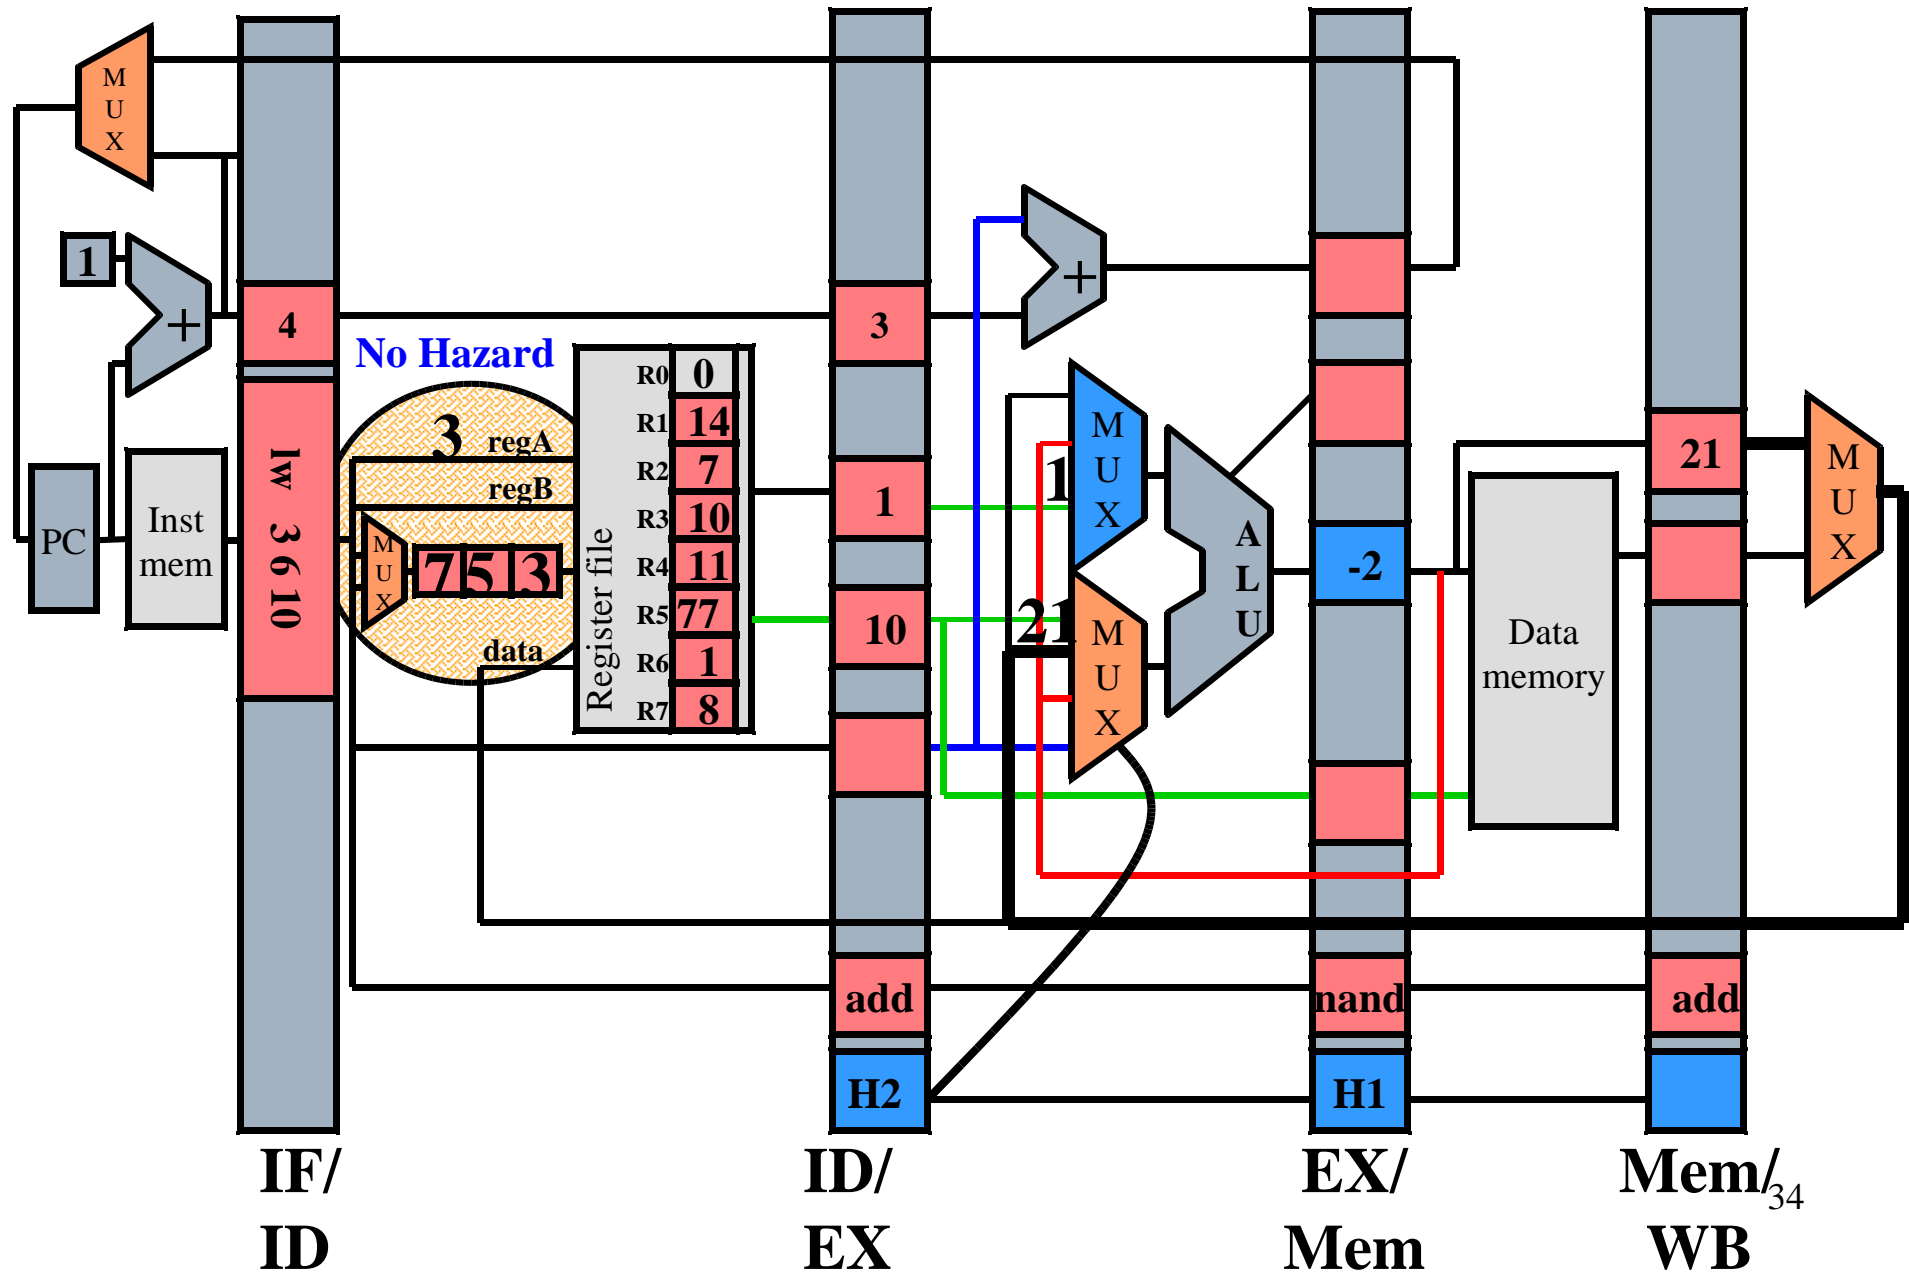
First half of cycle 4



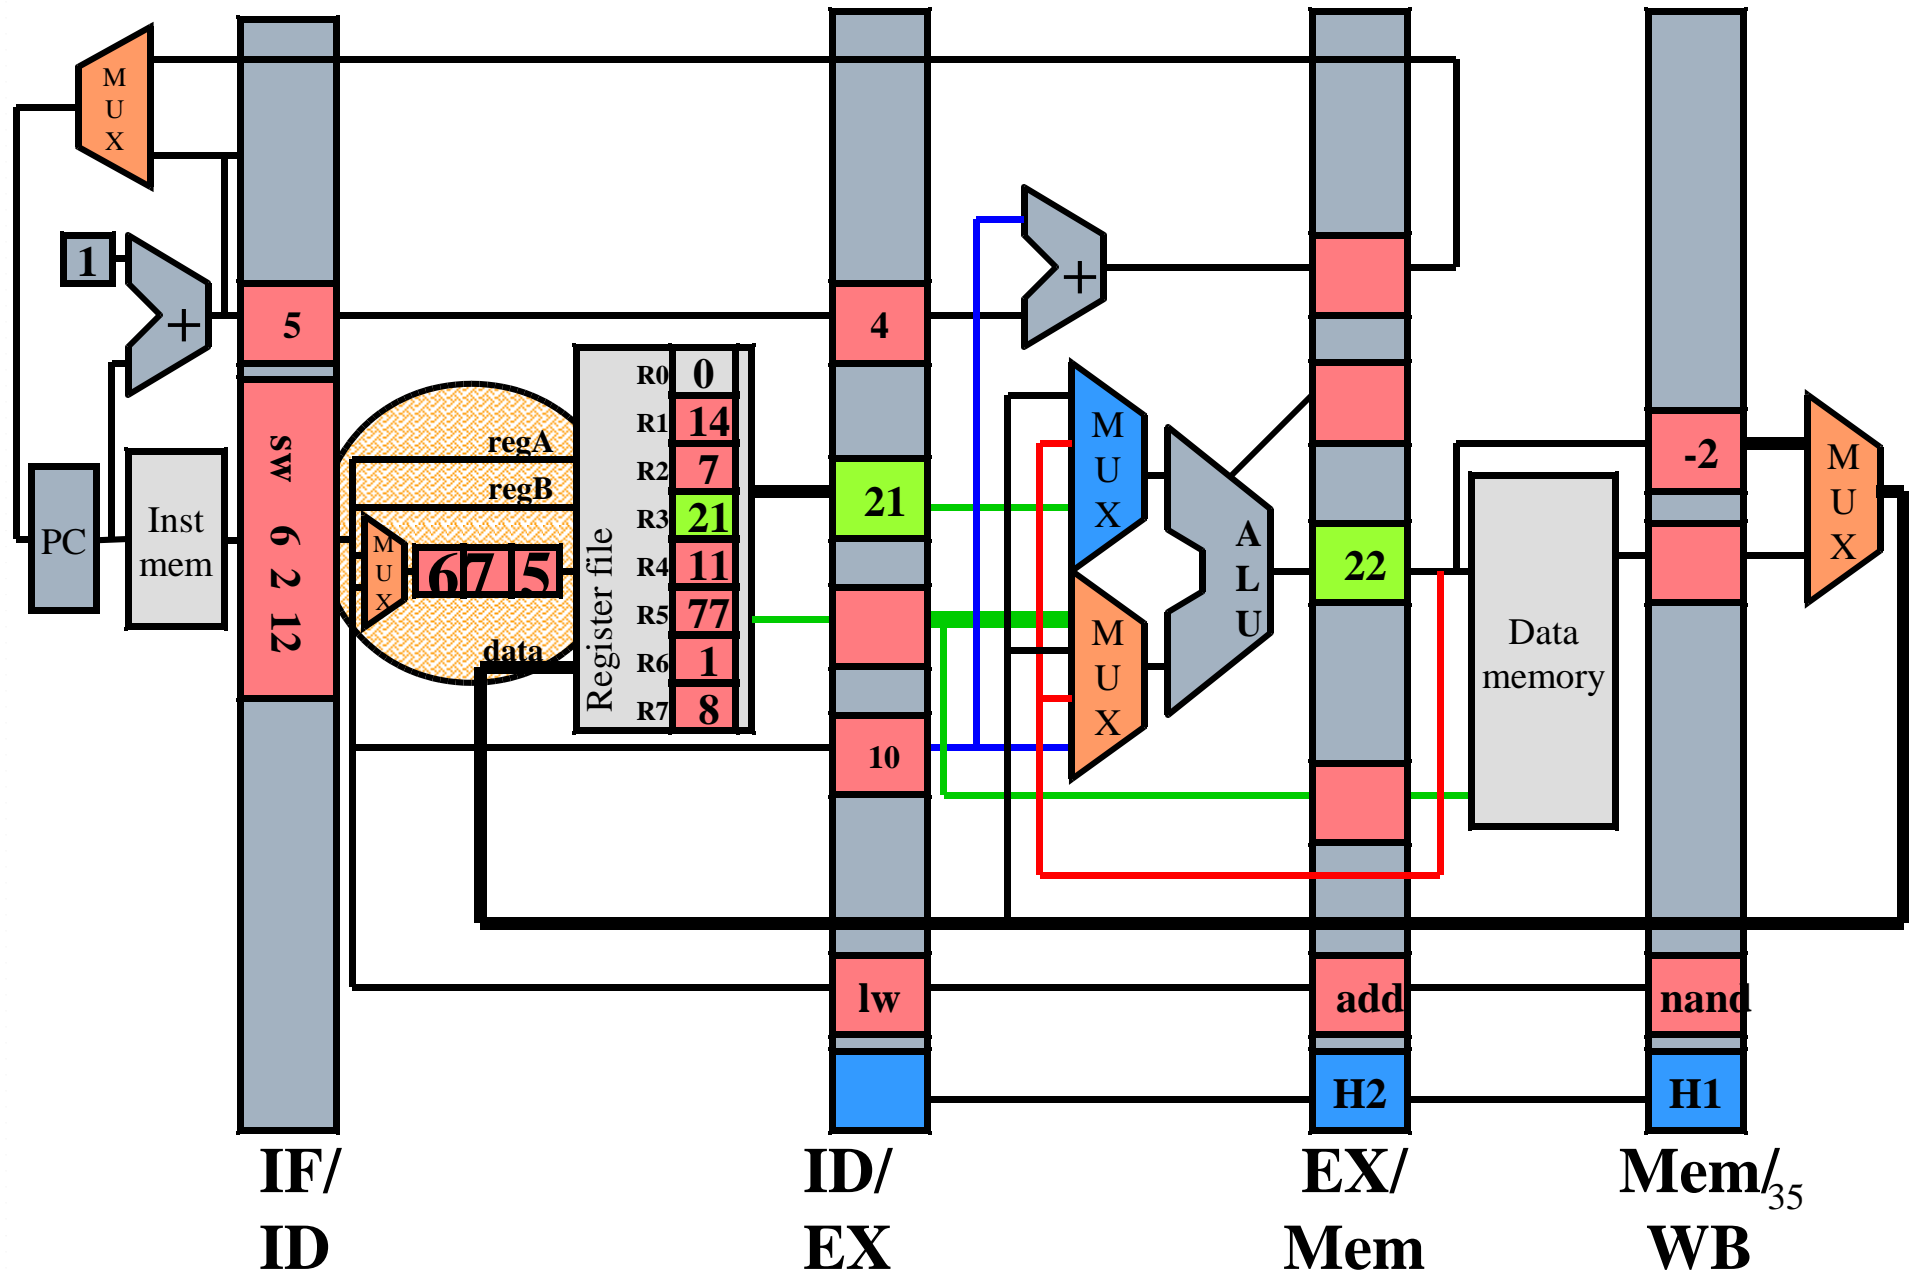
End of cycle 4



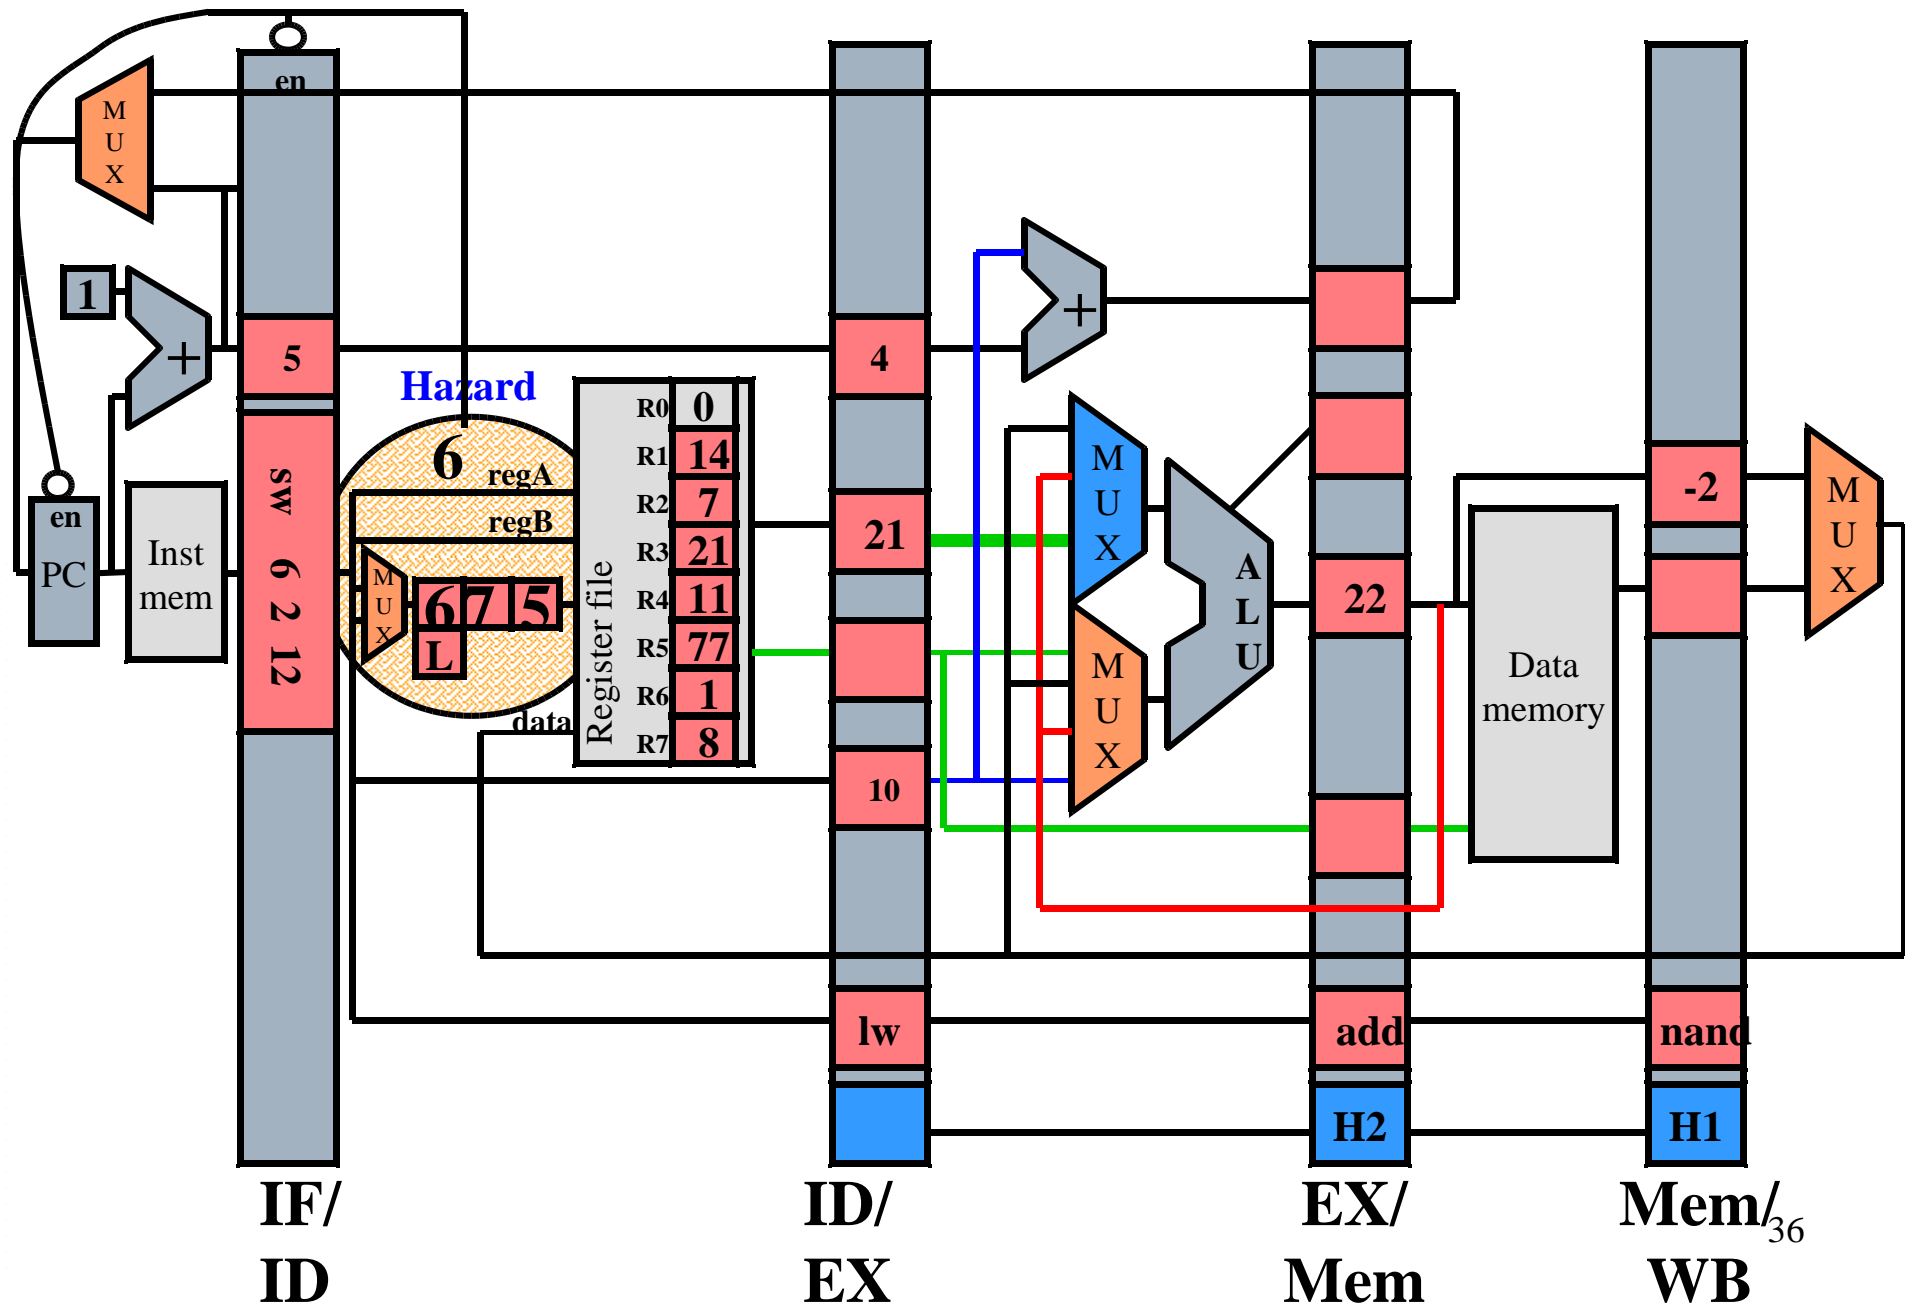
First half of cycle 5



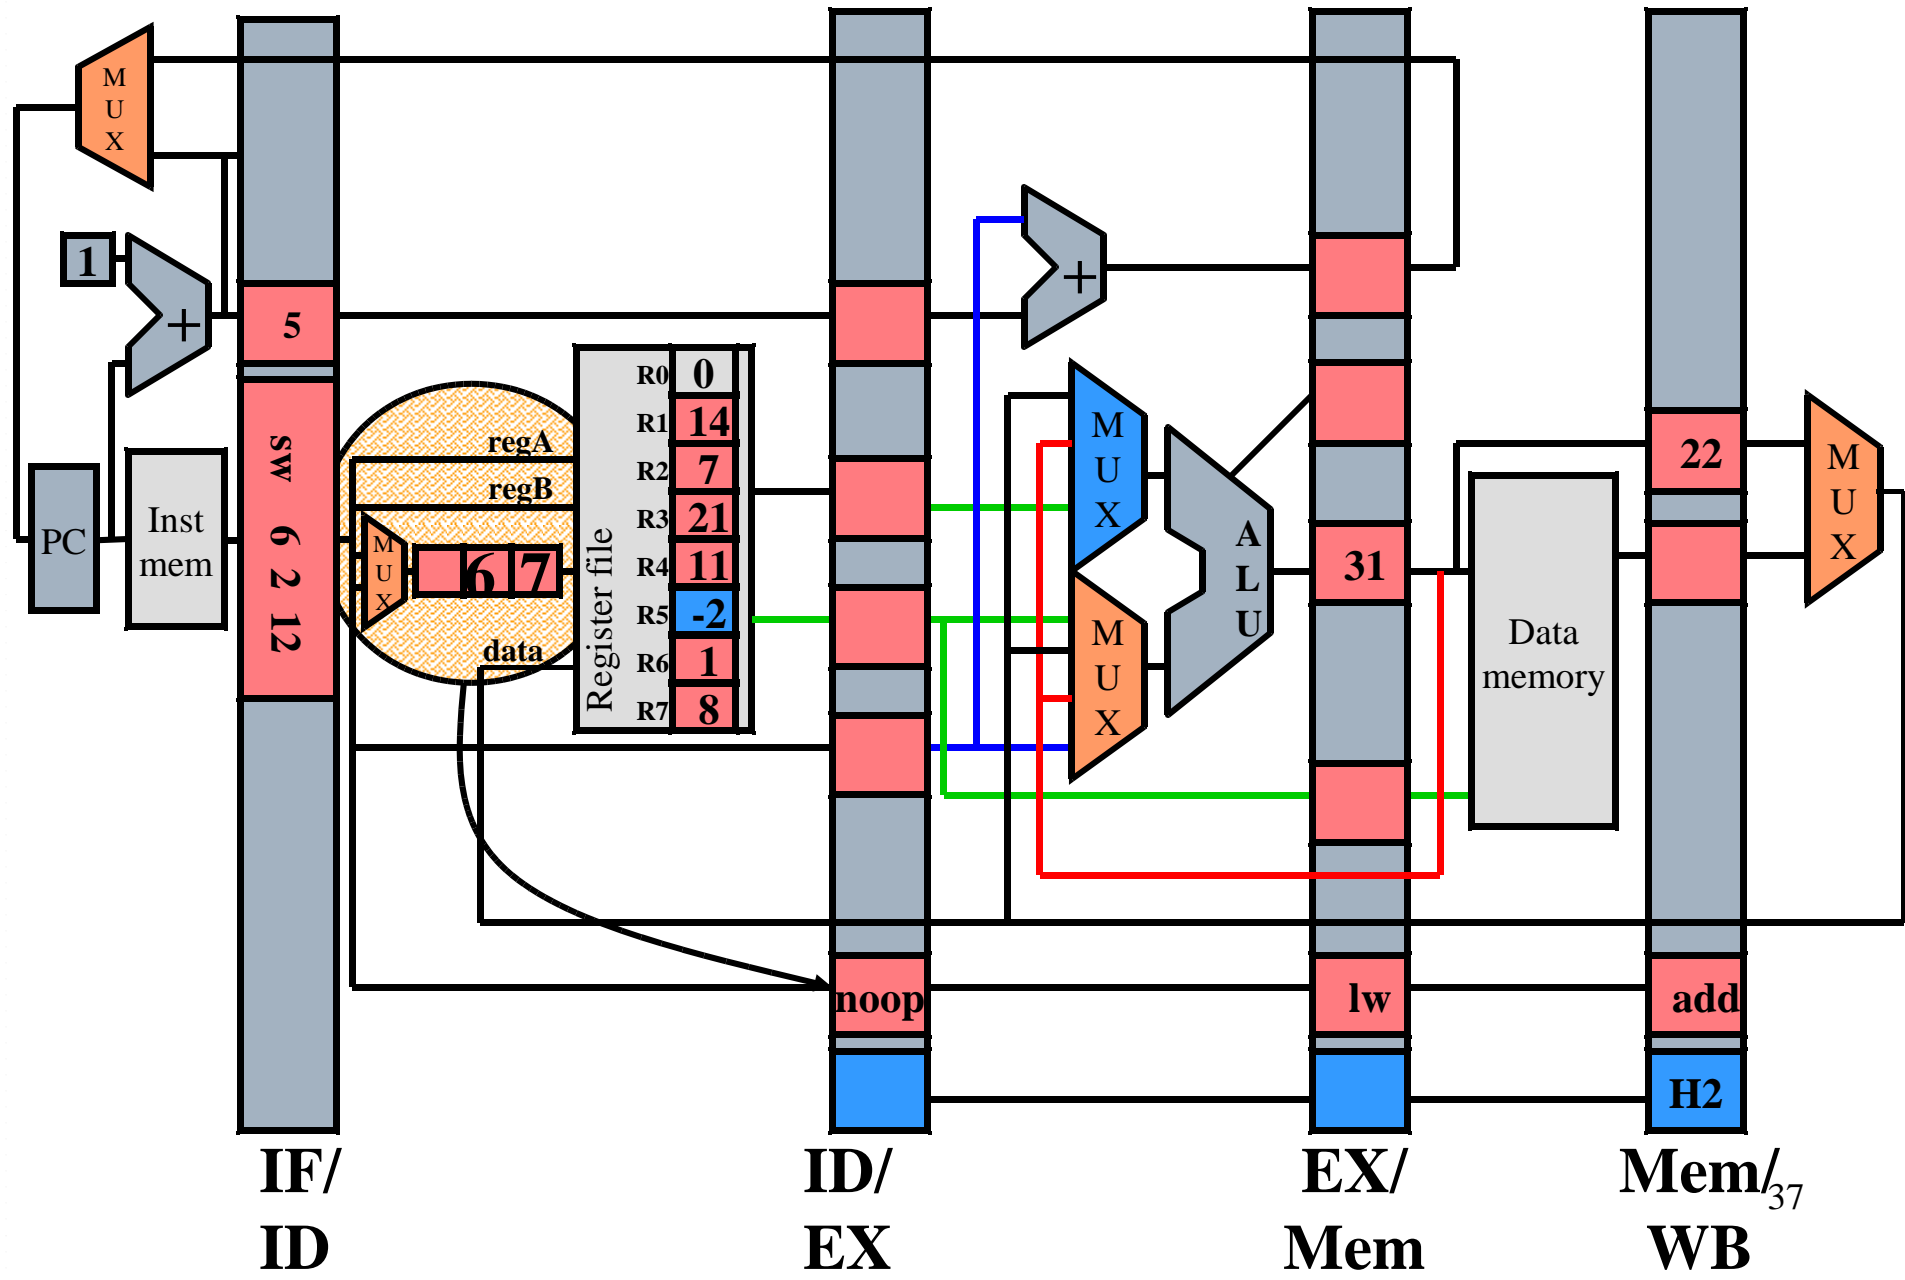
End of cycle 5



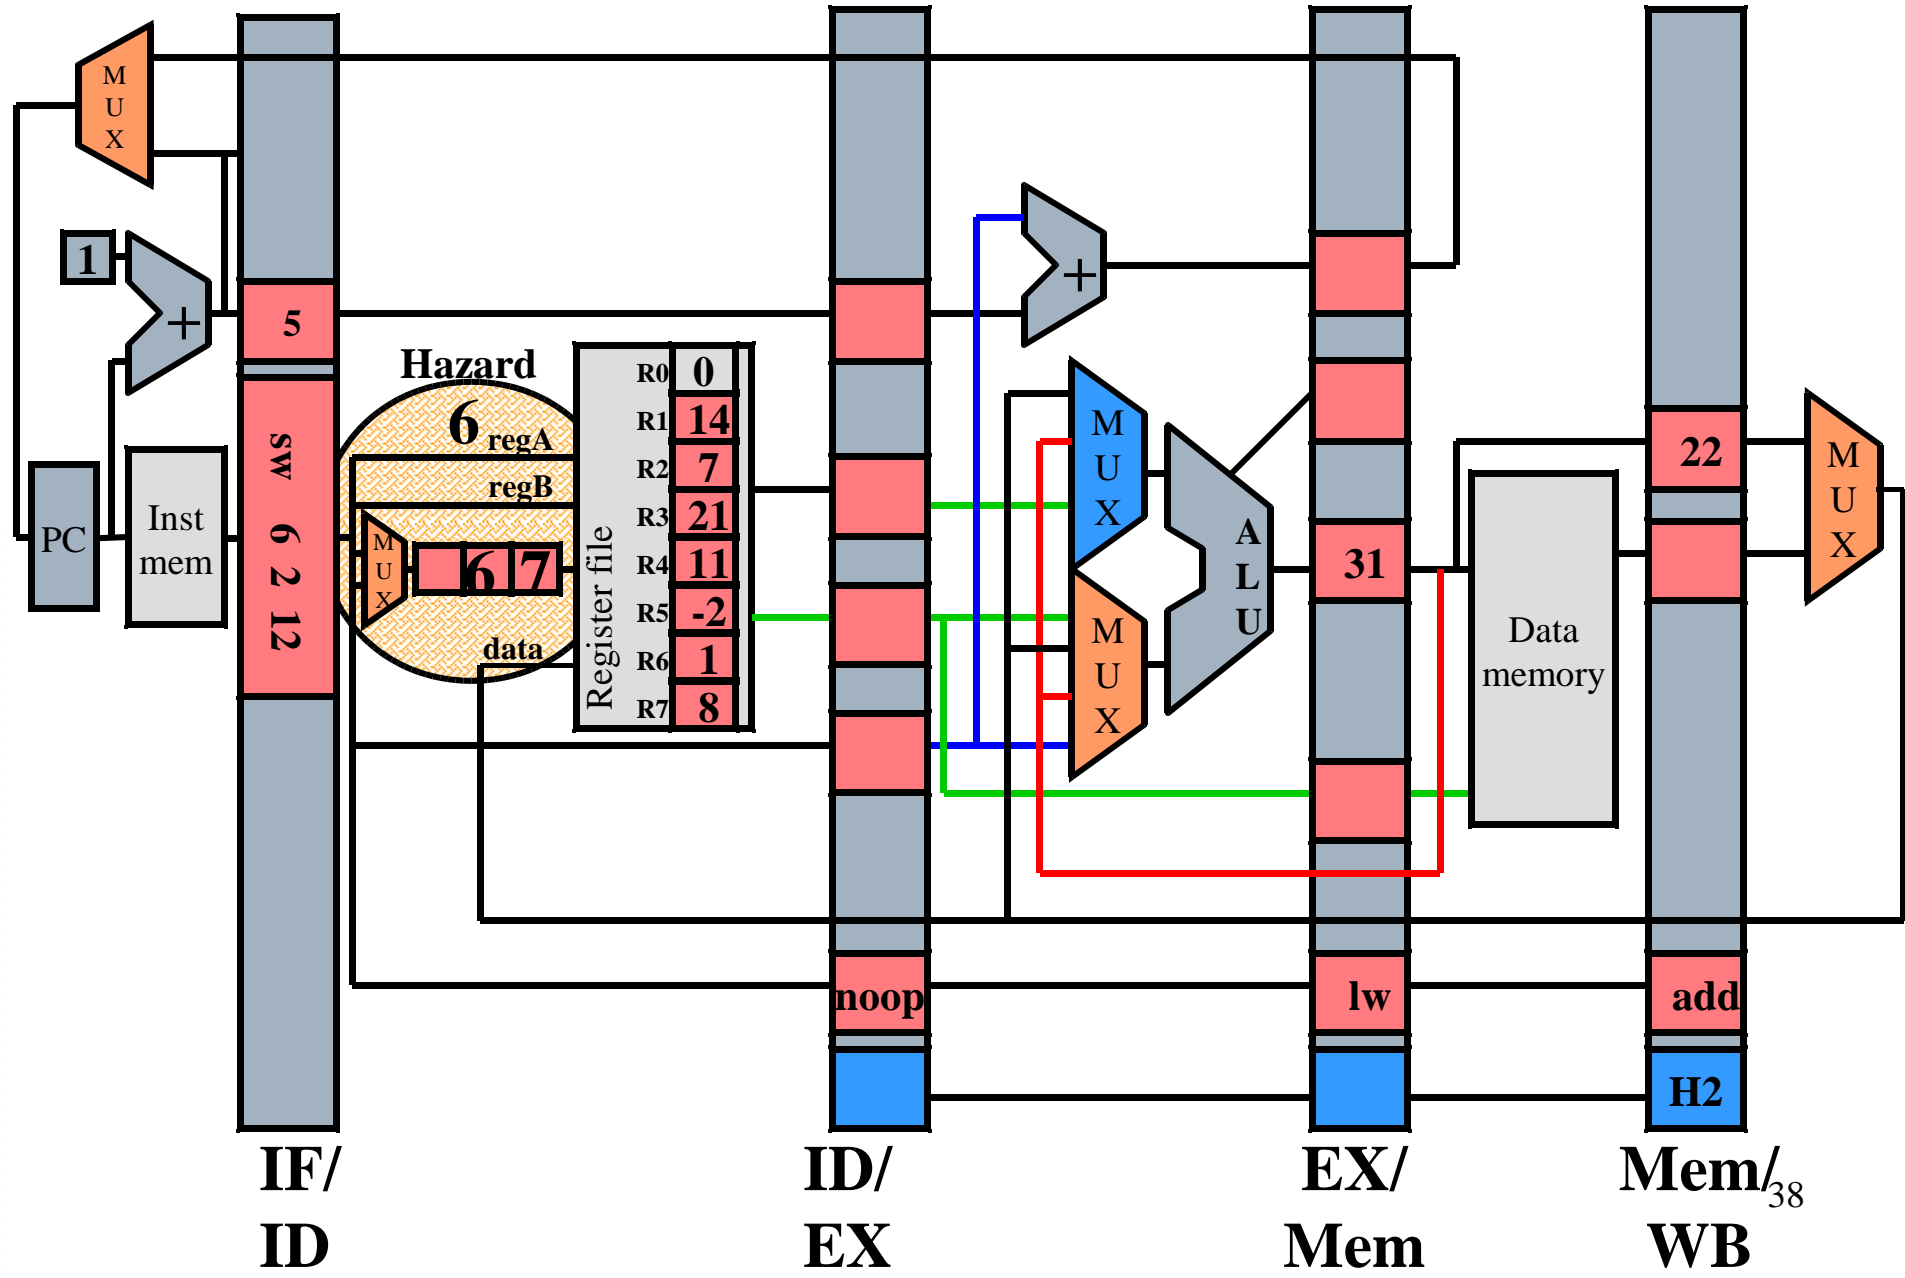
First half of cycle 6



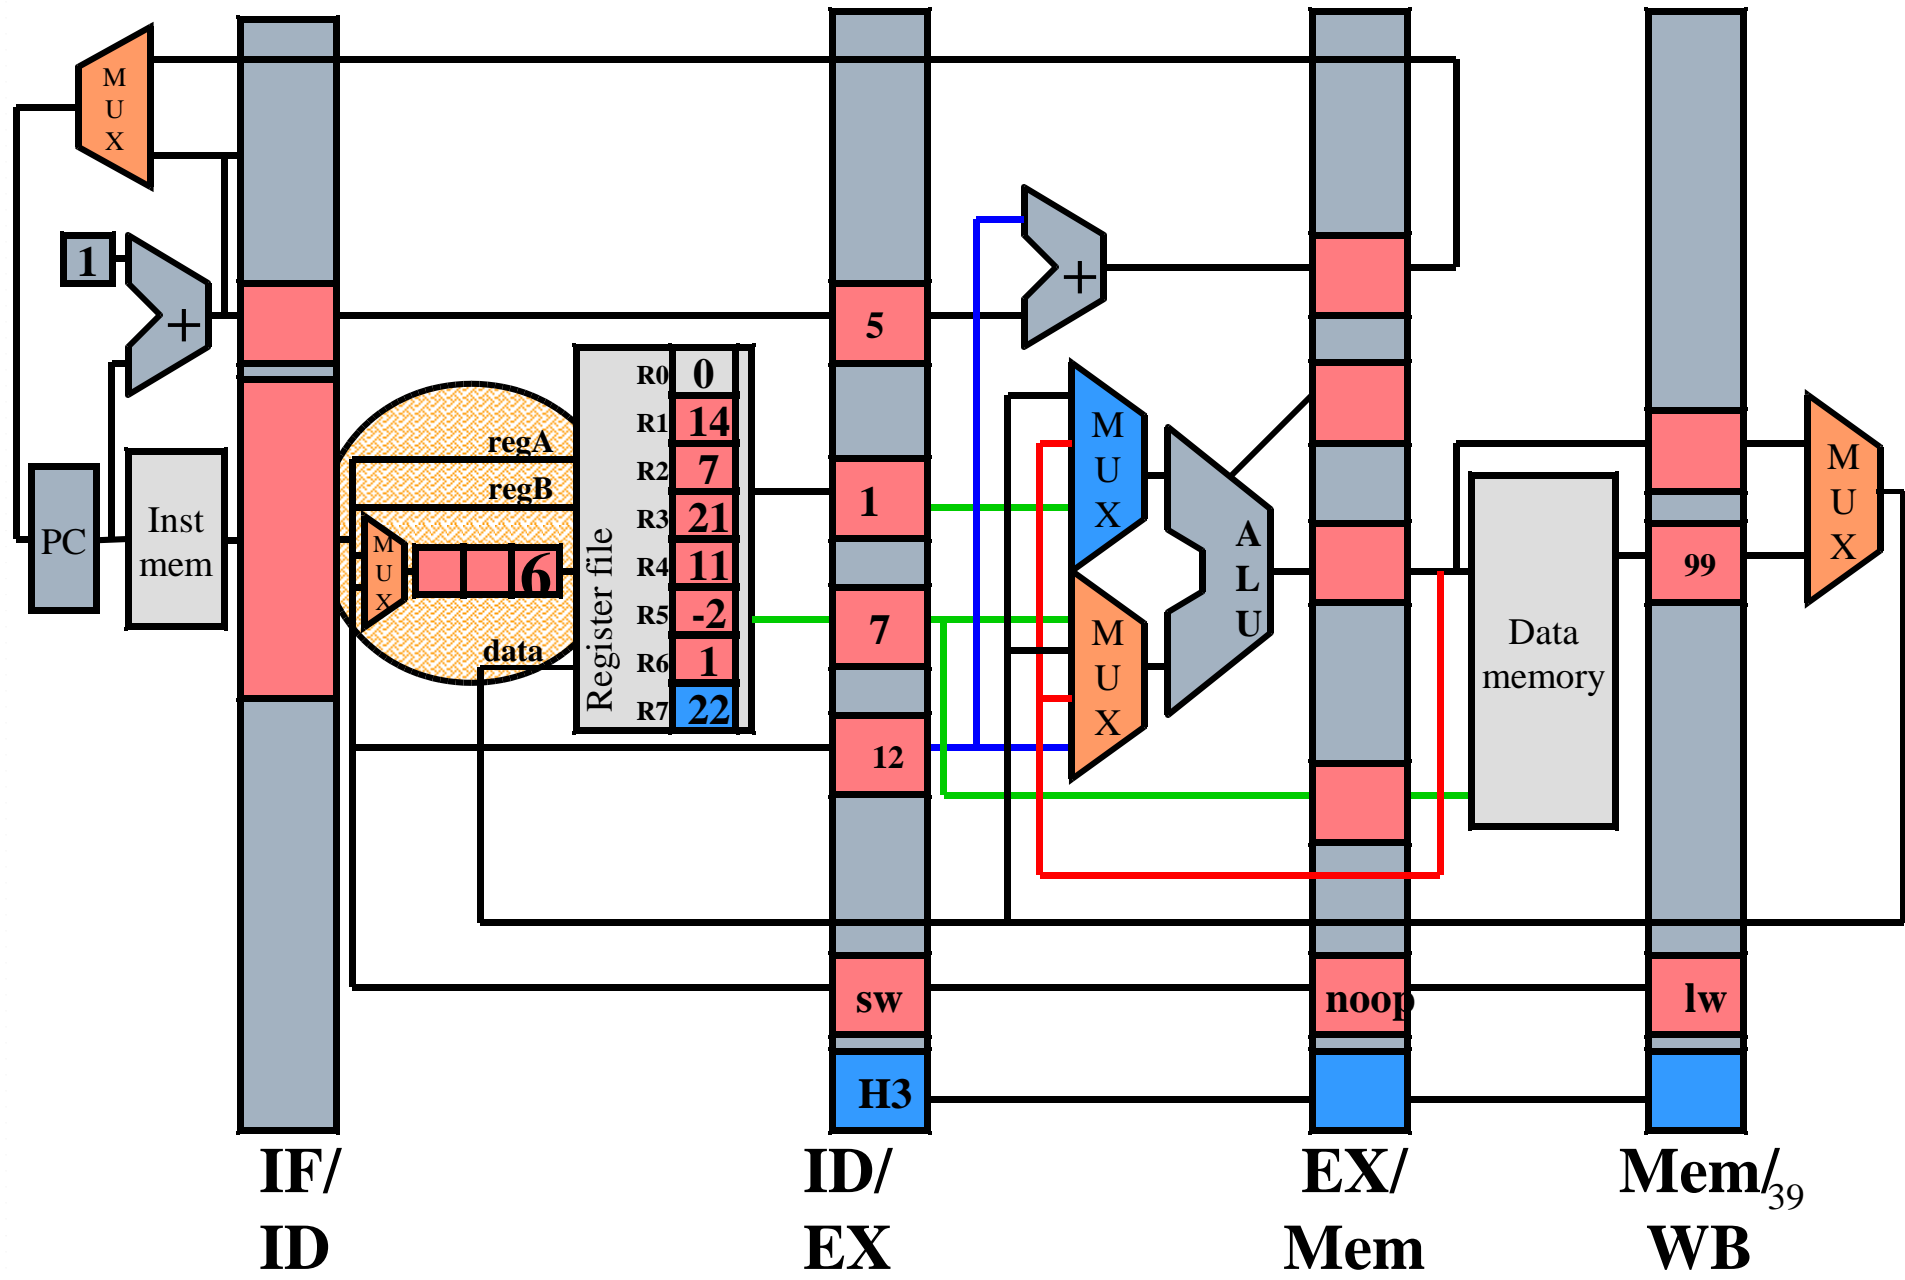
End of cycle 6



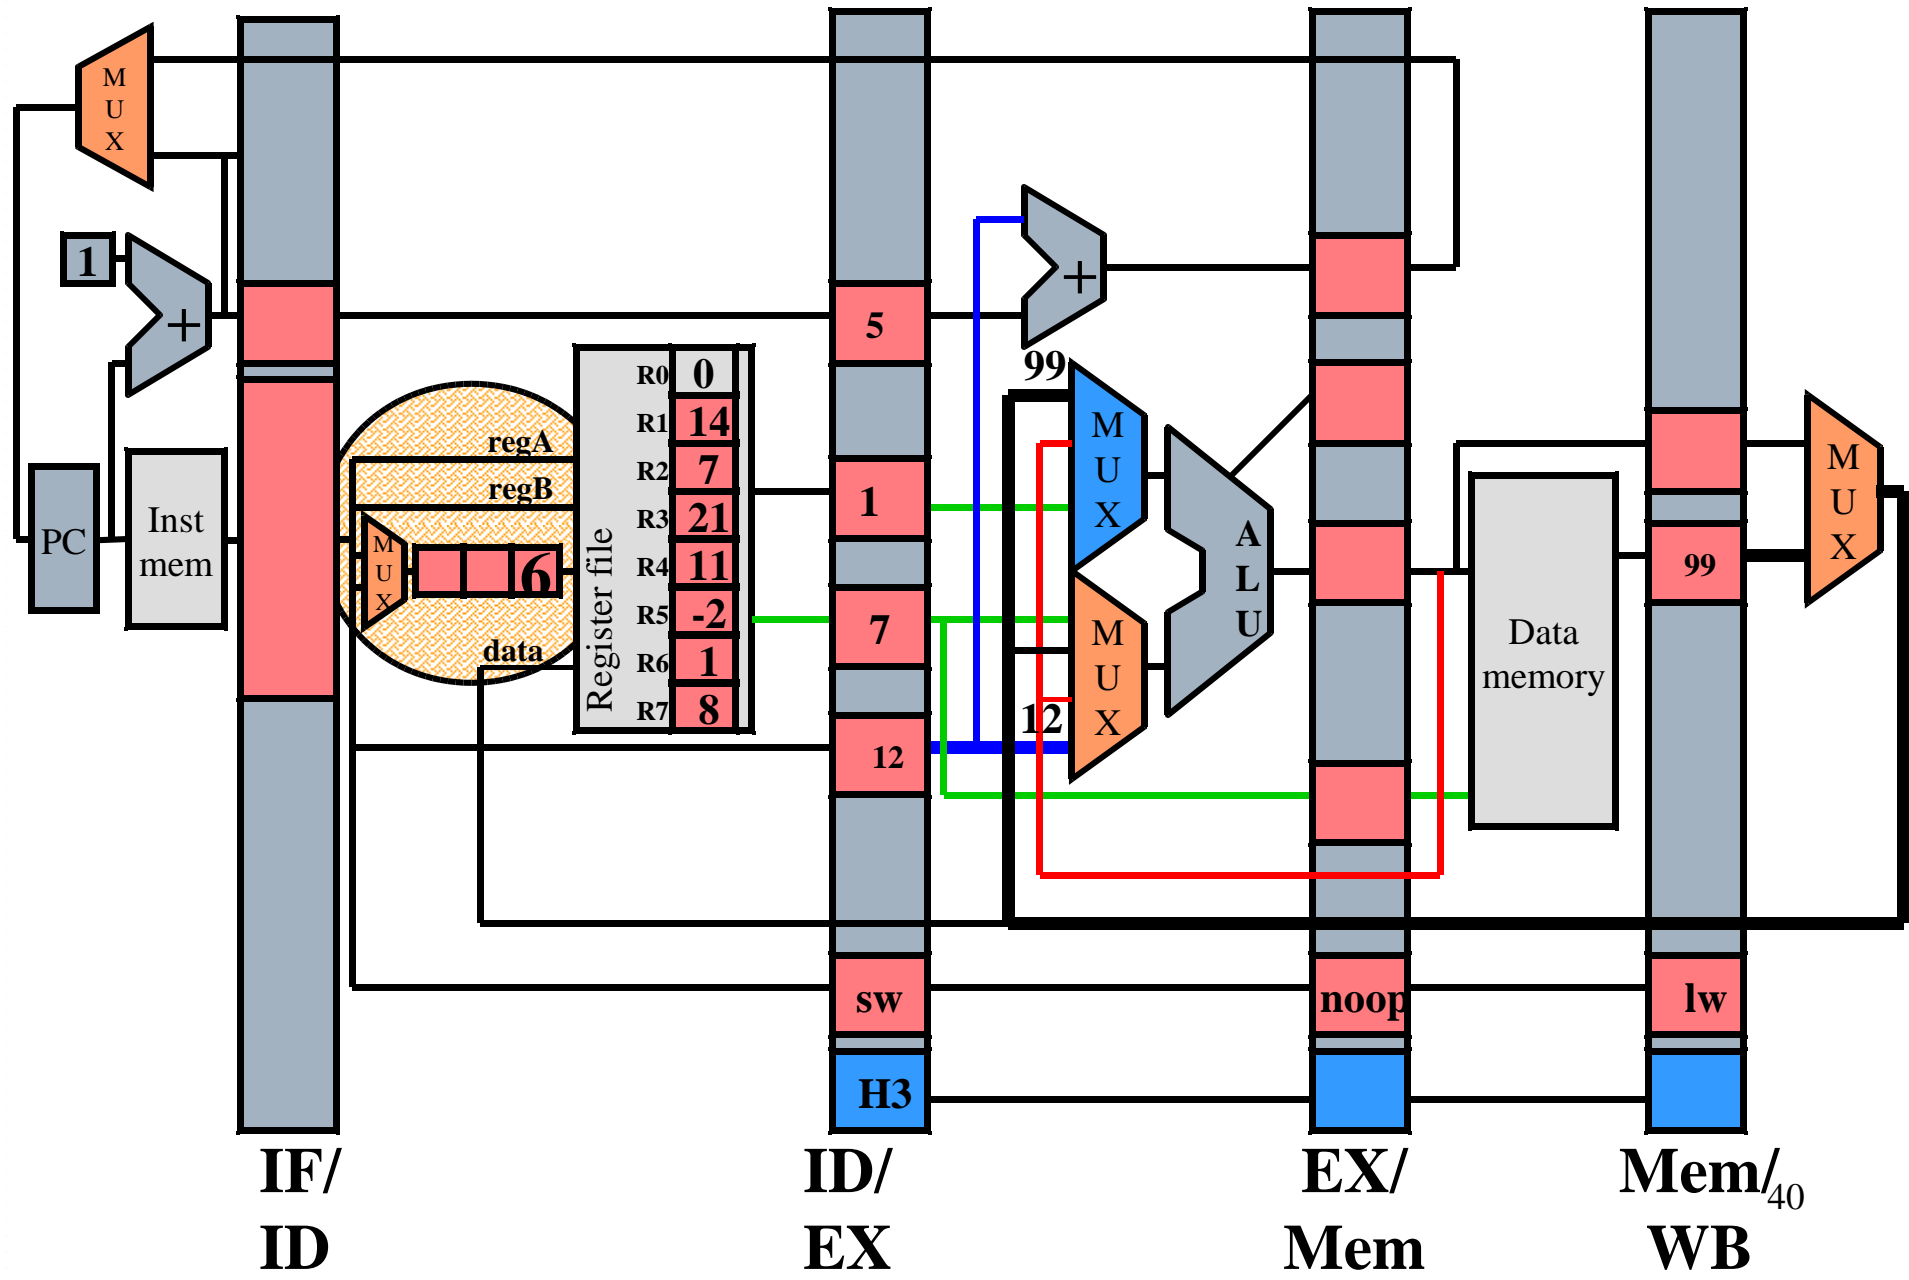
First half of cycle 7



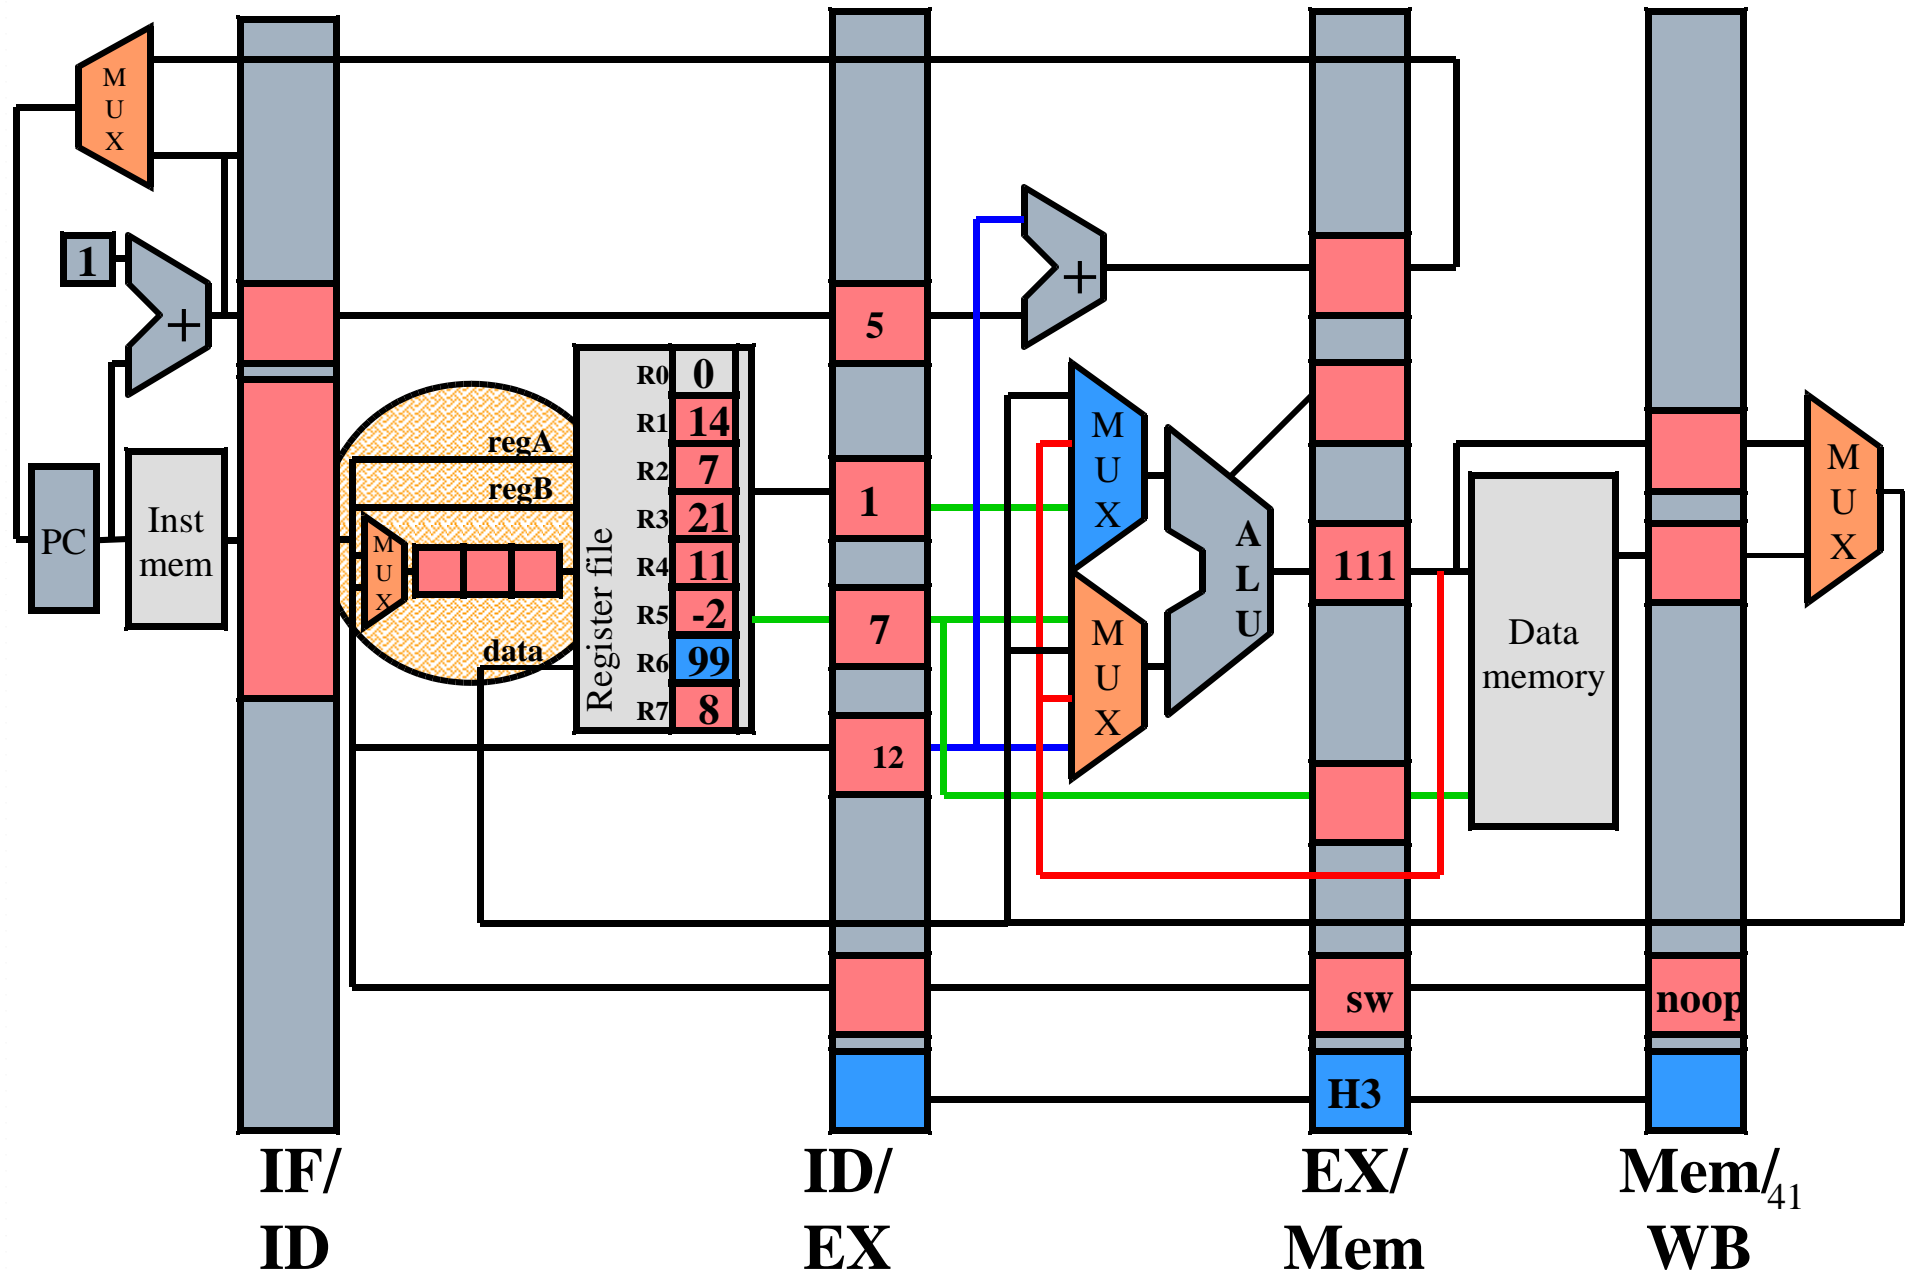
End of cycle 7



First half of cycle 8



End of cycle 8



Time Graph

Time:	1	2	3	4	5	6	7	8	9	10	11	12	13
add 1 2 3	IF	ID	EX	ME	WB								
nand 3 4 5		IF	ID	EX	ME	WB							
add 6 3 7			IF	ID	EX	ME	WB						
lw 3 6 10				IF	ID	EX	ME	WB					
sw 6 2 12					IF	no op	ID	EX	ME	WB			

Class Problem 2

```
add 1 2 3
lw 3 4 1
lw 4 5 6
add 6 1 7
sw 5 2 12
```

How many cycles to execute
this code using detect and stall?

How many cycles to execute
this code using detect and forward?

Next time

- ❑ Control hazards